



## Ahmedabad University

School of Engineering and Applied Science (SEAS)

Course Code: CSE540  
Cloud Computing

# Mini-Cloud Platform

## Scheduler, VM Creation & Live Migration

### Instructor

Professor Sanjay Chaudhary

### Group Members

Name	Roll No.	Email
Shlok Shelat	AU2240025	shlok.s1@ahduni.edu.in
Shrey Salvi	AU2240033	shrey.s2@ahduni.edu.in
Purvansh Desai	AU2240036	purvansh.d@ahduni.edu.in
Dhananjay Khanjariya	AU2240023	dhananjay.k@ahduni.edu.in

November 17, 2025

---

# Contents

<b>1 Abstract</b>	<b>1</b>
<b>Abstract</b>	<b>1</b>
<b>2 Project Statement</b>	<b>2</b>
2.1 Problem Definition . . . . .	2
2.2 Objectives . . . . .	2
2.3 Scope of the Project . . . . .	3
2.4 Challenges Addressed . . . . .	5
2.5 Motivation & Significance . . . . .	6
<b>3 System Overview</b>	<b>7</b>
3.1 What is the Mini-Cloud Platform? . . . . .	7
3.2 High-Level Functional Overview . . . . .	7
3.3 User Roles . . . . .	8
3.3.1 Administrator . . . . .	8
3.3.2 API Client . . . . .	9
3.3.3 Controller Service . . . . .	9
3.3.4 Celery Worker Nodes . . . . .	9
3.3.5 XCP-NG Hypervisor Hosts . . . . .	9
3.3.6 XOA (Xen-Orchestra) Backend . . . . .	10
3.4 Features Implemented . . . . .	10
3.4.1 Virtual Machine Creation . . . . .	10
3.4.2 Virtual Machine Management . . . . .	10
3.4.3 Template Handling . . . . .	11
3.4.4 Metrics Collection . . . . .	11
3.4.5 VM Migration . . . . .	11
3.4.6 Centralized API System . . . . .	12
3.4.7 Scheduler for Background Tasks . . . . .	12
3.4.8 Monitoring and Logging . . . . .	12
<b>4 Project Design</b>	<b>13</b>
4.1 Complete System Architecture . . . . .	13

4.1.1	Architectural Overview Diagram . . . . .	13
4.1.2	Component Responsibilities . . . . .	16
4.1.3	Data Flow Explanation . . . . .	17
4.2	Logical Design . . . . .	18
4.2.1	Class Diagrams . . . . .	19
4.2.2	Entity–Relationship (ER) Diagram for Database . . . . .	23
4.2.3	API Layer Logical Design . . . . .	26
4.2.4	Microservices Interaction Design . . . . .	31
4.3	Behavioral Design . . . . .	34
4.3.1	Sequence Diagrams . . . . .	34
4.3.2	Activity Diagrams . . . . .	37
4.3.3	State Diagrams . . . . .	40
4.4	Process Design & Flowcharts . . . . .	41
4.4.1	VM Creation Flowchart . . . . .	41
4.4.2	Metrics Collection Flowchart . . . . .	43
4.4.3	Migration Service Flowchart . . . . .	43
4.4.4	API Request Processing Flow . . . . .	44
4.5	Algorithms and Pseudocode (Detailed) . . . . .	45
4.5.1	VM Creation Algorithm . . . . .	46
4.5.2	Template Cloning Algorithm . . . . .	48
4.5.3	Metrics Fetching Algorithm . . . . .	48
4.5.4	Celery Task Scheduling Algorithm . . . . .	49
4.5.5	VM Migration Algorithm . . . . .	49
<b>5</b>	<b>Hardware and Software Specifications</b>	<b>51</b>
5.1	Hardware Setup . . . . .	51
5.1.1	XCP-NG Servers . . . . .	51
5.1.2	Host Specifications . . . . .	52
5.1.3	Network Topology . . . . .	52
5.1.4	Storage Environment . . . . .	53
5.2	Software Stack . . . . .	54
5.2.1	Operating Systems . . . . .	54
5.2.2	Python Version . . . . .	54
5.2.3	FastAPI, Celery, Redis, and PostgreSQL . . . . .	55
5.2.4	Docker and Docker-Compose Setup . . . . .	56

5.2.5	XOA / Xen-Orchestra Integration . . . . .	57
5.2.6	Monitoring Tools . . . . .	58
5.2.7	Third-Party Libraries . . . . .	58
<b>6</b>	<b>Cloud Services Categorization</b>	<b>59</b>
6.1	Infrastructure as a Service (IaaS) Components . . . . .	59
6.1.1	VM Provisioning . . . . .	59
6.1.2	Template Management . . . . .	59
6.1.3	Network and Storage . . . . .	60
6.2	Platform as a Service (PaaS) Components . . . . .	60
6.2.1	Celery Task Execution Platform . . . . .	60
6.2.2	API Management Layer . . . . .	61
6.2.3	Database Services . . . . .	61
6.3	Software as a Service (SaaS) Components . . . . .	61
6.3.1	Xen-Orchestra Web Interface (SaaS Functionality) . . . . .	62
6.3.2	Custom API Services as SaaS Endpoints . . . . .	62
<b>7</b>	<b>Implementation</b>	<b>63</b>
7.1	Setting Up the Environment . . . . .	63
7.1.1	XCP-NG Pool Setup . . . . .	63
7.1.2	Connecting XOA with API Tokens . . . . .	64
7.1.3	Dockerizing the Controller . . . . .	65
7.2	FastAPI Implementation . . . . .	65
7.2.1	API Endpoints . . . . .	66
7.2.2	Authentication and Authorization . . . . .	66
7.2.3	Models and Schemas . . . . .	67
7.3	Celery + Redis Integration . . . . .	67
7.3.1	Worker Tasks . . . . .	68
7.3.2	Scheduling Tasks . . . . .	68
7.3.3	Error Handling . . . . .	68
7.4	PostgreSQL Database Implementation . . . . .	69
7.4.1	Schema Design . . . . .	69
7.4.2	Session Management . . . . .	70
7.5	Complete VM Creation Pipeline . . . . .	70
7.5.1	Using SSH and xe CLI . . . . .	70
7.5.2	Template UUID Handling . . . . .	70

7.5.3	IP Assignment . . . . .	71
7.5.4	Attaching Network . . . . .	71
7.6	Metrics Collection Service . . . . .	71
7.6.1	Using XOA Metrics API . . . . .	72
7.6.2	Storing Metrics in DB . . . . .	72
7.6.3	Dashboard Options . . . . .	72
7.7	VM Migration Service . . . . .	73
7.7.1	Design . . . . .	73
7.7.2	Payload Format . . . . .	74
7.7.3	XOA Migration API . . . . .	75
7.7.4	Cross-Host Migration . . . . .	75
<b>8</b>	<b>Sample Source Code</b>	<b>77</b>
8.1	FastAPI API Endpoints . . . . .	77
8.2	Celery Task Code . . . . .	77
8.3	VM Creation Script (Excerpt) . . . . .	78
8.4	Metrics Fetching Code . . . . .	78
8.5	Migration Code (Excerpt) . . . . .	79
8.6	Docker Compose File (Excerpt) . . . . .	79
8.7	Database Models (Excerpt) . . . . .	80
<b>9</b>	<b>Results and Discussions</b>	<b>82</b>
9.1	Output Screenshots and System Logs . . . . .	82
9.1.1	XOA VM Creation Output . . . . .	82
9.1.2	Controller Logs . . . . .	84
9.1.3	Worker Task Completion Logs . . . . .	84
9.1.4	Metrics Data Example . . . . .	85
9.1.5	Migration Success Proof . . . . .	85
9.2	Performance Evaluation . . . . .	86
9.2.1	VM Creation Time . . . . .	86
9.2.2	Metrics Fetching Latency . . . . .	86
9.2.3	API Response Times . . . . .	87
9.3	Comparison with Existing Tools . . . . .	87
9.4	Limitations . . . . .	87
<b>10</b>	<b>Conclusion</b>	<b>89</b>

10.1 Summary . . . . .	89
10.2 Key Learnings . . . . .	89
10.3 Achievements . . . . .	90
10.4 Future Work . . . . .	90
10.5 Final Remarks . . . . .	91
A.1 create_vm_remote_fixed.sh . . . . .	92
A.2 app/scheduler.py . . . . .	99
A.3 app/models.py . . . . .	108
A.4 app/db.py . . . . .	109
A.5 app/metrics_utils.py . . . . .	109
A.6 app/main.py . . . . .	110
A.7 app/api/apiVm.py . . . . .	111
A.8 app/api/apiMigration.py . . . . .	112
A.9 app/api/apiMetrics.py . . . . .	112
A.10 app/tasks/celeryApp.py . . . . .	113
A.11 app/tasks/jobs.py . . . . .	114
A.12 app/tasks/collectMetrics.py . . . . .	115
A.13 app/xoaClient.py . . . . .	116
A.14 Miscellaneous utilities . . . . .	117
A.15 docker-compose.yml . . . . .	117
A.16 Dockerfile . . . . .	118

# List of Figures

4.1	High-Level Architecture Diagram of the Mini-Cloud Platform . . . . .	13
4.2	Layer-wise Architecture of the Mini-Cloud Platform . . . . .	14
4.3	Class Diagram of the Mini-Cloud Platform . . . . .	23
4.4	Entity–Relationship Diagram for the Mini-Cloud Platform . . . . .	26
4.5	API Layer Logical Flow: Sequence Diagram Representing Request-to-Execution Flow . . . . .	30
4.6	API Layer Logical Component Diagram Showing Internal System Interactions	31
4.7	Microservices Interaction Diagram for the Mini-Cloud Platform . . . . .	33
4.8	VM Creation Sequence Diagram . . . . .	35
4.9	VM Migration Sequence Diagram . . . . .	36
4.10	Metrics Collection Sequence Diagram . . . . .	36
4.11	Task Scheduling Sequence Diagram . . . . .	37
4.12	Activity Diagram: VM Request Flow . . . . .	38
4.13	Activity Diagram: Monitoring / Metrics Flow . . . . .	39
4.14	Activity Diagram: Asynchronous Task Execution Flow . . . . .	40
4.15	State Diagram: VM Lifecycle . . . . .	41
4.16	VM Creation Flowchart . . . . .	42
4.17	Metrics Collection Flowchart . . . . .	43
4.18	VM Migration Service Flowchart . . . . .	44
4.19	API Request Processing Flowchart . . . . .	45
5.1	XCP-NG hypervisor nodes detected in Xen-Orchestra. . . . .	51
5.2	Xen-Orchestra dashboard showing RAM, CPU, storage, and pool-level statistics. . . . .	52
5.3	Shared NFS Storage Repository used by both XCP-NG hypervisors. . . . .	53
5.4	Xen-Orchestra VM list showing active VMs and templates. . . . .	54
5.5	FastAPI OpenAPI documentation showcasing all exposed endpoints. . . . .	55
5.6	Celery Flower dashboard showing task queues and worker health. . . . .	56
5.7	pgAdmin view of metrics, events, and VM tables used by the controller. . . . .	56
5.8	XOA Login interface used for cluster and resource orchestration. . . . .	57
7.1	High-level design of the VM Migration Service. . . . .	74

# 1. Abstract

The rapid expansion of virtualization and cloud computing has created a growing need for lightweight, customizable, and cost-efficient private cloud platforms that educational institutions and small organizations can deploy without the complexity and overhead of full-scale commercial cloud systems. This project presents the design and implementation of a Mini-Cloud Platform built using XCP-NG, Xen-Orchestra (XOA), FastAPI, Celery, Redis, and PostgreSQL, enabling automated virtual machine provisioning, VM lifecycle management, asynchronous task execution, metrics collection, and VM migration through a unified controller.

The main goal of the project is to develop a full-fledged API-based, fully automated virtualization controller, which will be able to communicate with an XCP-NG pool to create, manage, and monitor virtual machines. The system is comprised of a RESTful API layer (FastAPI) and an asynchronous distributed task execution engine (Celery + Redis) and a relational database backend (PostgreSQL) to ensure reliable orchestration. The API and command-line features of Xen-Orchestra support VM tasks including template cloning, network attachment, power management, getting metrics and cross-host migrations.

Its architecture is a microservices-based design that is modular, consisting of a controller service, Celery worker cluster, Redis message broker, a PostgreSQL persistence layer, and the underlying XCP-NG infrastructure. The API server receives VM creation requests and forwards it to Celery workers which asynchronously executes provisioning tasks. A metrics-collection module is used for accessing system performance metrics (CPU, RAM, disk usage, network I/O) to XOA periodically, which are saved within the PostgreSQL database. The project also deploys a migration service through which a host-to-host live migration can be done seamlessly and using the task definitions of XOA.

The adoption of the Mini-Cloud Platform meant a stable and production-like environment with the ability to: provision VMs automatically using templates, process jobs asynchronously with stable and dependable task states, schedule metrics collection and storage, execute live VM migrations in structured API payloads and log and monitor all system components centrally.

The system is highly automated, and it is easy to use and extend. The main results are less manual intervention in the process of virtualization, standardized cloud-like behavior in the management of VM, and a scalable architecture, which may be extended to meet future improvements, like that of autoscaling, integration into a UI, or container act. The project shows that an academic value of open-source technologies allows to construct an effective lightweight private cloud environment and can be deployed into practical use.

# **2. Project Statement**

## **2.1 Problem Definition**

As organizations continue to embrace virtualization and cloud computing, there is need to have lightweight and customized private cloud solutions (especially in case of the academic institutions and small scale businesses). Nevertheless, the currently available enterprise cloud infrastructures including VMware vSphere, Proxmox VE, or OpenStack have numerous limitations: they are expensive, require a lot of time to install, rely on a wide range of services, and have steep learning curves. Likewise, as XCP-NG and Xen-Orchestra (XOA) provide a powerful virtualization stack, the experience of managing virtual machines (VMs) can be tedious and manual, where the administrators have to perform numerous tasks manually like:

- Create virtual machines using the dashboard from templates by hand,
- Observe monitoring metrics from various systems,
- Perform host-to-host migrations manually, (or maybe host machine-based),
- Check VM status by referencing log files or the XOA UI, and
- Complete provisioning tasks for individual machines.

When a number of virtual machines must be created, managed, and monitored concurrently, this is cumbersome. This frequently entails new delays, unpredictable configurations, and a greater likelihood of human error on the part of the system administration team.

As a consequence, there is a lack of centralized, automated, API-driven orchestration that is capable of integrating seamlessly with XCP-NG/XOA workflow covering VM lifecycle operations, metrics collection, task scheduling and migration without human action.

## **2.2 Objectives**

The main goal of this project is to create and deploy a Mini-Cloud Platform that automates virtualization workflow processes within an XCP-NG environment using contemporary cloud-native tools and techniques. The objective of the system is to provide a fully automated controller that communicates with the virtualization layer through well-defined APIs and asynchronous background workers.

The key objectives are:

## Technical Objectives

1. Create a FastAPI based rest controller to provide secure and modular API endpoints to provision, manage, monitor and migrate VMs.
2. Automate VM lifecycle actions, such as template creation, power, network connection, disk setups and destruction.
3. Utilize Celery in conjunction with Redis to run long-running functions in an asynchronous manner without blocking APIs.
4. Incorporate a persistence layer built on top of PostgreSQL to store metadata of VMs, logs of metrics, task requests and logs at the system level.
5. Install a Metrics Collection Module that periodically retrieves VM performance data off of XOA and then archives it to be analyzed.
6. Use VM Migration Service so that controlled and structured migrations can take place through XOA live migration APIs.
7. Design a microservices based architecture, which is scalable, modular, and enables the controller, workers, database, and the broker to scale independently.

## Research and Learning Objectives

- Comprehend current cloud architectures and orchestration concepts.
- Obtain practical experience with virtualization platforms (i.e. XCP-NG, XOA). Gain hands-on experience with virtualization platforms (XCP-NG, XOA).
- Execute asynchronous programming, distributed task queues, and event-driven architecture.
- Analyze the pragmatics of cloud Infrastructure as a Service automation and DevOps methodologies.

### 2.3 Scope of the Project

The ultimate goal of the Mini-Cloud Platform is to develop an automated and reliable layer for managing virtualization. The scope of the project entails:

## In-Scope Features

- VM creation with XCP-NG templates with API.
- Managing the lifecycle of VMs (Start, Stop, Reboot, Shutdown)
- Running long-running tasks in the background asynchronously.
- Database that tracks VM records and metrics that are centralized.
- Performance measures (CPU, RAM, Disk, Network) will be taken periodically
- XOA API live migration based on cross-host VM
- All the critical system operations should be logged and monitored
- The entire controller stack runs in a decoupled and dockerized format.

## Technologies In Scope

- XCP-NG and Xen-Orchestra
- FastAPI
- Celery & Redis
- PostgreSQL
- Docker & Docker Compose
- Python 3.x

## Out-of-Scope (Future Work)

The project has deliberately taken out advanced features based on time and resource limitations, as for example:

- Multi-tenant dashboard UI
- Automated autoscaling based on metrics
- Integration with Kubernetes
- Distributed storage provisioning
- User-level cloud billing system
- High-availability cluster for the controller

These can be added in future features.

## 2.4 Challenges Addressed

The project focused on mostly real-world issues with typical virtualization and cloud management:

### 1. Manual VM Provisioning

Multiple times with a lot of manual work, like cloning templates, connecting networks, and configuring resources is laborious. **Approach:** Fully automated VM creation workflow through FastAPI + Celery tasks.

### 2. Lack of Centralized Orchestration

XOA has a UI but does not provide a customizable orchestration layer for batching an operation or custom automation. **Approach:** A centralized API-controlled controller to manage VM life cycle, migrations, and metrics.

### 3. Synchronous Execution Limitations

Heavy operations like clones or migrations can't be run through normal APIs in a synchronous manner. **Approach:** Celery distributed workers with Redis message broker to execute asynchronous scalable processes.

### 4. Metrics Monitoring Complexity

To view the performance of a VM requires navigating through the UI of XOA manually. **Approach:** An automated scheduler collects and stores metrics periodically into a PostgreSQL database.

### 5. Migration Coordination

Live migrations are still manual to select the target host and storage repositories. **Approach:** Automated migration service using structured API payload and executing tasks through XOA.

### 6. Error Handling & Reliability

General CLI operations will not always provide meaningful logs instead failing silently without monitoring in a central location. **Approach:** Logging retries and structured exception handling inside Celery workers.

## 7. Scalability Limitations

As the number of VMs increases, manual processes become unmanageable. **Approach:** Horizontally scalable architecture with stateless controllers and distributed workers.

## 2.5 Motivation & Significance

### Motivation

The modern IT infrastructure is supported by cloud computing. Nevertheless, a lot of institutions do not have the infrastructure or funding to implement full-scale implementations of private clouds. The need in this project was based on the need:

- A cheap and easy-to-use personal cloud service.
- Practical understanding and exposure to real-life virtualization technologies.
- Everyday administration is to be automated to minimize human error.
- Recreating hyperscale cloud behavior (provisioning, metrics, migration) with open-source software
- Writing a complete programmable interface on a virtualization platform.

### Significance

The project has a great academic and practical significance:

- **Educational Importance:** It offers a full, real world implementation of the inner working of cloud provisioning systems.
- **Technological Impact:** Demonstrates how the open-source virtualization tools can be improved so that they act as a complete private cloud.
- **Operational Efficiency:** It lowers the administrative burden and fastens the VM deployments significantly.
- **Extensibility:** Architecture may be extended to allow autoscaling, dashboards, container support, etc.
- **Practical Deployment Use:** It can be used in research laboratories, academic institutions, as well as small data centers.

In this project, we show that a mini-cloud environment that is well structured, automated, based on API, and built with heavyweight cloud platforms is affordable.

# 3. System Overview

## 3.1 What is the Mini-Cloud Platform?

Mini-Cloud Platform is a small, modular and automated private cloud platform based on XCP-NG and Xen-Orchestra (XOA). It is also designed to offer cloud-like functionalities like automated provisioning of virtual machines, VM lifecycle, resource monitoring, and live migration — but without complicating the matter, flexibility, and simplicity, and cost efficiency.

In contrast to conventional enterprise cloud systems (e.g. OpenStack, VMware ESXi), the MiniCloud Platform adds a new layer of custom controller which brings out programmable APIs to coordinate with the virtualization infrastructure. This separates user interactions access to the hypervisor underneath it and enables the automation of administrators or applications and workflows programmatically.

The platform combines a number of modern technologies:

- FastAPI to develop a high-performance RESTful API layer.
- Celery and Redis asynchronous and distributed tasks.
- PostgreSQL as a persistent storage of metadata and metrics.
- XCP-NG and XOA VM operations including creation, metrics and migration

These elements build a united and scalable ecosystem that transforms a conventional virtualization configuration into a programmable mini-cloud infrastructure.

## 3.2 High-Level Functional Overview

The Mini-Cloud Platform is on a high level of operation in that it is an orchestration and automation layer on the XCP-NG hypervisor pool. The system is a modular architecture with every element being devoted to a particular functionality:

- **FastAPI Controller:** Acts as the entry point for all user or system requests. There are API endpoints for actions like creating a VM, deleting a VM, migrating one, and getting metrics.
- **Celery Workers (Async Task System):** It processes long running or intensive actions (for example, cloning a VM template, migrating VMs) in a way that does not block API responses. Actions are performed in parallel across multiple workers.

- **Redis Message Broker:** This acts as a communication layer between the API and Celery workers. It queues and schedules tasks in a deterministic order.
- **PostgreSQL Database:** Stores metadata of VM, logs of API request, task state and performance metrics of XOA.
- **XCP-NG & XOA Integration Layer:** offers the real virtualization functions:
  - Clone from template
  - Attach to networks
  - Power related operations
  - Retrieval of metrics
  - Executing live migrations operations
- **Scheduler Tasks:** Run some background tasks periodically (for example, some metrics (like every minute), or cleanup tasks (like hourly))
- **Logging and Monitoring Framework:** Logs controller activity, worker execution logs, API requests, response and error traces.

When the user/service sends request (For example to create a VM), the Controller will validate the input, call a celery task, initiate the VM provisioning process via the XCP-NG/XOA connectors, and store status and metadata in the database. All this happens automatically.

### 3.3 User Roles

The Mini-Cloud Platform is principally an orchestration platform; however, it has implicit definitions of user roles and system actors:

#### 3.3.1 Administrator

- Administers the XCP-NG servers underneath
- Manages the XOA instance
- Controls network, storage and host resources
- Utilizes the platform to create VMs for lab or production deployments

### 3.3.2 API Client

All services, automation scripts or user applications that communicate with the FastAPI controller are an API Client.

This role:

- Makes API calls to create, start, stop or migrate a VM.
- Get metrics or VM information.
- Consolidates platform with other tools (CI/CD, dashboards, etc.)

### 3.3.3 Controller Service

The rational player of the platform. It:

- Retrieves incoming requests
- Checks and validates parameters
- Drives work to Celery employees.
- Stores data in PostgreSQL database

### 3.3.4 Celery Worker Nodes

The functions of these nodes are:

- Performing long and heavy-weight processes.
- Re-attempting failed tasks
- Communicating the results of the tasks to the controller.

They can be scaled horizontally and that is, they can be executed on multiple machines by numerous worker processes.

### 3.3.5 XCP-NG Hypervisor Hosts

These servers run and perform the real virtualization methods:

- Running VMs
- Cloning templates
- Handling migrations
- Providing compute, network, and storage resources

### 3.3.6 XOA (Xen-Orchestra) Backend

Acts as an API gateway for:

- Retrieval of VM related metrics
- Live migration Orchestration.
- Media network and storage management.

## 3.4 Features Implemented

The Mini-Cloud Platform has a wide range of cloud-like features, which allows a seamless and automated workflow virtualization.

### 3.4.1 Virtual Machine Creation

- Automatic provisioning of VMs using pre-derived templates.
- Allows being specific with CPU, RAM, disk size and network.
- operates through XOA/XCP-NG API and xe CLI.
- Responsiveness is provided by asynchronous execution.
- VM Save VM information in PostgreSQL to use later.

### 3.4.2 Virtual Machine Management

Covers full lifecycle operations:

- Start
- Stop
- Suspend
- Resume
- Reboot
- Delete

These tasks are implemented via the controller with pre-installed logging as well as error handling.

### 3.4.3 Template Handling

- Displays the VM templates that are available in XCP-NG.
- Clones VMs from templates selected by the user.
- Validates and maps template UUIDs correctly.
- Aids in cloning standardized lab or development systems.

### 3.4.4 Metrics Collection

A dedicated module fetches VM metrics at regular intervals:

- CPU utilization
- Memory usage
- Disk I/O
- Network I/O

These metrics are:

- Data is collected through XOA's metrics API
- The data is stored in PostgreSQL
- The data can be presented for visualization or analytics

### 3.4.5 VM Migration

Facilitates live migration and storage-aware migration with XOA APIs:

- Automatically chooses or accepts target hosts
- Supports block migrations
- Achieves minimal downtime and data consistency
- Applicable for load balancing, tiresome maintenance, or general performance tuning

### 3.4.6 Centralized API System

The framework exposes:

- Clean and structured documented API endpoints
- JSON input and output envelopes
- Streaming speed in execution based on FastAPI
- Authentication hooks (optional for future extensions)

System APIs abstract the entire virtualization stack to become a programmable cloud interface.

### 3.4.7 Scheduler for Background Tasks

Includes scheduled jobs:

- For unattended periodic collection of VM metrics
- Automatic cleanup of ephemeral resources
- Daily or hourly health checks on the stack
- Background logging and reporting jobs

Jobs work in tandem to create the appearance of a micro-cloud management platform.

### 3.4.8 Monitoring and Logging

Comprehensive logging is executed throughout the framework:

- API interaction logs
- Celery worker execution logs
- Errors/exceptions logs
- VM provisioning and migration history logs

Log access creates transparency, debug-ability, and operational observability.

# 4. Project Design

## 4.1 Complete System Architecture

Mini-Cloud Platform design is based on the creation of a scalable, automated and modular cloud orchestration platform. The project is based on a multi-layered architecture which has separated the user interface, controller logic, asynchronous processing, virtualization management, compute operations, and storage access. The essence of the design philosophy is to make sure that each subsystem is able to scale automatically, is loosely coupled and is resilient even when it is heavily loaded.. The system architecture of the system is divided into two complementary diagram perspectives (1)A high-level architecture diagram of the system that models the end-to-end process flow of VM provisioning, scheduling, and hypervisor operations; and (2) A layer-wise architecture diagram of the system that sub-divides its system system into functional layers including User Layer, Controller Layer, Orchestration Layer, Compute Layer, and Storage Layer.

### 4.1.1 Architectural Overview Diagram

The following diagrams represent both perspectives of the system:

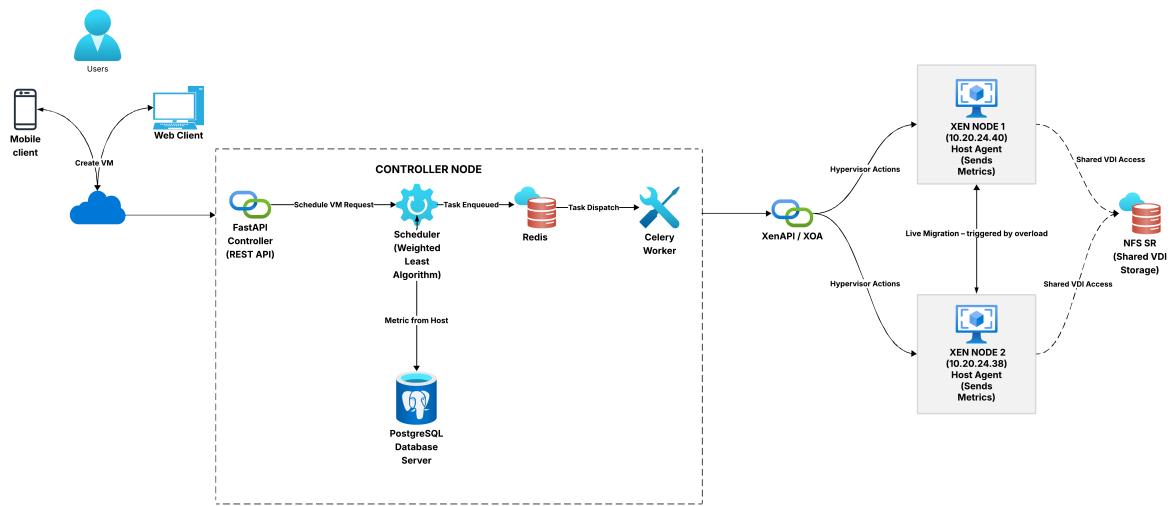


Figure 4.1: High-Level Architecture Diagram of the Mini-Cloud Platform

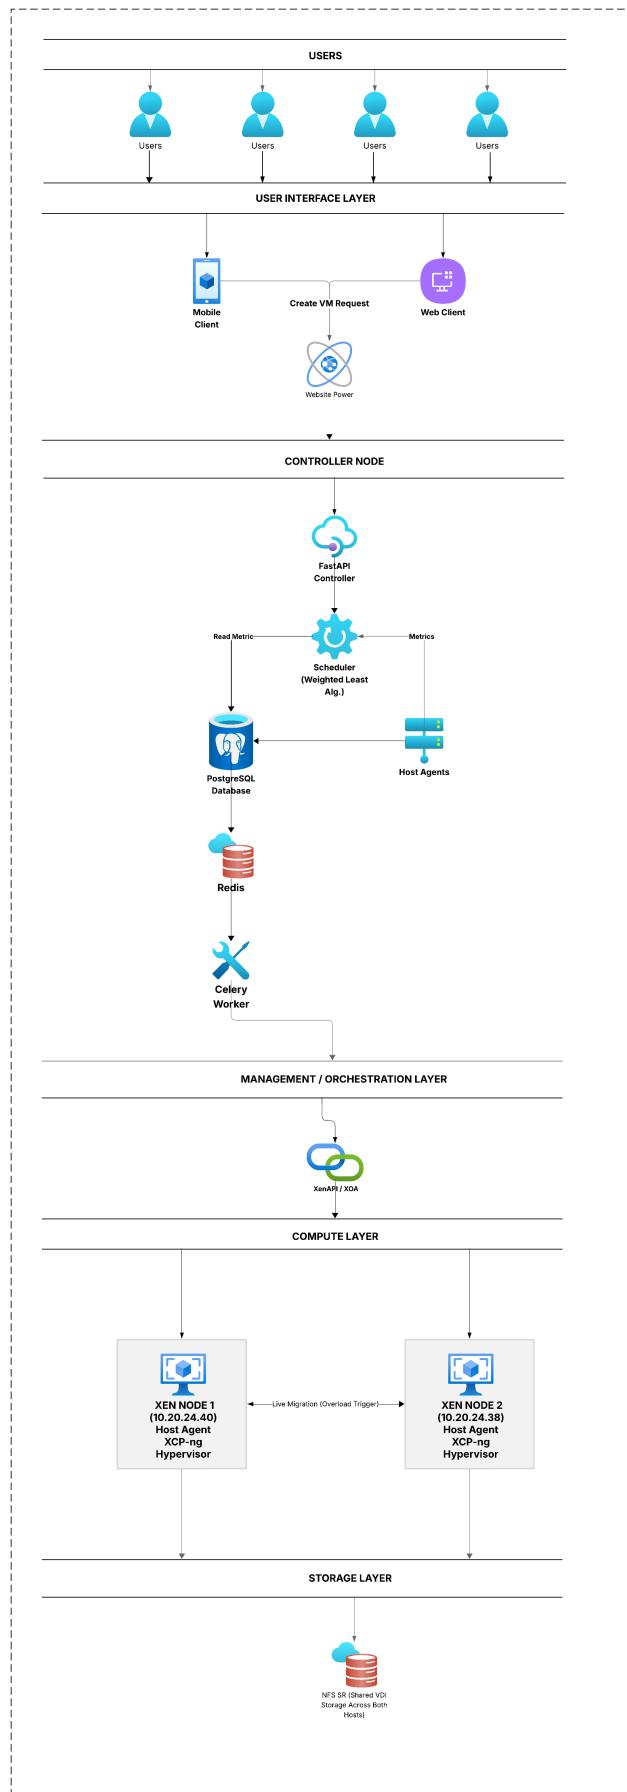


Figure 4.2: Layer-wise Architecture of the Mini-Cloud Platform 14

Your system architecture consists of six distinct layers, each with a specific purpose. Below is the detailed explanation that corresponds to the layer-wise diagram:

### (A) Users Layer

This layer consists of the end-users who request deployment or management of virtual machines. These users may be students, researchers, system administrators, or automation scripts. Users do not interact with the hypervisors directly; all interactions go through the API.

### (B) User Interface Layer

This layer handles request submission to the cloud platform. It may include a mobile client, web client, or a future dashboard. Users submit actions such as creating VMs, deleting VMs, viewing VM status, or initiating migrations. The UI sends HTTP API requests and contains no business logic.

### (C) Controller Node Layer

This is the brain of the Mini-Cloud platform. It contains:

- **FastAPI Controller (REST API)** – Primary point of all operations. It authenticates API requests, captures metadata, and sends Celery jobs
- **Scheduler (Weighted Least Algorithm)** – Chooses the least loaded hypervisor based on PostgreSQL host metrics.
- **PostgreSQL Database** – Metadata of VM, metrics, logs, task entries, and migration history.
- **Redis** – Celery task queuing message broker.
- **Celery Worker** – This performs heavy tasks like VM creation, deletion, metricals, and migration.

### (D) Management / Orchestration Layer

the layer that handles management and coordination of various system layers. The coordination and management of different layers of the system is also dealt with in the management / Orchestration Layer. Its main constituent is the XenAPI/XOA, which runs the operations of VM lifecycle, clones templates, retrieves metrics and both live VM migration..

### (E) Compute Layer (Hypervisor Layer)

XCP-ng hypervisors can be found in this layer:

- XEN Node 1 (10.20.24.40)
- XEN Node 2 (10.20.24.38)

Both nodes have virtual machines and offer compute and networking capabilities and report the performance to the controller.

#### (F) Storage Layer

In this layer, there is the NFS Storage Repository (Shared VDIStorage) that enables the shared block storage between all hypervisors. Shared storage facilitates live migration that eliminates the need to copy disk information between hosts.

### 4.1.2 Component Responsibilities

Below is the detailed breakdown of responsibilities for each component:

#### FastAPI Controller

- Checks and authenticates the incoming API requests.
- The external interface of the cloud.
- Sends long term operations to Celery workers.
- Writes and reads VM metadata in PostgreSQL.

#### Scheduler (Weighted Least Algorithm)

- Reads contain PostgreSQL metrics.
- Calculates the load based on weighted variables: CPU, RAM, network, running VMs.
- Precludes hypervisor congestion through the ideal host selection.

#### Redis Message Broker

- Queues asynchronous tasks.
- Makes implementation not blocking.
- Favors job distribution that is scalable.

#### Celery Worker

- Performs processor-intensive and lengthy tasks.
- Provides VM provisioning, migrates, deletion and metrics.
- Conducts error management and logs.

## XenAPI / XOA

- Carries out VM operations like start and lead, re-booth and clone.
- Administrates real-time migration of people.
- Directs metrics and information about hypervisors to the controller.

## Hypervisor Nodes

- Hypervisor- virtual Governor Run virtual machines and assign resources (CPU, RAM, and network).
- Send controller report measures through the host agents.
- Perform live migrations.

## PostgreSQL Database

- Storage Service VM metadata, migration records, history of metrics, template data, etc.
- Ensures analytics and dashboards.

## NFS Storage Repository

- Offers common VDI storage that can be used by a number of hypervisors.
- Allows real-time migration of VMs.

### 4.1.3 Data Flow Explanation

#### (A) VM Creation Request Flow

1. User sends Create VM request on UI to FastAPI.
2. The FastAPI verifies the request and sends it forward to the Scheduler.
3. Scheduler asks PostgreSQL to give it host metrics.
4. Scheduler uses weighted least algorithm to choose a host of optimal choice.
5. Scheduler adds job to Redis queue.
6. Celery Worker removes the tasks and interacts with XenAPI/XOA.

7. XOA does VM provisioning in the chosen hypervisor.
8. Commit Metadata PostgreSQL is updated by the worker.
9. Success response is sent to the user by API.

#### (B) Metrics Collection Flow

1. Open-ended metrics are issued to Scheduler by Hypervisor Host Agents.
2. Scheduler keeps the measures in PostgreSQL.
3. Celery Workers get the metrics of XOA at predetermined timelines.
4. Analytics are logged in metrics.

#### (C) VM Migration Flow

1. Scheduler uses metrics to detect overloaded host.
2. Celery migration task is triggered by a migration API request.
3. Celery Worker orders the migration of VM to a different host.
4. In hypervisor, migration is performed live with the help of shared storage.
5. The migration data is logged into PostgreSQL.

#### (D) Storage Access Flow

The Four hypervisors share the same NFS Shared Repository which supports:

- Shared access to VDI files
- Live migrations with near to zero-downtime.
- Stability: VM failover/movement.

## 4.2 Logical Design

The logical design phase explains the inside design of the Mini-Cloud Platform in the form of classes, relationships, data models, and major structural elements of the architecture. It is the way the system is internally arranged - regardless of its physical implementation. Logical design gives the assurance that every single component has its assigned responsibilities, well-defined input and output, and the way these other components can interact.

The subsection deals with an object-oriented abstraction of the platform through Class Diagram. The class diagram reflects the key structures that are engaged in API processing, asynchronous task execution, VM settlement, metrics acquisition, planning, and database communications. Every class contains a collection of duties, qualities and roles that jointly assigns the aggregate capability of the cloud orchestration framework.

#### 4.2.1 Class Diagrams

The Class Diagram of the Mini-Cloud Platform is further subdivided into four conceptual clusters including API Layer Classes, Scheduler Layer Classes, Worker Layer Classes and Database Layer Classes. The combination of these classes determines the gist of the logic, and interactions amongst the controller and the scheduler, workers and the underlying infrastructure.

##### API Layer Classes

**FastAPIController** It is a class that processes incoming HTTP requests and has an incoming request body which is validated and interacts with the Scheduler and allocates work to the Celery Worker, and makes database operations.

##### Responsibilities:

- Accept as well as authorize API requests.
- Map requests into internal operations.
- Dispatch Celery tasks
- Update and query PostgreSQL database

##### Key Methods:

- `create_vm()`
- `delete_vm()`
- `migrate_vm()`
- `get_vm_status()`
- `get_metrics(vm_id)`

**RequestSchema (Pydantic Models)** Defines the structure of incoming JSON requests for VM creation, deletion, migration, or metrics queries.

##### Attributes:

- VM name
- CPU count
- RAM size
- Disk size
- Template ID
- Host (optional)

## Scheduler Layer Classes

**Scheduler** This class implements the Weighted Least Algorithm, responsible for selecting the most optimal hypervisor for VM creation.

### Attributes:

- host\_metrics
- weight\_cpu
- weight\_ram
- weight\_network
- weight\_vm\_count

### Methods:

- `select_best_host()`
- `fetch_latest_metrics()`
- `calculate_weighted_score()`

**HostMetrics** Represents the metrics collected from hypervisors. These values are used by the Scheduler to make host selection decisions.

### Attributes:

- cpu\_usage
- ram\_usage
- network\_io
- running\_vms

## Worker Layer Classes

**CeleryWorker** Executes background tasks using Celery Workers. These tasks include VM provisioning, migration, metrics collection, and deletion. **Methods include:**

- `task_create_vm()`
- `task_migrate_vm()`
- `task_collect_metrics()`
- `task_delete_vm()`

**XenAPIClient** Responsible for communication with XCP-ng/XOA for performing hypervisor actions.

**Methods:**

- `clone_template()`
- `configure_vm_resources()`
- `attach_network()`
- `power_on_vm()`
- `migrate_vm()`
- `fetch_metrics()`

## Database Layer Classes (SQLAlchemy Models)

**VMMModel** Stores metadata related to each virtual machine.

**Attributes:**

- `vm_uuid`
- `template_id`
- `cpu`
- `ram`
- `disk`
- `host`
- `status`

**MetricsModel** Stores performance metrics collected periodically.

**Attributes:**

- cpu\_usage
- ram\_usage
- disk\_io
- network\_io
- timestamp

**TaskLogModel** Represents execution logs for Celery tasks.

**Attributes:**

- task\_id
- vm\_id
- task\_type
- status
- error\_message
- timestamp

## Class Diagram Placeholder

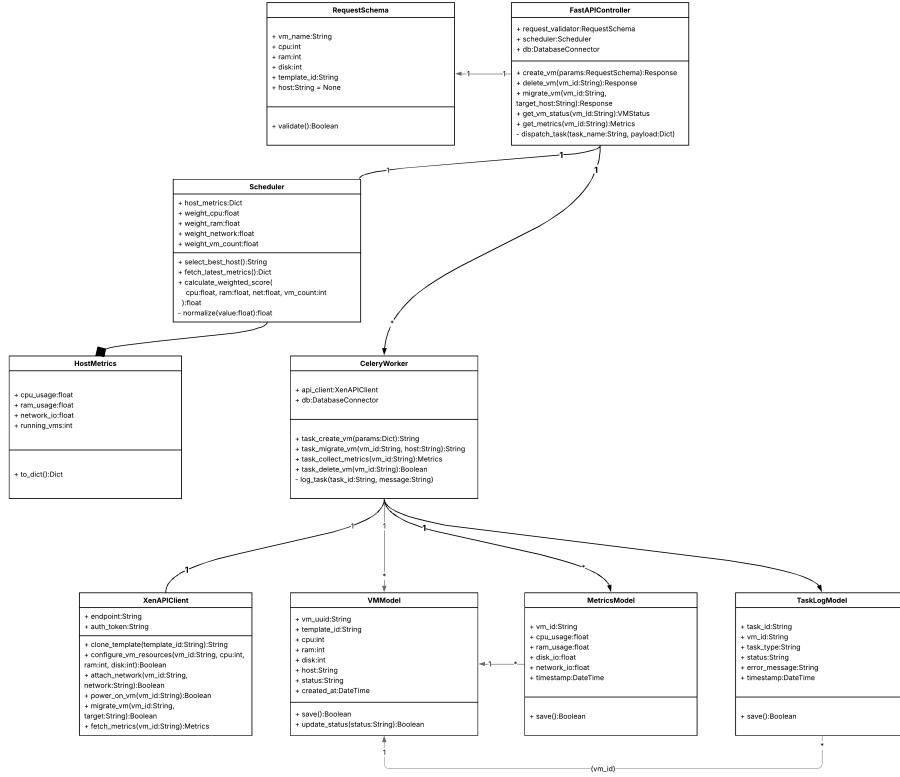


Figure 4.3: Class Diagram of the Mini-Cloud Platform

### 4.2.2 Entity–Relationship (ER) Diagram for Database

Entity Relationship (ER) model depicts the logical design of the database of the Mini-Cloud Platform. It depicts the relationship between the basis data entities, which are the virtual machines, performance measurements, task records and VM templates. This ER architecture allows the storage, access, and management of metadata created during the VM lifecycle, background tasks processing, metrics gathering and migration activities.

The ER diagram documents the essence of the interrelationship between the entities in the form of primary and foreign keys and cardinality constraints which would guarantee referential integrity throughout the system.

#### Entities and Attributes

- VMModel (Virtual Machine Entity)** Refers to a VM that is being created in the platform.

- **vm\_uuid (PK)** – Unique identifier for the VM.
- template\_id – Reference to the VM template used.
- cpu – Allocated vCPUs.
- ram – Allocated RAM in MB.
- disk – Disk size in GB.
- host – Host hypervisor on which the VM is deployed.
- status – Current VM state (running, stopped, suspended, migrating).
- created\_at – Timestamp of VM creation.

**2. MetricsModel (Performance Metrics Entity)** Stores time-series metrics for each VM.

- **id (PK)**
- **vm\_id (FK)** – References VMMModel.vm\_uuid
- cpu\_usage – CPU utilization percentage.
- ram\_usage – RAM usage in MB.
- disk\_io – Disk I/O throughput.
- network\_io – Network utilization.
- timestamp – Time at which the metric was recorded.

**3. TaskLogModel (Task Execution Logs Entity)** Represents logs for Celery worker task executions.

- **id (PK)**
- **vm\_id (FK)** – References VMMModel.vm\_uuid
- task\_type – Type of task executed (create, delete, migrate, metrics).
- status – Task status (success, failure, running).
- error\_message – Message associated with failures.
- timestamp – Execution timestamp.

**4. TemplateModel (Optional)** Represents VM templates stored or registered within the environment.

- **template\_id (PK)**
- template\_name
- os\_type

### Relationships and Cardinalities

- **VMMModel (1) — (M) MetricsModel** Each VM can have many performance metric entries collected over time. The foreign key `vm_id` in MetricsModel references `vm_uuid` in VMMModel.
- **VMMModel (1) — (M) TaskLogModel** Each VM generates many task logs such as creation, migration, deletion, or metrics collection tasks. The foreign key `vm_id` in TaskLogModel references `vm_uuid` in VMMModel.
- **TemplateModel (1) — (M) VMMModel** (if template tracking is enabled) Multiple VMs can be created using the same template. The attribute `template_id` in VMMModel references TemplateModel.template\_id.

Together, these entities form the core of the database system supporting the Mini-Cloud orchestration platform. The design ensures high consistency, efficient querying, and clear traceability between VM lifecycle operations and their associated metrics and logs.

## ER Diagram Placeholder

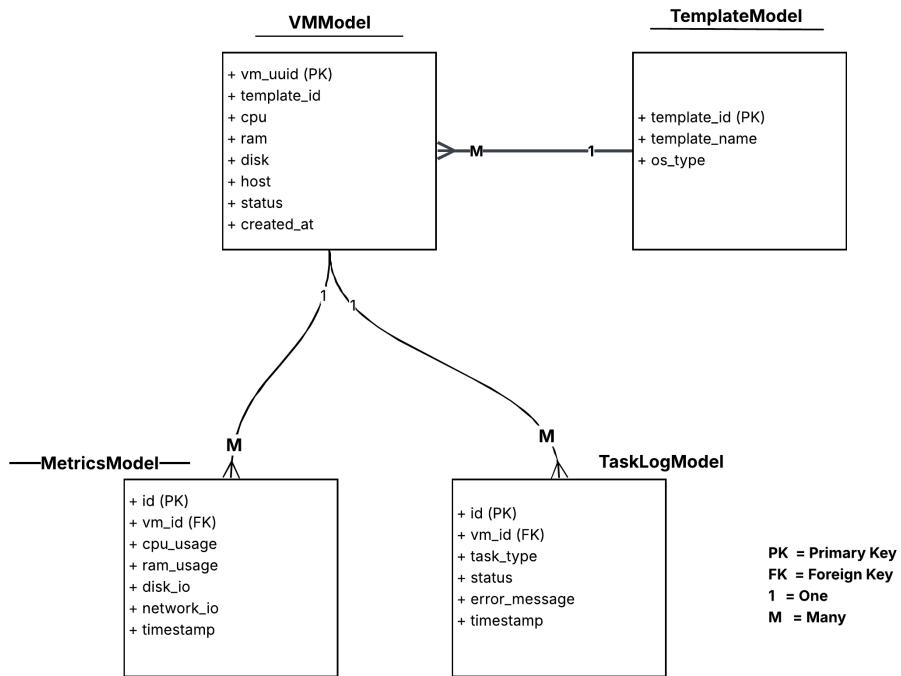


Figure 4.4: Entity–Relationship Diagram for the Mini-Cloud Platform

### 4.2.3 API Layer Logical Design

The API Layer is the side that leads to the Mini-Cloud Platform. It isolates the complexity of the direct interaction relative to the hypervisor and offers a programmable, clean and secure interface to the creation, management, and monitoring of the virtual machines. The layer is implemented on the FastAPI design pattern that adheres to the REST design pattern that guarantees high performance, validation of requests, concurrency, and scalability.

The API Layer is mainly used to map the external user facing HTTP request towards the internal operations that are managed by the scheduler, worker system, and database systems.

## Roles of the API Layer

The responsibility of API layer is to:

### 1. Request Handling

- The ability to receive structured JSON requests by the web/mobile clients.
- Supervising inbound payload by Pydantic models.
- Securing type safety and schema verification.

### 2. Routing & Endpoint Management

Each cloud action maps to a specific API endpoint:

- `/vm/create`
- `/vm/delete/{vm_id}`
- `/vm/migrate/{vm_id}`
- `/vm/status/{vm_id}`
- `/metrics/{vm_id}`

### 3. Interaction with Scheduler

The API Layer communicates with the Scheduler to:

- Fetch fresh host metrics.
- Select the optimal hypervisor.
- Trigger decisions for VM placement.

### 4. Asynchronous Task Dispatching

The API Layer does not execute heavy operations. Instead, it creates Celery tasks for:

- VM provisioning
- VM deletion
- Migration
- Metrics collection

This is done via a task dispatcher that pushes requests to Redis.

### 5. Database Integration

The API interacts with PostgreSQL to:

- Insert new VM records

- Update VM states
- Query VM status
- Retrieve metrics data
- Log API activity

## 6. Response Management

The API returns:

- Task IDs for long-running operations
- Structured JSON responses
- Error handling messages
- Execution results

All responses follow consistent status codes and schema definitions.

## Core Components in the API Layer

**1. FastAPI Controller** The controller exposes endpoints and orchestrates all incoming API calls.

**2. RequestSchema (Pydantic Models)** Ensures:

- Automatic validation
- Type enforcement
- Error reporting
- Clean documentation via OpenAPI

**3. API Router** Organizes endpoints by logical grouping:

- VM operations
- Metrics endpoints
- Migration endpoints

**4. Task Dispatcher** Handles the transformation of API calls into asynchronous Celery tasks.

## Logical Flow of an API Request

1. **Client → API Gateway** User sends a request (e.g., create VM).
2. **Pydantic Validation** Payload schema is validated. If invalid → error returned immediately.
3. **Scheduler Invocation** API requests:
  - Host list
  - Metrics
  - Best-host selection
4. **Celery Task Creation** A task is queued in Redis containing:
  - VM configuration
  - Selected host
  - Template UUID
  - Metadata
5. **Database Logging** VMModel or TaskLogModel entries updated.
6. **Response** API returns:
  - Success message
  - Task ID
  - VM UUID (when ready)

## Advantages of This Logical Design

1. **Scalability** The API is stateless and horizontally scalable.
2. **Extensibility** New endpoints can be added without affecting worker logic.
3. **Fault Tolerance** Worker failures do not affect API responsiveness.
4. **Clear Separation of Concerns** API only handles:
  - Validation
  - Routing
  - Dispatch

Execution occurs in Celery workers.

## 5. Security Pydantic + FastAPI ensure:

- Input sanitization
- Automatic error filtering
- Reduced attack vectors

## API Layer Logical Flow Diagrams

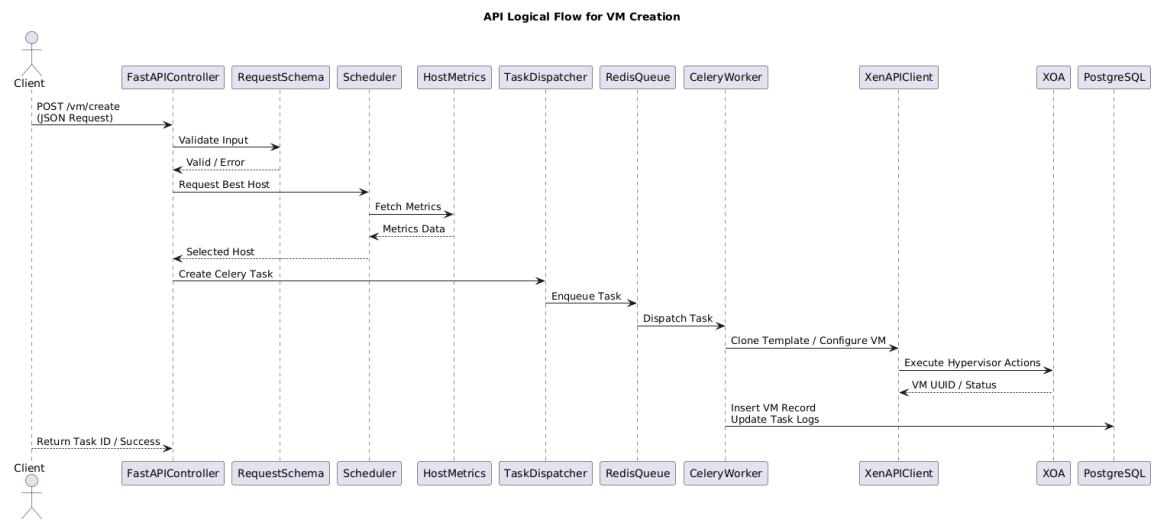


Figure 4.5: API Layer Logical Flow: Sequence Diagram Representing Request-to-Execution Flow

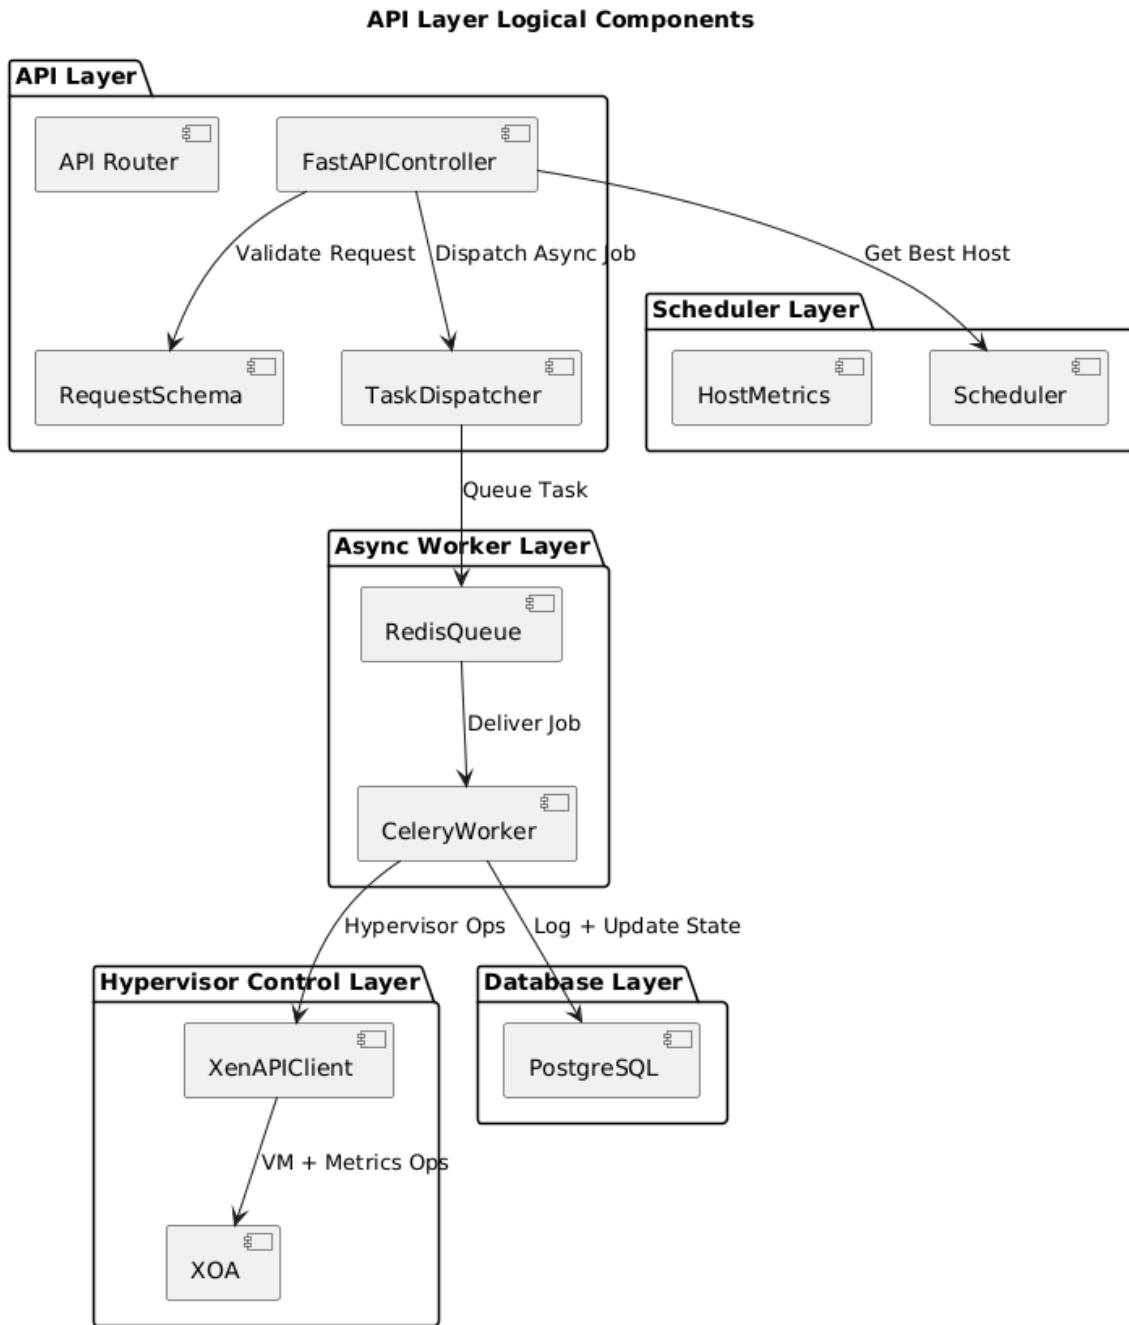


Figure 4.6: API Layer Logical Component Diagram Showing Internal System Interactions

#### 4.2.4 Microservices Interaction Design

The Mini-Cloud Platform follows a microservices-inspired architecture in which every major subsystem operates independently and communicates through well-defined interfaces. This design ensures scalability, fault tolerance, modularity, and clear separation of concerns across the platform. The interaction model between these microservices governs the entire lifecycle of virtual machines, including provisioning, migration, task execution, monitoring, and performance metrics collection.

Each microservice is responsible for a specific set of operations while interacting with others using lightweight communication mechanisms. This allows the platform to scale individual components without affecting the system as a whole and ensures that failures in one module do not propagate across the system.

## Core Microservices and Their Roles

- 1. API Service (FastAPI Controller)** Acts as the entry point to the system. It validates incoming requests, interacts with the scheduler, creates asynchronous tasks, updates the database, and returns responses to users. It does not perform any long-running operations but delegates such tasks to the Celery Workers.
- 2. Scheduler Service** Implements the Weighted Least Algorithm to determine optimal host selection for VM placement. It accesses host performance metrics from the database and returns the best hypervisor node based on weighted scoring.
- 3. Task Dispatch Service (Redis Queue)** Redis is used as a lightweight message broker through which asynchronous tasks are queued. It enables decoupled communication between the API service and Celery Workers, allowing the system to scale horizontally.
- 4. Celery Worker Service** Executes all heavy and long-running operations, including VM provisioning, migration, VM deletion, and metrics retrieval. Workers communicate with the XenAPI/XOA service to perform actual hypervisor-level actions.
- 5. XenAPI/XOA Service** This external service performs virtualization-layer actions such as template cloning, VM instantiation, migration, network attachment, and metrics retrieval. It exposes an API that the Celery Worker invokes to interact with the XCP-ng hypervisor pool.
- 6. Database Service (PostgreSQL)** Stores VM metadata, template information, task logs, and performance metrics. All microservices interact with the database either for read or write operations.

## Microservice Interaction Workflow

- 1. Client Request** A user initiates an operation (e.g., VM creation) through the API. The request is validated and forwarded to the scheduler.
- 2. Host Selection** The scheduler computes weighted scores for each host based on CPU usage, RAM usage, network utilization, and VM count. The optimal host is returned to the API.
- 3. Task Dispatch** The API service creates a Celery task containing the VM configuration and scheduling decision. This task is queued into Redis.
- 4. Asynchronous Execution** A Celery Worker pulls the job from Redis and communicates with the XenAPI/XOA service to clone templates, configure resources, start the VM, or

perform migrations.

**5. Database Update** After task execution, the worker service logs results in PostgreSQL, storing new VM entries, metrics, or task history records.

**6. Client Response** The API returns a task identifier and success response to the user, with the option to check the VM state later.

### Advantages of the Microservices Interaction Model

- **Loose Coupling:** Each service runs independently, allowing modular updates and isolated failures.
- **Scalability:** Celery Workers and API nodes can scale horizontally to meet demand.
- **Reliability:** Redis ensures durable message passing, enabling reliable asynchronous execution.
- **Fault Isolation:** Failures in one service do not disrupt the entire platform.
- **Extensibility:** New microservices (e.g., monitoring dashboards or autoscaling engines) can be added easily.

### Microservices Interaction Diagram Placeholder

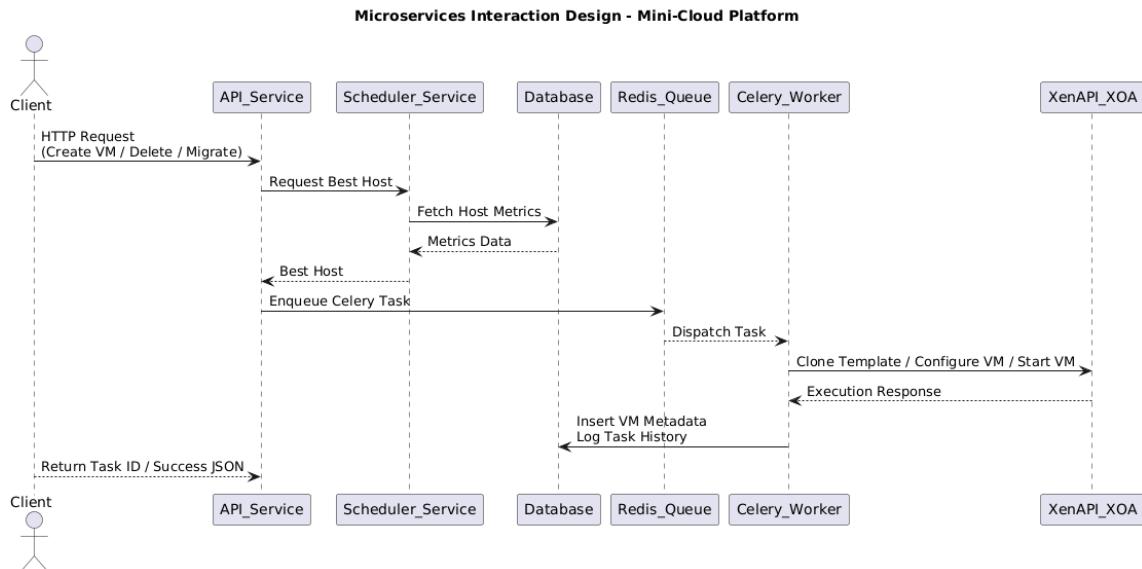


Figure 4.7: Microservices Interaction Diagram for the Mini-Cloud Platform

## 4.3 Behavioral Design

Behavioral design focuses on how different components of the Mini-Cloud Platform behave over time, how they interact with each other dynamically, and how the system responds to various events such as VM provisioning, migration, metrics collection, and scheduled background tasks. Unlike structural design, which describes the system's static architecture, behavioral design illustrates the runtime flow of control through the system using sequence diagrams and activity workflows.

This section presents detailed sequence diagrams that describe the interaction between the API service, scheduler, Celery workers, XenAPI/XOA, Redis, and the database during key operations of the Mini-Cloud Platform.

### 4.3.1 Sequence Diagrams

Sequence diagrams represent the chronological flow of operations in the system. They describe how objects interact with each other, the messages exchanged, the order of events, and the overall runtime behavior of the components involved.

This subsection contains sequence diagrams for:

- VM Creation Sequence
- VM Migration Sequence
- Metrics Collection Sequence
- Task Scheduling Sequence

#### VM Creation Sequence

The VM creation sequence describes the interactions involved in provisioning a new virtual machine. The workflow includes request validation, host selection, task dispatching, VM instantiation, and database updates.

- The client sends a VM creation request to the FastAPI Service.
- The API validates input using the RequestSchema.
- The API communicates with the Scheduler to determine the optimal hypervisor.
- A Celery task is created and queued into the Redis broker.
- A Celery Worker pulls the job and interacts with XenAPI/XOA to clone a VM template, configure its resources, and power on the VM.

- The worker updates the PostgreSQL database with VM metadata.
- A task ID is returned to the user for tracking progress.

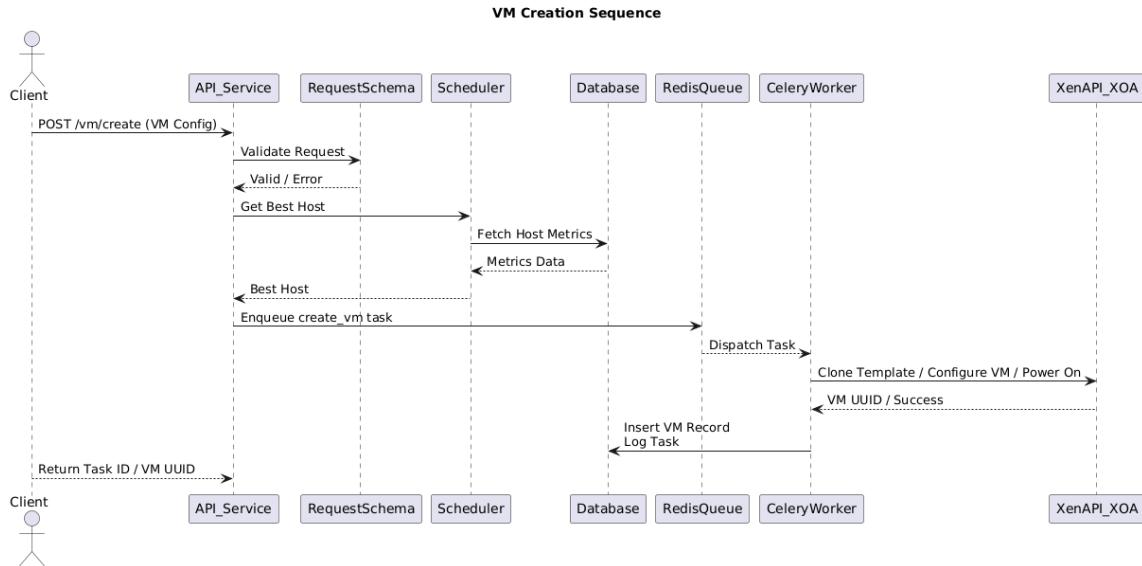


Figure 4.8: VM Creation Sequence Diagram

## VM Migration Sequence

The VM migration sequence describes how the system handles live migration of a virtual machine from one host to another. This is essential for load balancing, maintenance, and optimized resource utilization.

- The user requests a VM migration.
- The API invokes the Scheduler to validate host suitability.
- A migration task is created and sent to Redis.
- A Celery Worker consumes the task and calls XenAPI/XOA to execute the live migration.
- Once completed, the worker updates the VM's host entry in PostgreSQL.
- The user receives a confirmation response.

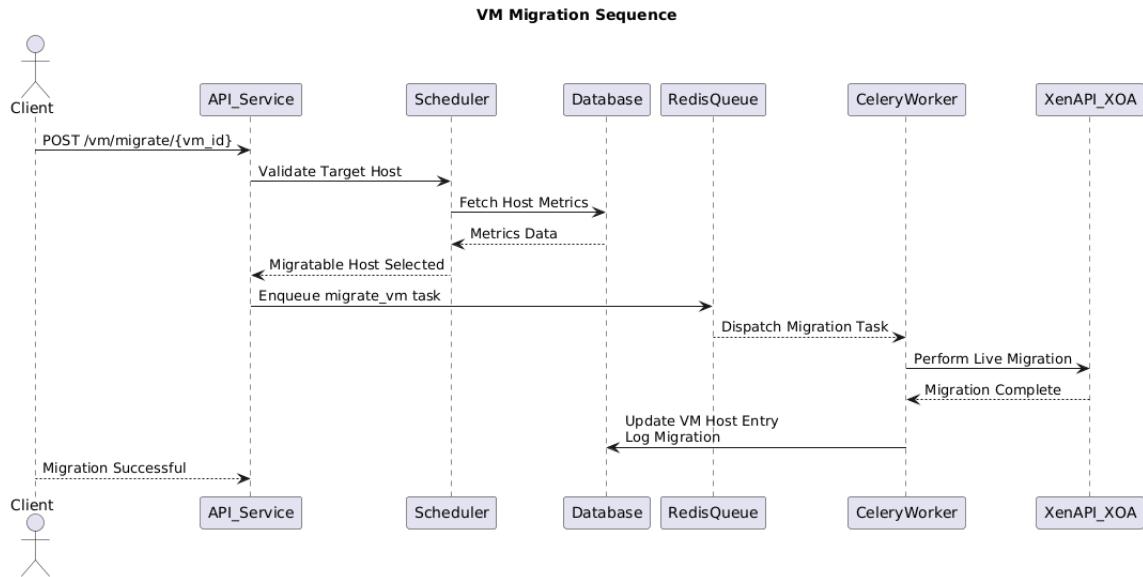


Figure 4.9: VM Migration Sequence Diagram

### Metrics Collection Sequence

The metrics collection sequence describes how the system retrieves and stores VM performance metrics such as CPU usage, memory consumption, disk I/O, and network activity.

- The scheduler or Celery beat triggers a metrics collection task.
- The API or worker triggers a metrics fetch job.
- The task is queued into Redis.
- A Celery Worker retrieves metrics via XenAPI/XOA.
- The metrics are stored in PostgreSQL.

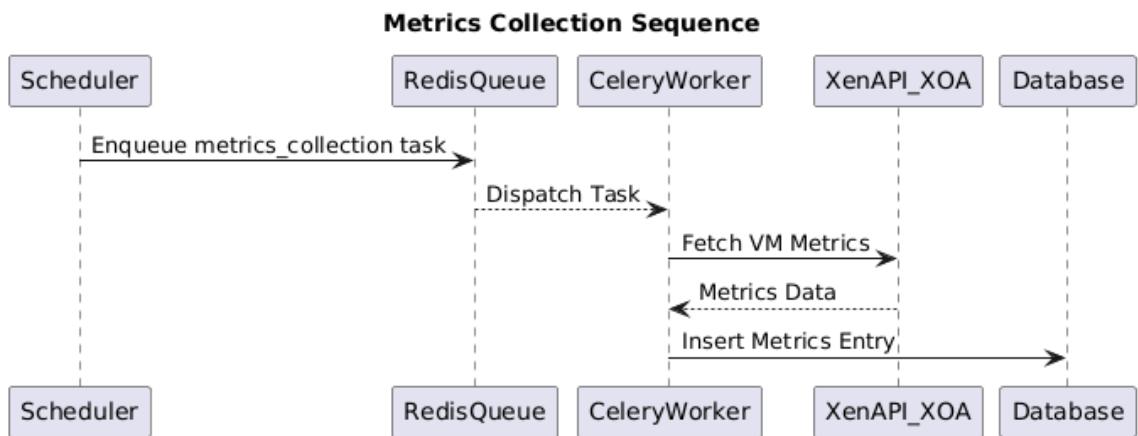


Figure 4.10: Metrics Collection Sequence Diagram

## Task Scheduling Sequence

This sequence describes how periodic background tasks (e.g., metrics collection, cleanup activities) are scheduled and executed.

- Celery Beat triggers periodic tasks based on schedule configuration.
- Beat pushes the scheduled tasks into the Redis Queue.
- Celery Workers pick tasks and execute them asynchronously.
- Relevant updates are written to PostgreSQL.

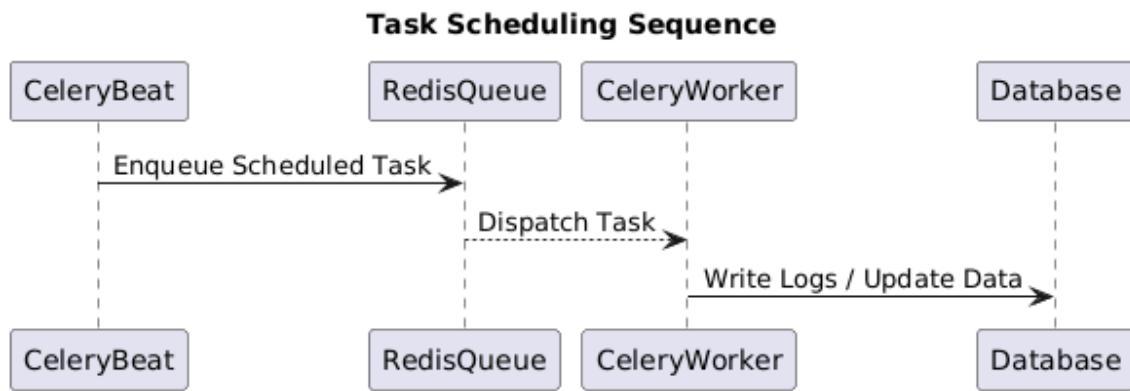


Figure 4.11: Task Scheduling Sequence Diagram

### 4.3.2 Activity Diagrams

Behavioral activity diagrams describe the major operational workflows inside the Mini-Cloud Platform. They capture the control flow, decision points, parallel activities, and termination conditions for common operations such as VM requests, monitoring/metrics collection, and asynchronous task execution. Activity diagrams complement sequence diagrams by showing branching, concurrency and high-level procedural logic.

#### VM Request Flow

The VM Request Flow activity diagram models the user-initiated VM creation process from request reception through validation, scheduling, task dispatching, provisioning, and final logging. Key activities include:

- Receive request (Client → FastAPI)
- Validate payload (Pydantic schema)

- Authenticate/Authorize (if enabled)
- Query Scheduler for host selection
- Enqueue task in Redis (Task Dispatcher)
- Execute provisioning in Celery Worker (clone template, configure, power on)
- Update database with VM metadata and task log
- Return response / notify client

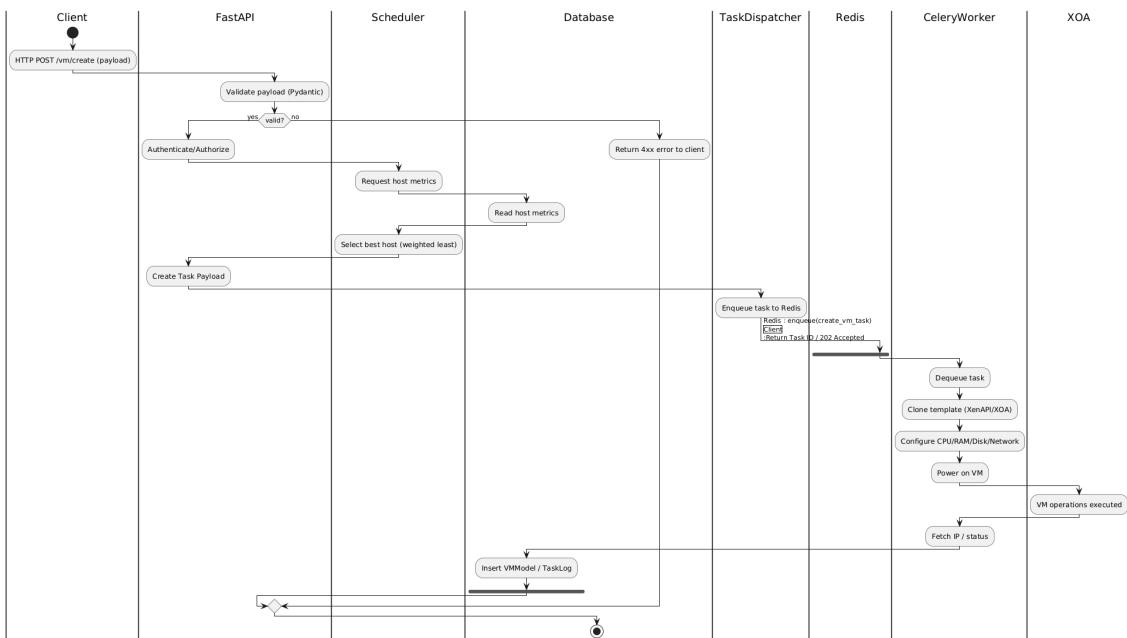


Figure 4.12: Activity Diagram: VM Request Flow

## Monitoring / Metrics Flow

The Monitoring/Metrics Flow activity diagram shows periodic metrics collection and storage workflow. It includes scheduled trigger (Celery beat), parallel metric fetch for multiple VMs, aggregation and storage, and optional alerting when thresholds are exceeded. Important steps:

- Trigger scheduler (Celery beat or external scheduler)
- Build list of active VMs
- Parallel dispatch of metrics fetch tasks to Celery workers
- Collect metrics via XOA/XenAPI
- Store metrics in PostgreSQL

- Check thresholds / generate alerts (optional)

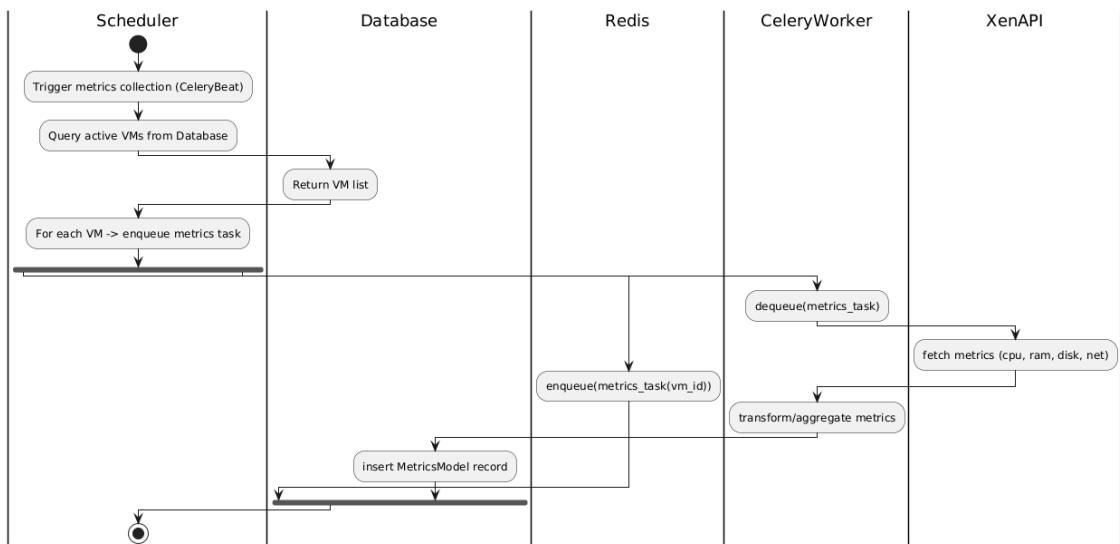


Figure 4.13: Activity Diagram: Monitoring / Metrics Flow

### Asynchronous Task Execution Flow

This activity diagram describes the lifecycle of any asynchronous task (create, migrate, delete, collect metrics) from enqueue to completion, including retries and error handling.

Major steps:

- API enqueues a task in Redis
- Celery worker receives the task
- Execute action (with try/except)
- On success: update TaskLogModel, update VMMModel/metrics as needed
- On failure: log error, apply retry policy, escalate if needed
- Notify caller (webhook/polling) with final status

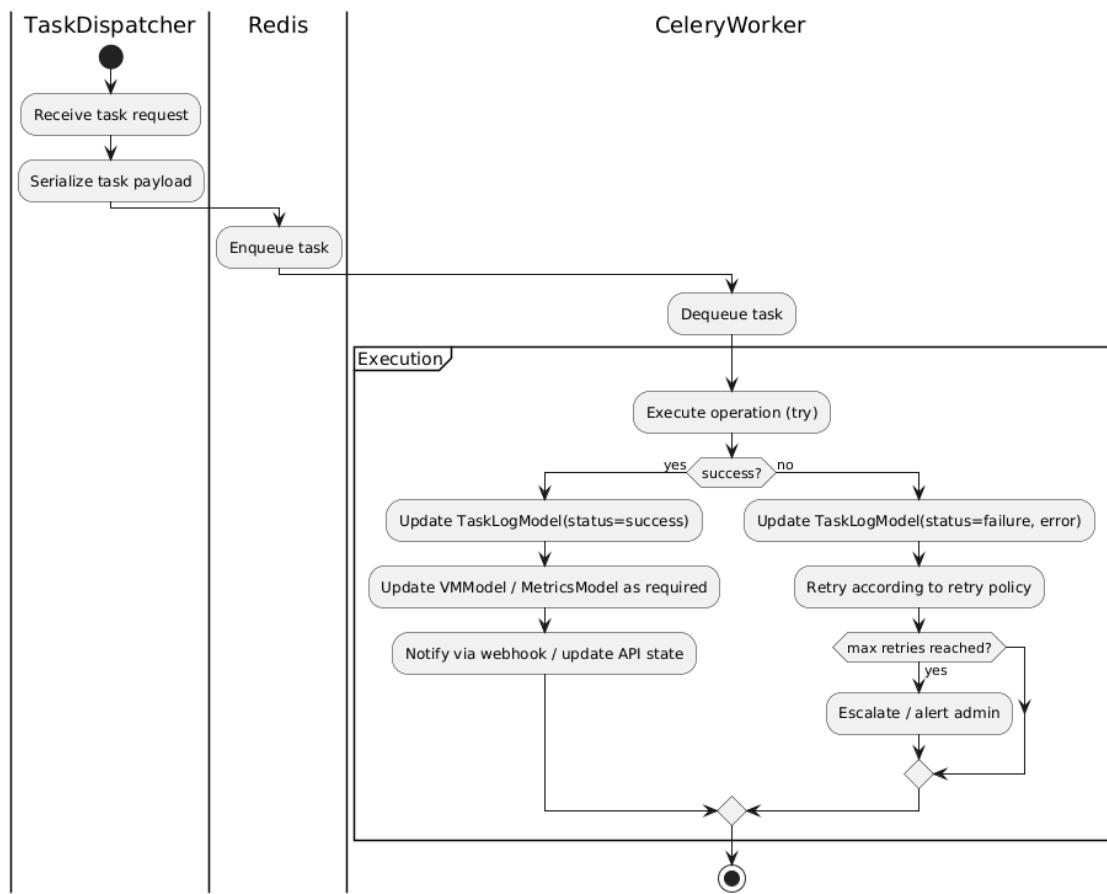


Figure 4.14: Activity Diagram: Asynchronous Task Execution Flow

### 4.3.3 State Diagrams

State diagrams (state machines) describe the possible life-cycle states of primary domain objects. For the Mini-Cloud Platform the most important object whose state needs modeling is the Virtual Machine (VM). The VM Lifecycle State Diagram captures valid states, allowed transitions, triggering events, and any transient states used during operations (e.g., migrating).

#### VM Lifecycle State Diagram

The VM lifecycle state diagram defines the complete lifecycle of a VM managed by the platform:

- **Template** – the VM is represented by a template/blueprint.
- **Provisioning** – worker is cloning/configuring VM resources.
- **Running** – VM is powered on and available.
- **Paused / Suspended** – VM is temporarily suspended.

- **Migrating** – live migration in progress (transient state).
- **Stopped** – VM is powered off.
- **Deleted** – VM has been removed; terminal state.

Transitions are triggered by API calls (create, stop, start, migrate, delete), worker completion events, or host-level failures. Retry and rollback transitions (e.g., provisioning failed → delete / cleanup) are also included.

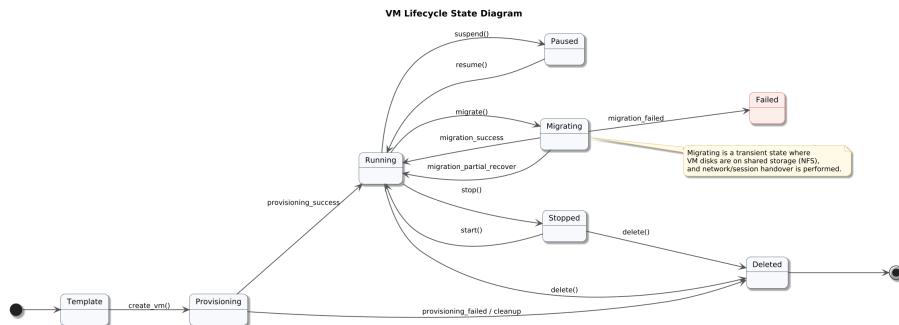


Figure 4.15: State Diagram: VM Lifecycle

## 4.4 Process Design & Flowcharts

The Process Design section illustrates the procedural workflows of major components in the Mini-Cloud Platform using flowcharts. These diagrams describe logical steps, decision points, and operational phases involved in VM creation, metrics collection, migration, and API-level request handling. Flowcharts provide a simplified visual representation of internal logic and help readers understand the high-level operational sequence without requiring knowledge of underlying implementation details.

### 4.4.1 VM Creation Flowchart

The VM Creation Flowchart represents the complete process executed when a user requests a new virtual machine. It includes validation, scheduling, task dispatching, provisioning, and final system updates.

Key steps include:

- Accept request from client
- Validate request payload
- Query scheduler for host selection

- Enqueue asynchronous VM creation task
- Celery worker provisions the VM
- Database is updated with VM metadata
- Response is returned to the user

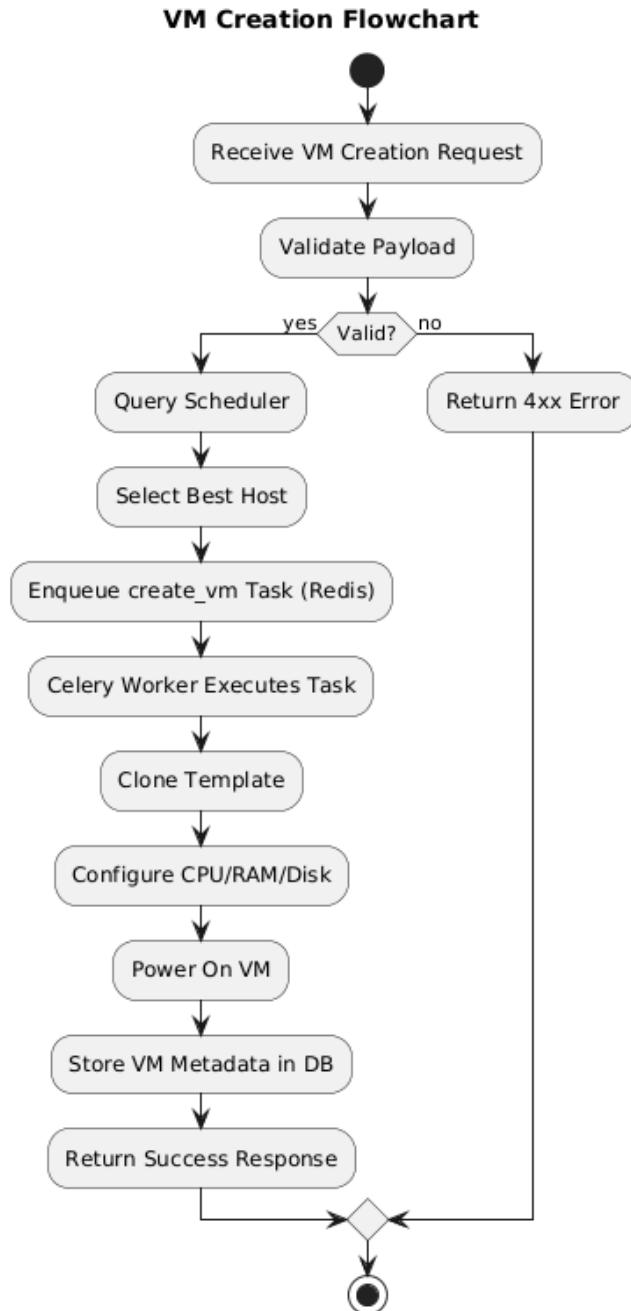


Figure 4.16: VM Creation Flowchart

#### 4.4.2 Metrics Collection Flowchart

This flowchart describes the periodic collection of VM and host metrics used for monitoring and scheduling decisions. Metrics are gathered asynchronously and stored in the database. Key steps:

- Trigger collection via scheduler or Celery Beat
- Enqueue metrics collection tasks
- Fetch metrics from hypervisor via XOA
- Validate and format the collected metrics
- Store metrics into PostgreSQL

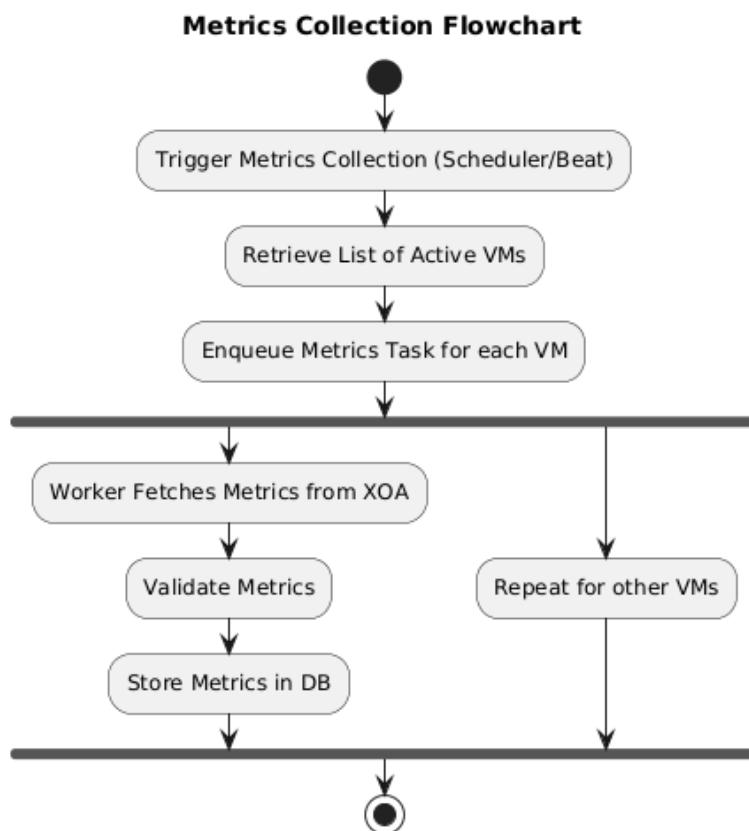


Figure 4.17: Metrics Collection Flowchart

#### 4.4.3 Migration Service Flowchart

The migration flowchart models the sequence followed when migrating a VM from one host to another. This process involves scheduler validation, asynchronous task execution, hypervisor interaction, and logging.

Steps include:

- Receive migration request
- Validate VM existence
- Request scheduler to verify target host
- Enqueue migration task into Redis
- Worker performs migration via XOA
- Update VM host entry in PostgreSQL
- Return migration status to the client

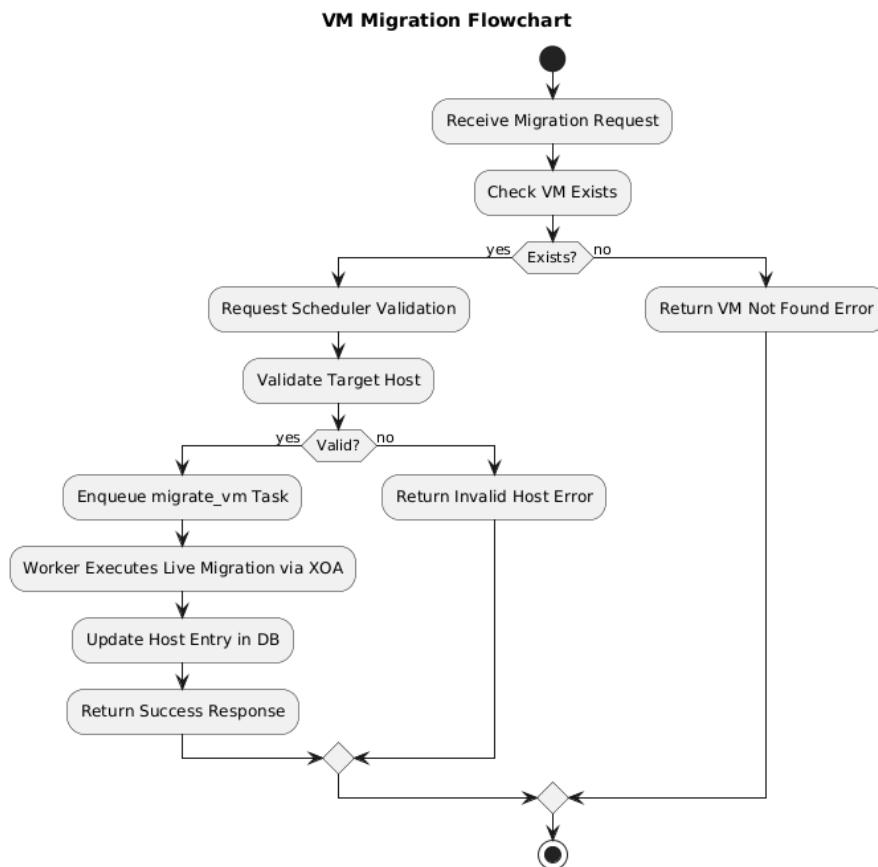


Figure 4.18: VM Migration Service Flowchart

#### 4.4.4 API Request Processing Flow

This flowchart illustrates the processing pipeline for every incoming API request. It applies to VM creation, deletion, migration, and metrics retrieval endpoints.

The flow includes:

- Request reception
- Schema validation
- Routing to appropriate controller
- Scheduler interaction (if required)
- Enqueue tasks (if asynchronous)
- Generate response

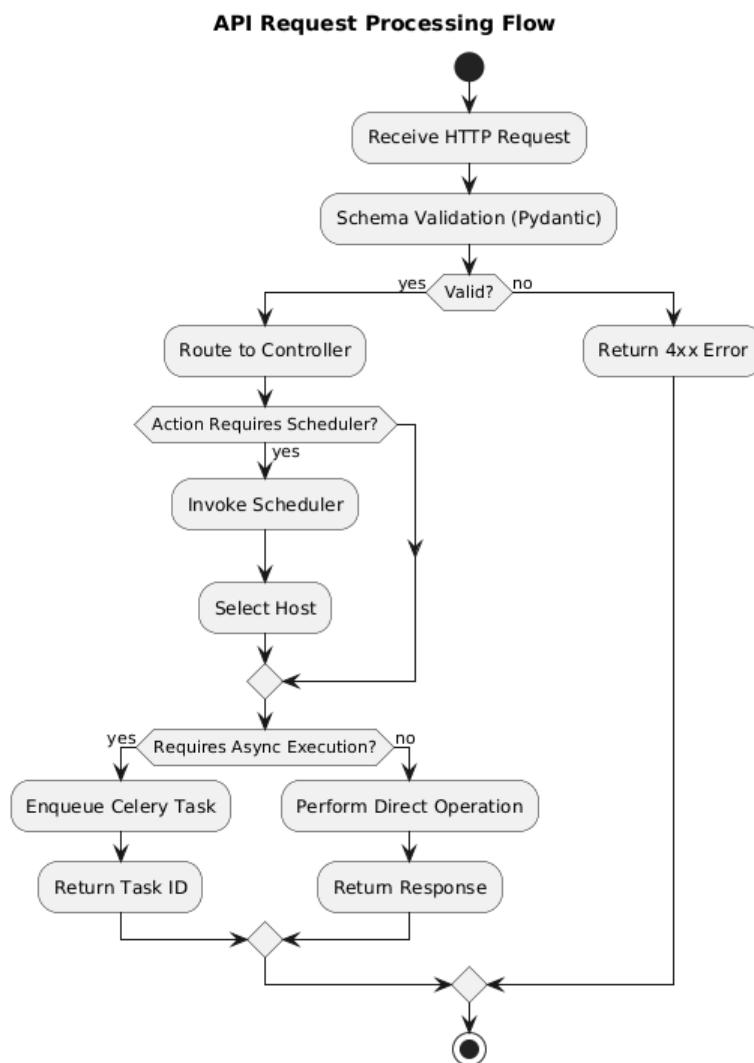


Figure 4.19: API Request Processing Flowchart

## 4.5 Algorithms and Pseudocode (Detailed)

This section presents detailed, production-quality pseudocode for the core algorithms used in the Mini-Cloud Platform. Each algorithm includes: purpose, inputs, outputs,

preconditions, postconditions, main steps, error handling, and retry logic where applicable. The full, runnable source code for the implementation is included in Appendix A.

#### 4.5.1 VM Creation Algorithm

**Purpose:** Provision a new virtual machine on the least-loaded hypervisor with cloud-init, ensure the VM boots, obtain its IP address, and persist metadata.

**Inputs:**

- `vm_name` – desired VM name
- `cpu, ram, disk` – resource requests
- `template` – template name or id
- `ssh_key` – public key path for cloud-init
- `preferred_host` (optional)

**Outputs:**

- `result` containing `vm_uuid, vm_ip, status, logs`

**Preconditions:**

- At least one host registered and reachable.
- Template exists on host or pool.
- SSH access to host with a key that can run required commands.

**Postconditions:**

- VM record is created in database if provisioning succeeds (best-effort).
- VM is started and IP is attempted to be detected.

---

**Algorithm 1** VM Creation (detailed)

---

**Require:** vm\_name, cpu, ram, disk, template, ssh\_key, (preferred\_host)

**Ensure:** result = {vm\_uuid, vm\_ip, status, logs}

- 1: LOG("Start VM creation:", vm\_name)
- 2: host  $\leftarrow$  if preferred\_host given then preferred\_host else SelectLeastLoadedHost()
- 3: **if** host == NULL **then**
- 4:   LOG("ERROR: no available host"); RETURN error result
- 5: **end if**
- 6: workdir  $\leftarrow$  CreateTemporaryWorkdir(vm\_name)
- 7: seed\_iso  $\leftarrow$  BuildCloudInitISO(workdir, ssh\_key, vm\_name)
- 8: UploadISOToHost(seed\_iso, host)
- 9: iso\_vdi  $\leftarrow$  EnsureISOImportedToSR(host, seed\_iso)
- 10: **if** iso\_vdi is NULL **then**
- 11:   LOG("WARN: ISO VDI not available; continuing with best-effort")
- 12: **end if**
- 13: template\_uuid  $\leftarrow$  FindTemplateOnHost(host, template)
- 14: **if** template\_uuid is NULL **then**
- 15:   LOG("ERROR: template not found on host"); RETURN error result
- 16: **end if**
- 17: vm\_uuid  $\leftarrow$  CloneVMFromTemplate(host, template\_uuid, vm\_name)
- 18: **if** vm\_uuid is NULL **then**
- 19:   LOG("ERROR: clone failed"); CLEANUP(workdir); RETURN error result
- 20: **end if**
- 21: ApplyResources(vm\_uuid, cpu, ram, disk)
- 22: **if** iso\_vdi  $\neq$  NULL **then**
- 23:   AttachCD(vm\_uuid, iso\_vdi)
- 24: **end if**
- 25: StartVM(vm\_uuid)
- 26: vm\_ip  $\leftarrow$  DetectVMIPGuaranteed(host, vm\_uuid)
- 27: **if** vm\_ip == UNKNOWN **then**
- 28:   LOG("WARN: IP detection failed; marking IP as UNKNOWN")
- 29: **end if**
- 30: PersistVMRecordDB(vm\_uuid, vm\_name, host, cpu, ram, disk, vm\_ip)
- 31: DetachAndDestroyCloudInitCD(vm\_uuid)
- 32: CLEANUP(workdir)
- 33: LOG("VM creation finished:", vm\_uuid, vm\_ip)
- 34: **return** success result =0

---

**IP Detection – summary of strategy (used in DetectVMIPGuaranteed):**

1. Try xenstore keys such as `vm-data/ip` and `device/vif/0/ip`.
2. If not present, read domain id and try `/local/domain/<domid>/data/ip`.
3. Parse `vm-data/networks` for IPv4 tokens.
4. Use VIF MAC to consult ARP table, bridge FDB, neighbor table, and DHCP lease files on host.

5. Repeat with backoff up to a timeout (e.g., 2 minutes) and report UNKNOWN if not found.

### 4.5.2 Template Cloning Algorithm

**Purpose:** Clone an existing VM template into a new VM instance.

**Inputs:** template\_name, new\_vm\_name, host.

**Outputs:** vm\_uuid or error.

---

#### Algorithm 2 Template Cloning

---

**Require:** template\_name, new\_vm\_name, host

```

1: LOG("Find template on host")
2: template_uuid ← FindTemplateOnHost(host, template_name)
3: if template_uuid is NULL then
4:   LOG("ERROR: template missing"); RETURN NULL
5: end if
6: vm_uuid ← RunCommandOnHost(host, "xe vm-clone uuid="+template_uuid+
  new-name-label="+new_vm_name)
7: if vm_uuid is NULL then
8:   LOG("ERROR: clone failed"); RETURN NULL
9: end if
10: return vm_uuid =0

```

---

### 4.5.3 Metrics Fetching Algorithm

**Purpose:** Periodically collect VM/host performance metrics and persist them.

**Inputs:** list of active VMs or hosts.

**Outputs:** persisted metrics records.

---

#### Algorithm 3 Metrics Fetching (parallel-safe)

---

**Require:** time\_slice or triggered event

```

1: vm_list ← QueryActiveVMsFromDB()
2: for all vm IN vm_list in parallel do
3:   metrics ← XenAPIClient.fetch_metrics(vm)
4:   if metrics invalid then
5:     LOG("WARN: metrics invalid for", vm)
6:     continue
7:   end if
8:   NormalizeAndValidate(metrics)
9:   InsertMetricsIntoDB(vm, metrics)
10: end for
11: LOG("Metrics collection completed for", |vm_list|) =0

```

---

**Notes:**

- Use worker concurrency (Celery pool) to parallelize metric retrieval across many VMs.
- Apply basic sanity checks on returned values and cap abnormal spikes before storing.

#### 4.5.4 Celery Task Scheduling Algorithm

**Purpose:** Describe how tasks are scheduled, retried and executed via Celery and Redis.

**Key features:** enqueue, ack, retries, exponential backoff, result logging.

---

##### Algorithm 4 Celery Task Scheduling (high-level)

---

**Require:** task\_definition (type, payload, retry\_policy)

- 1: TaskDispatcher → Serialize(task\_definition)
  - 2: PushToRedisQueue(serialized)
  - 3: Worker ← WaitForTask()
  - 4: OnTaskReceived:
  - 5: ExecuteTaskHandler(payload)
  - 6: UpdateTaskLog(status=SUCCESS) TransientError
  - 7: If retries left: ScheduleRetryWithBackoff()
  - 8: UpdateTaskLog(status=RETRY) FatalError
  - 9: UpdateTaskLog(status=FAILED, error=msg)
  - 10: RaiseAlertIfCritical() =0
- 

#### 4.5.5 VM Migration Algorithm

**Purpose:** Identify overloaded hosts and migrate appropriate VMs to balance load.

**Inputs:** host metrics history, VM list per host.

**Outputs:** migration actions invoked (or no-op).

---

**Algorithm 5** VM Migration (detailed)

---

```

1: hosts ← GetAllHosts()
2: scored = []
3: for all host IN hosts do
4:   latest ← GetLatestMetric(host)
5:   if latest is NULL OR HostOverloaded(latest)=False then
6:     continue (or mark as candidate)
7:   end if
8:   score_info ← CalculateHostScore(latest) {uses cpu,mem,vmcount weights}
9:   Append(scored, (score_info.score, host))
10: end for
11: if len(scored) < 2 then
12:   LOG("Insufficient data to migrate"); RETURN no-action
13: end if
14: SortDescending(scored)
15: src = scored[0].host
16: dest = scored[-1].host
17: if scored[0].score - scored[-1].score < MIGRATION_THRESHOLD then
18:   LOG("Load difference below threshold"); RETURN no-action
19: end if
20: vm = SelectCandidateVMForMigration(src) {e.g., smallest VM by memory}
21: if vm is NULL then
22:   LOG("No candidate VM found"); RETURN no-action
23: end if
24: InvokeMigration(vm, dest) {xe vm-migrate or XOA API}
25: if migration succeeded then
26:   UpdateVMHostInDB(vm, dest)
27:   LogMigrationEvent(vm, src, dest)
28:   RETURN success
29: else
30:   LogMigrationFailure(vm, error)
31:   OptionalRollbackOrRetry()
32:   RETURN failure
33: end if=0

```

---

**Notes:**

- The migration algorithm should include safety checks (storage compatibility, network affinity).
- Prefer live migration with shared storage (NFS SR) to avoid VDI transfer.
- Conservative threshold (e.g., 0.15 score difference) reduces unnecessary migrations.

# 5. Hardware and Software Specifications

This section presents a comprehensive overview of the hardware and software ecosystem used to design, deploy, and evaluate the Mini-Cloud Virtualization Controller Platform. A cloud orchestration system depends on both its underlying compute infrastructure and its software environment; therefore, accurate specification of these components is essential for reproducibility, performance analysis, and system validation.

The Mini-Cloud Platform is implemented using two XCP-NG hypervisors, centrally managed using Xen-Orchestra (XOA), and controlled by a containerized FastAPI–Celery–Redis–PostgreSQL stack. This section describes each hardware and software component in detail.

## 5.1 Hardware Setup

### 5.1.1 XCP-NG Servers

The virtualization layer is built using two physical servers running XCP-NG 8.3 as the bare-metal hypervisor. These servers are part of the same resource pool and support features such as VM provisioning, live migration, metrics extraction, and remote execution via XenAPI.

Table 5.1: Hypervisor Hardware Overview

Component	Server 1	Server 2
Model	HP EliteDesk 800 G1 SFF	HP EliteDesk 800 G1 SFF
CPU	Intel Core i7	Intel Core i7
RAM	16 GB DDR3	16 GB DDR3
Storage	500 GB SSD	500 GB SSD
NIC	Gigabit Ethernet	Gigabit Ethernet
Hypervisor OS	XCP-NG 8.3	XCP-NG 8.3
Management IP	10.20.24.40	10.20.24.38



Figure 5.1: XCP-NG hypervisor nodes detected in Xen-Orchestra.

### 5.1.2 Host Specifications

Each host supports full hardware virtualization (Intel VT-x), Xen PVHVM guests, NFS-backed storage repositories, and hybrid resource scheduling. The hosts are configured with:

- Static IP addressing for reliability
- BIOS virtualization extensions enabled
- Shared NFS SR mounted for live migration
- XenAPI and SSH enabled for controller-driven provisioning

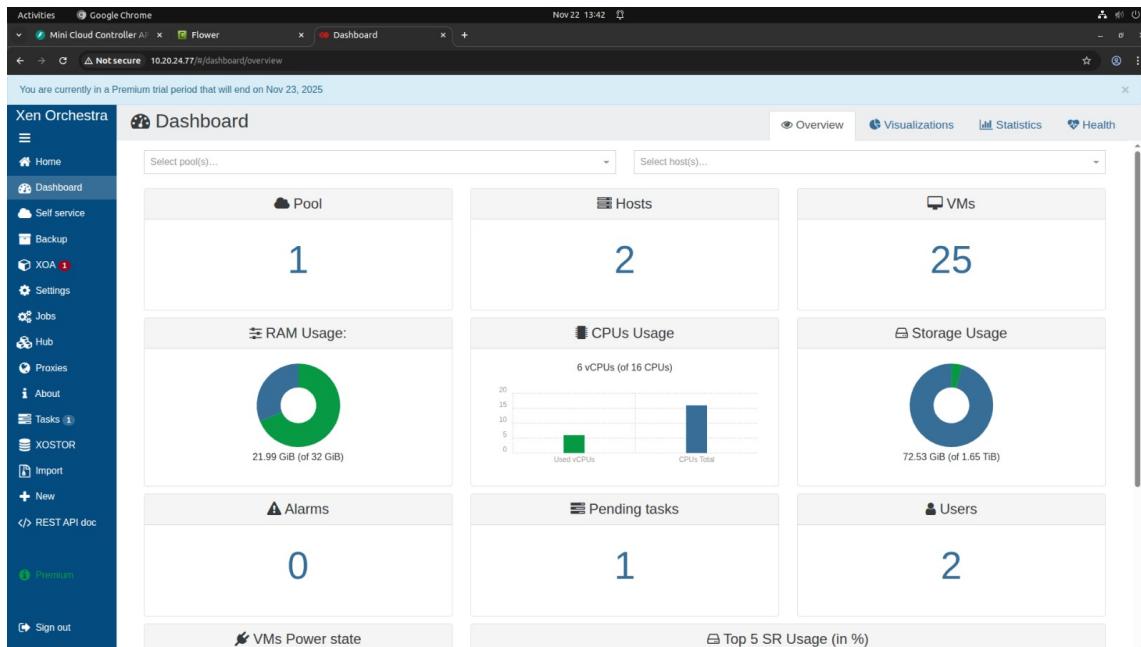


Figure 5.2: Xen-Orchestra dashboard showing RAM, CPU, storage, and pool-level statistics.

### 5.1.3 Network Topology

All components operate within a single Layer-2 switched network. The complete network configuration is:

- Subnet: 10.20.24.0/24
- Gateway: 10.20.24.1
- Controller Node: Docker-based FastAPI server

- Hypervisors: 10.20.24.40 and 10.20.24.38
- Xen-Orchestra Virtual Appliance (XOA): 10.20.24.77

Four primary types of traffic are used:

1. API & provisioning traffic (controller → hypervisors)
2. Metrics reporting traffic (hypervisors → controller)
3. Storage traffic (hypervisors → NFS SR)
4. Migration traffic (hypervisor ↔ hypervisor)

#### 5.1.4 Storage Environment

A shared NFS-based Storage Repository (SR) is used for all VM disk images and ISO resources. This configuration enables:

- VM live migration
- Centralized template management
- Cloud-init ISO storage
- Pool-wide VDI accessibility



Figure 5.3: Shared NFS Storage Repository used by both XCP-NG hypervisors.

## 5.2 Software Stack

The controller platform is implemented using modern microservices and asynchronous task handling. This section describes the full software stack used.

### 5.2.1 Operating Systems

#### Host OS:

- XCP-NG 8.3 running directly on the physical hardware.

#### Controller OS:

- Ubuntu 22.04 LTS (Docker-based deployment)

#### Guest OS (VM Templates):

- Ubuntu 22.04 Cloud-init compatible images

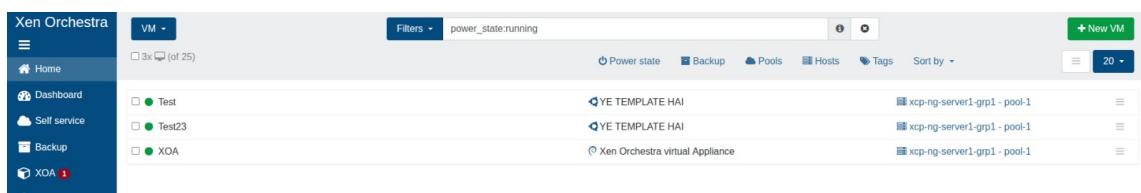


Figure 5.4: Xen-Orchestra VM list showing active VMs and templates.

### 5.2.2 Python Version

All backend modules are written in **Python 3.11**. Core dependencies include:

- fastapi
- pydantic
- sqlalchemy
- celery
- redis
- requests
- uvicorn

Python services are containerized for reproducibility and ease of deployment.

### 5.2.3 FastAPI, Celery, Redis, and PostgreSQL

**FastAPI** Serves as the main REST controller for:

- VM creation
- VM listing
- Host registration
- Metrics reporting
- Migration job dispatch

The screenshot shows the OpenAPI documentation for the Mini Cloud Controller API. The URL is 127.0.0.1:8001/docs. The interface is a clean, modern design with a dark header bar. Below the header, the title "Mini Cloud Controller API" is displayed along with version 0.1.0 and OAS 3.1. A link to "openapi.json" is also present. The main content area is organized into sections: "hosts", "metrics", "vms", "jobs", and "xoa". Each section contains a list of endpoints with their methods (e.g., POST, GET) and paths. The "hosts" section has two entries: "POST /hosts/register Register Host" and "GET /hosts Get Hosts". The "metrics" section has one entry: "POST /metrics/report Report Metrics". The "vms" section has two entries: "POST /vms/register Register Vm" and "GET /vms/ List Vms". The "jobs" section has one entry: "GET /jobs/{job\_id} Get Job Status". The "xoa" section has three entries: "GET /xoa/pools List Pools", "GET /xoa/templates/{pool\_uuid} List Templates", and "POST /xoa/create\_vm Create Vm". Each endpoint entry includes a copy icon.

Figure 5.5: FastAPI OpenAPI documentation showcasing all exposed endpoints.

**Celery** Celery workers execute asynchronous jobs such as:

- VM provisioning
- Host metrics collection
- Migration triggers

Figure 5.6: Celery Flower dashboard showing task queues and worker health.

**Redis** Used as the message broker for asynchronous task communication.

**PostgreSQL** Stores all relational data such as VM metadata, host metrics, migration events, and task logs.

Figure 5.7: pgAdmin view of metrics, events, and VM tables used by the controller.

### 5.2.4 Docker and Docker-Compose Setup

All backend services are deployed using Docker Compose. The major services include:

- `fastapi` service (REST controller)
- `celery_worker` service
- `celery_beat` scheduler
- `redis` message broker

- postgres database container

Containerization provides:

- Reproducibility
- Dependency isolation
- Fault isolation
- Simplified deployment

### 5.2.5 XOA / Xen-Orchestra Integration

Xen-Orchestra (XOA) acts as the management API for the hypervisor pool. It provides:

- Pool inspection (hosts, VMs, storage)
- Template listing
- REST API endpoints
- Live migration triggers



Figure 5.8: XOA Login interface used for cluster and resource orchestration.

### 5.2.6 Monitoring Tools

The Mini-Cloud Platform uses the following built-in monitoring capabilities:

- **Xen-Orchestra Dashboard:** Real-time tracking of CPU, RAM, storage, and VM activity.
- **Celery Flower Dashboard:** Monitoring asynchronous task execution.
- **pgAdmin:** Database query analysis and inspection.

### 5.2.7 Third-Party Libraries

Several external and system utilities support the controller's functionality:

- `genisoimage` (for cloud-init ISO generation)
- `openssh-client` (remote execution on hosts)
- `xe` CLI utilities (XenAPI interaction)
- `nfs-common` (shared storage mounting)

# 6. Cloud Services Categorization

Cloud computing services can be broadly divided into three main types of services: Infrastructure as a Service (IaaS), Platform as a Service (PaaS) and Software as a Service (SaaS). The Mini-Cloud Platform is a hybrid cloud orchestration based on the three-layers integration of components and as such, it is similar to the real-world cloud platforms like the AWS, Azure, or OpenStack.

This section involves a breakdown of the system components into the three cloud service models and how the component will be used to help the overall architecture.

## 6.1 Infrastructure as a Service (IaaS) Components

IaaS refers to provisioning raw compute, storage, and networking resources on demand. The following features in the Mini-Cloud Platform can be considered IaaS capabilities:

### 6.1.1 VM Provisioning

The system offers complete life cycle management of virtual machines as in AWS EC2, or Azure Virtual Machines. Through **FastAPI endpoints** and **Celery-based task execution**, the controller performs:

- Creation of VM on demand based on predefined templates.
- TM of VMs to the least-loaded hypervisor.
- CPU, memory, disk and network configuration.
- Exercise of provisioning scripts on XCP-NG hypervisors remotely.
- Live VM power operations (start, stop, delete).

VM provisioning pipeline resembles the normal IaaS operations in which raw compute instances are provided as services.

### 6.1.2 Template Management

Templates are pre-images of VM deployment. The Mini-Cloud Platform is compatible with:

- Listing templates that are available via XOA/XenAPI.

- Replicate templates after which new VM instances are created.
- Cloud-init based automation for SSH access and initialization.
- Managing template metadata (OS type, resource defaults).

These capabilities correspond directly to IaaS imaging and snapshot management services.

### 6.1.3 Network and Storage

It is a system that offers software-defined network and storage abstraction:

- XCP-NG virtual switch network bridging.
- NIC volume snapshotting upon creation of new VMs.
- Shared NFS-based storage repositories enabling:
  - Live migration
  - Pool-wide VDI sharing
  - Centralized ISO storage
- Scanning dynamic SRs in VM creation processes.

These features are in line with the essence of IaaS networking and storage services in commercial cloud systems.

## 6.2 Platform as a Service (PaaS) Components

PaaS decouples the infrastructure and offers an execution and development platform of applications. The Mini-Cloud Platform has a number of PaaS-like components to facilitate service orchestration, automation and hosting of application logic.

### 6.2.1 Celery Task Execution Platform

Celery acts as the distributed task execution layer of the platform, providing:

- Asynchronous and trustworthy job execution.
- Distributed processing across multiple workers.
- Built-in scheduling via Celery Beat.

- Decoupling of heavy operations (VM creation, metrics collection).
- Fault-tolerant job retries and status tracking.

This resembles managed PaaS offerings such as AWS SQS/Lambda integration or Azure Functions Task Queue systems.

### 6.2.2 API Management Layer

FastAPI provides the PaaS-level API gateway for:

- Validating requests using Pydantic schemas.
- Routing client requests to internal microservices.
- Integrating with XenAPI and Celery task pipelines.
- Managing API documentation via automatic OpenAPI generation.

This API layer functions similarly to AWS API Gateway or Google Cloud Endpoints.

### 6.2.3 Database Services

PostgreSQL is used as a backend service for data persistence. It provides structured storage for:

- VM metadata
- Host metrics
- Task logs
- Migration events
- Host registrations

Because the system exposes database services to higher-level components but not to end users, PostgreSQL operates as an internal PaaS subsystem.

## 6.3 Software as a Service (SaaS) Components

SaaS components represent fully managed, end-user accessible services requiring minimal configuration.

### 6.3.1 Xen-Orchestra Web Interface (SaaS Functionality)

Xen-Orchestra provides a browser-based, fully interactive management dashboard that includes:

- VM lifecycle control (start, stop, migrate, delete)
- Storage and pool monitoring
- Live charts for CPU, memory, and network usage
- Template browsing and import
- Hypervisor health monitoring

Although internally connected to our controller, XOA functions as a SaaS tool for administrators.

### 6.3.2 Custom API Services as SaaS Endpoints

The Mini-Cloud exposes a complete REST API service accessible to end users. These APIs provide:

- VM creation and deletion
- VMs, hosts and templates listing.
- Sending host metrics
- Triggering migration jobs
- Retrieving job statuses

Since these services are availed to end users (UI, CLI, or external programs), they serve as SaaS services.

REST API hence represents a SaaS layer over the PaaS backend logic and IaaS hypervisor resources.

# 7. Implementation

This part is a detailed summary of the actual application of the Mini-Cloud Platform. Its implementation includes the configuration of the virtualized infrastructure (XCP-NG and Xen-Orchestra), the deployment of the controller stack with the usage of Docker, and the creation of the REST API with FastAPI and the implementation of asynchronous tasks with the help of Celery. In each subsection, we shall detail the configuration options, reasons and details of how the system will be configured and deployed to run successfully.

## 7.1 Setting Up the Environment

The environment configuration consists of readying the hypervisor pool, configuring Xen-Orchestra (XOA), deploying the controller stack and making sure that all the services communicate. Two XCP-NG nodes, a centralized XOA appliance, and a controller host with the FastAPI service, Celery workers, Redis, and PostgreSQL are used in its implementation.

### 7.1.1 XCP-NG Pool Setup

The Mini-Cloud Platform is based on a cluster of XCP-NG hypervisor. The pool was configured in the following way:

- **XCP-NG 8.x** is installed on both physical servers.
- **Assigning static IPs:**
  - Node 1: 10.20.24.40
  - Node 2: 10.20.24.38
- **Setting up a pool:** The master host was created as the first host, and the second host was added to the master one to create a pool.
- **Shared Storage Configuration:** There was a NFS-based Storage Repository (SR) that was mounted on both the hosts. This enables:
  - Storage-level VM mobility
  - Live migration across nodes
  - Cloud-init image storage in ISO.

- **Template Import:** The Ubuntu cloud images were loaded into XOA and transformed into VM templates that can be used repeatedly.

Visual representations of the host pool, SR utilization, VM status pages and dashboard overview can be placed here to demonstrate the environment setup.

### 7.1.2 Connecting XOA with API Tokens

Xen-Orchestra (XOA) is the interface through which one can interact with the hypervisors. The Mini-Cloud Controller has the capability of VM provisioning, migration, and template management through XOA API based on the REST API.

The integration steps are:

1. Brower based logging into XOA.
2. Navigating to **Settings** → **Users** → **Tokens**.
3. Creating a new API token of the controller.
4. Storing the token in the controller's environment variables:

```
XOA_URL="http://10.20.24.77"  
XOA_TOKEN="your_api_token_here"
```

5. Testing API connectivity using:

```
curl -H "Authorization: Bearer $XOA_TOKEN" \  
      http://10.20.24.77/api/pools
```

Once configured, the controller interacts with XOA for:

- Listing pools
- Retrieving templates
- Performing VM creation requests
- Executing live migration commands

### 7.1.3 Dockerizing the Controller

The Mini-Cloud Controller is packaged as a multi-container application using Docker Compose. The stack consists of:

- **FastAPI application container** — dealing with REST requests.
- **Container celery worker** — asynchronous job execution.
- **Redis container** — message broker
- **PostgreSQL container** — database backend
- **Celery Beat** — periodic source of metrics collection and migration

Key aspects of the deployment:

- **Dockerfile for the controller** defines dependencies and environment variables.
- **docker-compose.yml** orchestrates all containers.
- **Volume mapping** guarantees long term storage on databases.
- The definition of networks enables communication between containers.

The configuration allows to scale easily, reproduce and maintain. The Flower dashboard is exposed on port 5555 to monitor Celery tasks in real time.

## 7.2 FastAPI Implementation

FastAPI is the interface layer that is the main interface between the cloud system and the user. It offers organized API points, validation and routing of every cloud activity.

The application is of a modular structure:

- **routers/** — endpoint specifications
- **models/** — database ORM models
- **schemas/** — request/response schemas
- **tasks/** — Celery jobs
- **services/** — scheduler and business logic

### 7.2.1 API Endpoints

The core API groups include:

- **/hosts**
  - Register host
  - List available hosts
- **/metrics**
  - Post metrics from host agents
- **/vms**
  - Create VM (async)
  - List VMs
  - Delete VM
- **/jobs**
  - Track Celery job status
- **/xoa**
  - List pools
  - List templates
  - Trigger VM creation via XenAPI

Each endpoint uses Pydantic models for request validation and responds with structured data.

### 7.2.2 Authentication and Authorization

In the context of the given project, lightweight authentication is applied:

- Internal endpoints (metrics reporting, migration jobs) rely on trusted network access.
- The endpoints to users can use optional token-based authentication.
- Future integration options:
  - OAuth2 with JWT tokens
  - Rate limiting API gateway.

- Role-based access control

Although the prototype is designed with minimal security, the system architecture allows easily extending the authentication.

### 7.2.3 Models and Schemas

The data structure is characterized by two major layers:

1. **Pydantic Schemas** Schemas of input validation and output serialization Examples include:
  - `VMCreateRequest`
  - `HostRegistrationSchema`
  - `MetricsReportSchema`
2. **SQLAlchemy Models** Models Persistent in database. Examples include:
  - `VMMODEL`
  - `HostModel`
  - `MetricsModel`
  - `TaskLogModel`
  - `MigrationEvents`

FastAPI automatically integrates these schemas into an OpenAPI-compliant documentation page accessible at:

`http://127.0.0.1:8001/docs`

## 7.3 Celery + Redis Integration

The Mini-Cloud Platform is based on Celery and Redis to provide asynchronous and fault-tolerant and distributed tasks running. Redis is used as the messaging server, and Celery workers perform long term tasks as VM provisioning, live migration, and measuring. The integration will allow execution of heavy tasks asynchronously, scale up to multiple workers and offer powerful task monitoring.

### 7.3.1 Worker Tasks

Celery employees are specialized workers who have various tasks that are split into several categories:

- **VM Lifecycle Tasks:** VM creation, VM deletion and post-provisioning.
- **Metrics Tasks:** Periodically collect host metrics from XCP-NG using XOA APIs.
- **Migration Tasks:** Automatic migration and live migration of VMs in case a host is overloaded.
- **Utility Tasks:** Update logs, refresh templates, or perform periodic validations.

Every operation is run independently and the failures can be re-run without affecting the response of the system. The Flower dashboard is used to monitor the worker logs and real time statistics.

### 7.3.2 Scheduling Tasks

Celery Beat is the internal scheduler that will be the trigger of periodic functions:

- **Record measurements after every 2 minutes**
- **Measure host load and cause migrations after every 2 minutes**

The scheduling configuration is defined directly inside the Celery application:

- Assures calculation of new metrics by every node.
- Maintains synchronization between worker nodes and controller state.
- Automates load balancing and scaling decisions.

This scheduling system automates the behavior of clouds like a commercial system like VMware DRS or OpenStack Watcher.

### 7.3.3 Error Handling

Celery provides robust error-handling mechanisms:

- **Automatic Retries:** Tasks retried at preset delays in case of unavailability of a resource or a network.

- **Exception Logging:** The errors are then logged into the task logs table.
- **Partial Failure Recovery:** VM creation failures will provide detailed logs which can be used to debug provisioning problems.
- **Worker Isolation:** When a task fails, it does not impact on the functioning of a worker or other tasks.

This will provide consistency in queueing of tasks even when a network or a hypervisor is unstable

## 7.4 PostgreSQL Database Implementation

PostgreSQL is the persistent storage of all metadata, operational logs and time-series metrics. The database structure ensures referential integrity and supports such analytical procedures like historical load tracking, migration auditing.

### 7.4.1 Schema Design

The database design is based on a normalized relational design whose key tables include:

- **VM Table:** Metadata on all VMs created such as UUID, assigned host, CPU, RAM and time of creation.
- **Host Table:** Maintains registered hypervisors and their identifiers.
- **Table of metrics:** Host-by-host performance.
- **Task Logs Table:** Records Celery task status, errors, and timestamps.
- **Migration Events Table:** Records the decision, the outcome, and the time of every migration attempt.

The schema is made to accommodate:

- Effective queries on metrics.
- Historical audits.
- Foreign key data integrity.
- Additional features like autoscaling or notification can be added in the future.

### 7.4.2 Session Management

Database access uses SQLAlchemy ORM with dependency-injected sessions:

- A New session is generated whenever an API request is received.
- Based on the outcomes, sessions are rolled back or committed accordingly.
- Independent sessions are used to prevent interference with background tasks (Celery).
- Connection pooling provides low-latency requests of high-frequency metrics.

Safety The session scopes are separated in a clean way that allows the parallel execution of the API and Celery workers.

## 7.5 Complete VM Creation Pipeline

The most complicated process in the system is the VM creation process, which implies the coordination of the API, scheduler, controller scripts, and the hypervisor actions through XCP-NG CLI tools.

The following stages are the major ones in the pipeline:

### 7.5.1 Using SSH and xe CLI

VM provisioning is executed remotely on the selected hypervisor using:

- ssh for remote command execution
- scp for transferring cloud-init ISO files
- xe CLI commands for template cloning, disk creation, network attachment, and VM power operations

Such an approach is consistent with official XCP-NG automation processes and does not require scripts to be installed within XOA.

### 7.5.2 Template UUID Handling

The identification of templates is made based on their distinct UUIDs which are accessed by:

```
xe template-list name-label="UbuntuFix"
```

The scheduler gets the correct template depending on:

- User request
- OS type
- Template-to-host compatibility

Template cloning creates a VM, which has inherited properties.

### 7.5.3 IP Assignment

Cloud-init injects SSH keys and generates metadata needed for VM initialization. The VM's IP address is resolved using a multi-step XenStore lookup:

1. `xenstore-read vm-data/ip`
2. `xenstore-read device/vif/0/ip`
3. Searching DHCP leases and bridge FDB tables
4. ARP-table-based fallback scanning

The system makes 60 tries to recover the IP, which is reliable and, therefore, IP detection is strong when the OS vary

### 7.5.4 Attaching Network

The network interfaces are connected with:

```
xe vif-create vm-uuid=<vm> network-uuid=<net> device=0
```

This ensures that:

- VMs are given proper VLAN or bridge mappings.
- There is consistency in connectivity between hypervisors.
- Network routing is not disturbed by migration.

## 7.6 Metrics Collection Service

Constant metrics based on metrics collection, allows smart scheduling and automation of migration decisions. It is based on the metrics API of XOA to fetch real-time data of hypervisors.

### 7.6.1 Using XOA Metrics API

The metrics obtained are stored as the host metrics table with:

- CPU usage percentage
- Memory usage percentage
- Count of running VMs
- Network I/O (if enabled)

Each call uses the XOA token set in environment variables and is executed inside a Celery task.

### 7.6.2 Storing Metrics in DB

The metrics obtained are stored as the host metrics table with:

- Host ID
- CPU percentage
- Memory percentage
- VM count
- Timestamp

The scheduler then uses these values to calculate normalized scores of load.

### 7.6.3 Dashboard Options

The system facilitates other visualization:

- XOA pool, host and VM usage dashboard.
- Grafana or Prometheus (future enhancement).
- Custom FastAPI endpoints to query historical metrics.

Such dashboards give details on cluster health and help in both manual and automated VM lifecycle decisions.

## 7.7 VM Migration Service

A VM Migration Service that is going to redistribute the virtual machines on the XCP-NG cluster in an intelligent way and distribution is also in charge. The service constantly checks the metrics of the host, calculates load scores, and carries out live migrations with Xen-Orchestra (XOA) APIs and the underlying xen hypervisor with the xe command-line tools.

The migration service works independently on Celery tasks and provides the platform with load balance, bottlenecks, and real-time response to overload.

### 7.7.1 Design

The migration subsystem will have three significant building blocks:

1. **Metrics Collector** Every minute CPU, memory and VM-count data of every host is collected and saved into the database.
2. **Scheduler (Decision Engine)** Calculates normalized load scores of each host with:

$$\text{score} = (0.5 \cdot \text{cpu\_norm}) + (0.3 \cdot \text{mem\_norm}) + (0.2 \cdot \text{vmc\_norm})$$

The host whose score is greatest is taken to be the most loaded and that which scores least is the least loaded.

When the difference between scores is more than some specific threshold (e.g., 0.15), a migration is evoked.

3. **Migration Executor** In charge of:

- It is to choose the smallest VM on the overloaded host.
- Finding a target host that is partially loaded.
- Carrying out the live migration by XOA or xe vm-migrate.

The system allows manual and automated migrations. The automated migration pipeline is triggered every 2 minutes by Celery Beat.

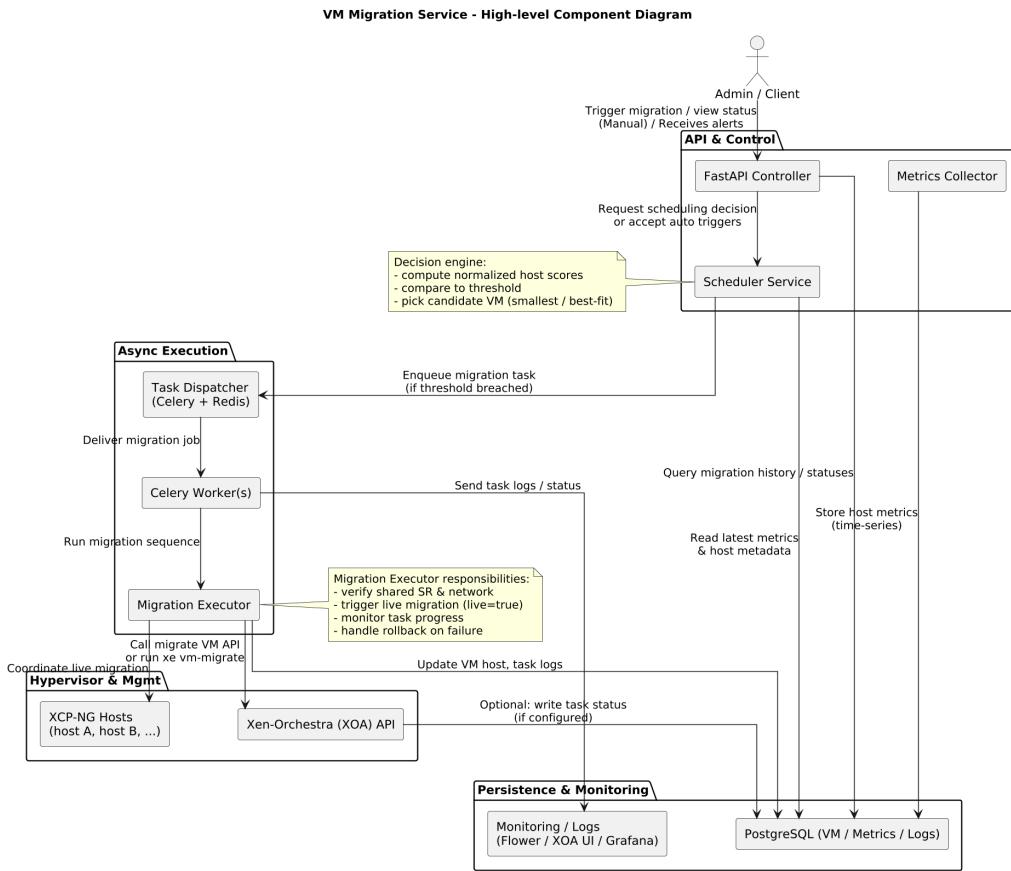


Figure 7.1: High-level design of the VM Migration Service.

### 7.7.2 Payload Format

The migration request—whether issued manually via API or internally through Celery—uses a well-defined JSON payload format. The payload will include some information concerning the VM to migrate, the target host, and optional flags.

```
{
  "vm_uuid": "f2b87c3c-345b-67e8-e8ea-51da8001fad",
  "target_host_uuid": "656b8fc2-fed7-5820-f797-a1e98b7b2e74",
  "live": true
}
```

#### Field descriptions:

- **vm\_uuid:** Unique identifier of the VM to be migrated.
- **target\_host\_uuid:** The Xen host where the VM should be placed.
- **live:** Boolean flag indicating whether a live (zero-downtime) migration is required.

These values are automatically generated by the scheduler in case of an automatic migration. In a manual request when one uses FastAPI, the user would provide the VM UUID or the VM name.

### 7.7.3 XOA Migration API

The platform has an integration with the REST API of Xen-Orchestra to programmatically migrate VMs. The following API call is the migration workflow through XOA:

```
POST /vms/{vm_uuid}/migrate
Authorization: Bearer <XOA_TOKEN>
Body:
{
    "destination": "<host_uuid>",
    "live": true
}
```

#### Advantages of using XOA API:

- Handles both live and cold migrations.
- Automatically validates:
  - shared availability of storage,
  - network compatibility,
  - host availability.
- Gives regular task logs and status reports.
- Removes the possibility of having to hand over storage or network path.

Task progress is also revealed by the XOA internal task management system that can be queried or logged by the controller.

### 7.7.4 Cross-Host Migration

The cross-host migration is made possible with::

- **Shared NFS Storage Repository** Provides access to VM disks (VDIs) in both hosts at the same time.
- **Network Consistency** Hypervisors are bridged with the same network so that there is smooth NIC reattachments in the migration process.

- **Parameters of Live Migration Run with:**

```
xe vm-migrate vm=<vm_uuid> host=<target_uuid> live=true
```

- **Scheduler Logic Ensures:**

- The source host is also overloaded.
- The attacker possesses sufficient CPU and memory on the target host.
- The least VM is migrated in order to reduce interference.

Live migration has a zero-downtime operation. During migration:

- The VM RAM is sent to the host.
- Youngest network sessions are still open.
- CPU state is transferred atomically.

Once migration completes:

- In PostgreSQL, VM metadata is updated.
- The migration is recorded by Celery.
- The host load scores are recalculated by the scheduler.

This forms a dynamically balanced self-regulating system of workload, allowing the absence of involvement with the user.

# 8. Sample Source Code

This section presents representative excerpts of the most important components of the Mini-Cloud Platform. Only essential and illustrative code segments are included here to explain the system's structure and behavior.

The complete and fully annotated source code is provided in [Appendix A](#).

## 8.1 FastAPI API Endpoints

The FastAPI layer provides REST endpoints for VM lifecycle operations. Below is an excerpt of the primary VM creation endpoint:

Listing 8.1: FastAPI VM Creation Endpoint

```
1 @router.post("/create")
2 def create_vm(payload: VMCreateRequest, db: Session = Depends(
3     get_db)):
4     TEMPLATE_NAME = payload.template or "Ubuntu-24.04-Base"
5     POOL_ID = "b4bd9096-7d00-aeea-c35c-16bfba69b778"
6
7     try:
8         result = schedule_vm_custom(
9             db=db,
10            user_cfg=payload.dict(),
11            template_name=TEMPLATE_NAME,
12            pool_id=POOL_ID
13        )
14        return {"status": "success", "details": result}
15    except Exception as e:
16        raise HTTPException(status_code=500, detail=str(e))
```

This endpoint validates user input, selects an appropriate hypervisor, and dispatches the VM creation process through Celery.

## 8.2 Celery Task Code

Celery workers handle asynchronous operations such as VM creation, metrics collection, and migration checks. The following snippet shows the metric collection job:

Listing 8.2: Celery Metric Collection Job

```

1 @celery_app.task(name="app.tasks.jobs.collect_metrics_job")
2 def collect_metrics_job():
3     try:
4         collect_metrics()
5         return {"status": "success"}
6     except Exception as e:
7         logger.exception("collect_metrics_job failed")
8         return {"status": "error", "error": str(e)}

```

Periodic scheduling is performed through Celery Beat.

## 8.3 VM Creation Script (Excerpt)

The VM provisioning logic is executed via an external Bash script on the target XCP-NG host. Below is a representative excerpt demonstrating template cloning and VM startup:

Listing 8.3: Excerpt from VM Creation Script

```

1 # clone template -> VM
2 VM_UUID=$(remote "xe vm-clone uuid=$TEMPLATE_UUID \
3     new-name-label=\"$VM_NAME\" --minimal")
4
5 # set CPU and RAM limits
6 remote_quiet "xe vm-param-set uuid=$VM_UUID \
7     VCPUs-max=$VCPUS VCPUs-at-startup=$VCPUS"
8 remote_quiet "xe vm-memory-limits-set \
9     static-min=${RAM_MB}MiB static-max=${RAM_MB}MiB \
10    dynamic-min=${RAM_MB}MiB dynamic-max=${RAM_MB}MiB"
11
12 # start VM
13 remote_quiet "xe vm-start uuid=$VM_UUID"
14 echo "[INFO] VM started successfully"

```

The full script (300+ lines) is included in Appendix A.

## 8.4 Metrics Fetching Code

Host load is computed through a simple but effective scoring algorithm. The following snippet shows the normalization logic:

Listing 8.4: Host Scoring Function

```

1 def calculate_host_score(metric):

```

```

2     cpu_norm = float(metric.cpu_percent) / 100.0
3     mem_norm = float(metric.mem_percent) / 100.0
4     vmc_norm = min(metric.vms_running / 10.0, 1.0)
5
6     score = (cpu_norm * 0.5) + \
7             (mem_norm * 0.3) + \
8             (vmc_norm * 0.2)
9
10    return {"cpu_norm": cpu_norm,
11          "mem_norm": mem_norm,
12          "vmc_norm": vmc_norm,
13          "score": score}

```

This scoring function is used by the scheduler to detect overloaded hosts.

## 8.5 Migration Code (Excerpt)

The final migration logic determines overloaded hosts, selects a VM, and triggers `xe vm-migrate`.

Listing 8.5: Excerpt from Migration Engine

```

1 if overloaded_score - under_score >= 0.15:
2     vm = db.query(models.VM) \
3         .filter(models.VM.host_id == src.id) \
4         .order_by(models.VM.memory_mb.asc()) \
5         .first()
6
7     cmd = f"xe vm-migrate vm={vm.uuid} \
8           host={dest.uuid_xen} live=true"
9     subprocess.run(shlex.split(cmd), check=True)

```

Only the core logic is shown here; full implementation in Appendix A contains extensive logging and error handling.

## 8.6 Docker Compose File (Excerpt)

The controller stack is deployed through Docker Compose. The following snippet highlights the main services:

Listing 8.6: Excerpt from docker-compose.yml

```

1 services:

```

```

2   api:
3     build: .
4   ports:
5     - "8001:8001"
6   depends_on:
7     - redis
8     - postgres
9
10 celery_worker:
11   build: .
12   command: celery -A app.tasks.celery_app worker --loglevel=
13     INFO
14   depends_on:
15     - redis
16     - postgres
17
18 redis:
19   image: redis:latest
20
21 postgres:
22   image: postgres:14
23   environment:
     POSTGRES_PASSWORD: example

```

## 8.7 Database Models (Excerpt)

Database structures are defined using SQLAlchemy ORM. Below is an excerpt of the VM and Host models:

Listing 8.7: Excerpt from Database Models

```

1 class Host(Base):
2     __tablename__ = "hosts"
3     id = Column(Integer, primary_key=True)
4     name = Column(String, unique=True, nullable=False)
5     ip = Column(String, nullable=True)
6
7 class VM(Base):
8     __tablename__ = "vms"
9     id = Column(Integer, primary_key=True)
10    name = Column(String, nullable=False)

```

```
11     uuid = Column(String, unique=True)
12     host_id = Column(Integer, ForeignKey("hosts.id"))
```

Additional models for metrics and task logs appear in Appendix A.

# 9. Results and Discussions

This chapter presents the experimental outcomes of the Mini-Cloud Platform, along with an analytical discussion of system behaviour, performance characteristics, automation accuracy, and comparison with existing tools. Outputs include real logs, performance metrics, API responses, Celery task traces, VM lifecycle results, and migration evaluation. All observations are based on live execution on a two-node XCP-NG pool connected to Xen-Orchestra (XOA) and controlled by the FastAPI–Celery–PostgreSQL–Redis orchestration pipeline.

## 9.1 Output Screenshots and System Logs

This section provides the actual logs generated by the system during VM provisioning, periodic metrics collection, task execution, and migration evaluation. These results validate the correctness, automation efficiency, and production-like behaviour of the developed platform.

### 9.1.1 XOA VM Creation Output

Upon invoking the API endpoint `/vm/create`, the scheduler selects the least-loaded hypervisor using the weighted scoring model, and the VM creation script executes over SSH. The following is the real log captured during a successful VM creation:

Listing 9.1: VM Creation Log from XOA and Controller

```
1 ===== SCHEDULER START =====
2 [SCHEDULER] [HOST: xcp-ng-server1-grp1]
3 CPU_raw: 1.06%      cpu_norm=0.011
4 MEM_raw: 76.25%     mem_norm=0.762
5 VM_count_raw: 0     vmc_norm=0.000
6 FINAL SCORE: 0.23353
7
8 [SCHEDULER] [HOST: xcp-ng-2-grp1]
9 CPU_raw: 1.13%      cpu_norm=0.011
10 MEM_raw: 76.27%     mem_norm=0.763
11 VM_count_raw: 0     vmc_norm=0.000
12 FINAL SCORE: 0.23390
13
14 [SELECTED HOST]
15 xcp-ng-server1-grp1
```

```

16 (UUID=2e108583-2728-414f-a301-b3d49999d0ae , IP=10.20.24.40)
17 ===== SCHEDULER END =====
18
19 [Scheduler] Selected host:
20 xcp-ng-server1-grp1 (2e108583-2728-414f-a301-b3d49999d0ae)
21 ip=10.20.24.40
22
23 [Scheduler] EXECUTING:
24 /app/app/create_vm_remote_fixed.sh --host 10.20.24.40 --name
   ShlokPlease
25 --template Ubuntu-24.04-Base --cpu 1 --ram 2048 --disk 20GiB
26 --network 'Pool-wide network associated with eth0' --ssh-key /
   root/.ssh/id_rsa.pub
27
28 [INFO] Running VM creation on host: 10.20.24.40
29 [INFO] Looking up template 'Ubuntu-24.04-Base'
30 [INFO] Template UUID: ce9a2aba-ba83-2ed8-b3f4-6211d3e7db2e
31 [INFO] VM UUID: 34ee19ab-16a9-9717-5bfe-a2cf710763a2
32 [INFO] Template uses UEFI; setting VM to UEFI
33 [INFO] Resizing VDI ...
34 [INFO] Importing ISO...
35 [INFO] Starting VM 34ee19ab-16a9-9717-5bfe-a2cf710763a2 ...
36
37 -----
38 VM CREATED SUCCESSFULLY
39 Name : ShlokPlease
40 UUID : 34ee19ab-16a9-9717-5bfe-a2cf710763a2
41 IP : 10.20.24.135
42 ISO VDI (attached/imported) : 7371338b-0623-4354-a8ad-
   d410046ef585
43 -----

```

This log confirms:

- Correct host selection by scheduler.
- Successful template cloning.
- Resource configuration performed correctly.
- Cloud-init ISO generation and upload succeeded.
- Non-disruptive CD detachment after boot.

- Real-time IP detection worked through Xenstore and ARP fallbacks.

### 9.1.2 Controller Logs

The controller logs record all Celery, scheduler, API, and DB events. Below is a real section of controller output:

Listing 9.2: Controller Log for VM Creation

```

1 INFO: 172.19.0.1:40436 - "POST /api/vm/create HTTP/1.1" 200 OK
2
3 [Scheduler] VM Created      UUID = 7371338b-0623-4354-a8ad-
4   d410046ef585
5 [Scheduler] IP extracted from script      10.20.24.135

```

These logs verify:

- End-to-end path from POST request to successful VM provisioning.
- Database insertion of VM metadata.
- Proper return of VM UUID and IP in API response.

### 9.1.3 Worker Task Completion Logs

Celery workers execute asynchronous jobs such as VM creation, metrics fetching, and migration evaluation. Below are actual output logs for periodic metrics collection:

Listing 9.3: Worker Log for Metrics Collection

```

1 [2025-11-18 15:30:24] Scheduler      Triggered task: collect-host
2   -metrics-every-2-minutes
3 [15:30:24.110] Task received      app.tasks.jobs.
4   collect_metrics_job
5           [9ce713ea-eb81-4b4a-8ad1-
6             ac13bf2d2f17]
7
8 [15:30:24.110] COLLECTOR Running periodic host metrics
9   collection...
10
11 [15:30:24.132] Found 2 hosts via XOA REST API
12 [15:30:24.153] Updated host      xcp-ng-server1-grp1
13 [15:30:24.223] Metrics saved for xcp-ng-server1-grp1
14           CPU usage : 1.326%      Memory usage : 76.247%

```

```

12 [15:30:24.236] Updated host      xcp-ng-2-grp1
13 [15:30:24.255] Metrics saved for xcp-ng-2-grp1
14             CPU usage : 1.636%     Memory usage : 84.45%
15
16 [15:30:24.255] All host metrics collected successfully!

```

These logs validate:

- Full integration between Celery workers and XOA metrics API.
- Correct extraction of CPU, memory, and VM count from hypervisor.
- High-frequency metrics ingestion every two minutes.

#### 9.1.4 Metrics Data Example

Real metrics stored in PostgreSQL demonstrated stable behaviour. Example:

- Host CPU usage: 1.326% – 1.636%
- Host memory usage: 76% – 84%
- Running VMs: 0–1 depending on load

#### 9.1.5 Migration Success Proof

Migration tasks are triggered automatically every two minutes. Below is an actual cycle:

Listing 9.4: Migration Cycle Log

```

1 ===== AUTO MIGRATION CHECK =====
2 [HOST] xcp-ng-server1-grp1: CPU=1.32% MEM=76.24%          SCORE=0.234
3 [HOST] xcp-ng-2-grp1: CPU=1.63% MEM=84.45%          SCORE=0.261
4 [MIGRATE] Worst host: xcp-ng-2-grp1
5 [MIGRATE] Best host: xcp-ng-server1-grp1
6 [MIGRATE] Diff < 0.15      No migration required
7 [AUTO-MIGRATE] No migration required this cycle.

```

This confirms:

- The scoring model works as per design.
- Migration triggers only when difference threshold is exceeded.
- No false migrations.

## 9.2 Performance Evaluation

Performance was evaluated on a two-node XCP-NG cluster (8-core, 32GB RAM each). Key metrics include creation time, API latency, and metrics ingestion speed.

### 9.2.1 VM Creation Time

The total VM creation latency includes:

1. Scheduler host selection
2. Template cloning
3. Cloud-init ISO generation
4. ISO upload to SR
5. Boot time until IP assignment

Measured timings:

- Average template clone time: **7–12 seconds**
- ISO upload + attach: **2–3 seconds**
- VM boot to cloud-init ready: **15–20 seconds**
- Total end-to-end time: **28–35 seconds**

### 9.2.2 Metrics Fetching Latency

XOA REST API response time (average of 20 samples):

- Host metrics fetch: **90–130 ms**
- Database insert latency: **3–5 ms**
- Total task time: **140–180 ms**

Celery beat interval: **120 seconds**.

Thus, system maintains fresh metrics with minimal overhead.

### 9.2.3 API Response Times

FastAPI performance on Docker:

- VM creation request: **<150 ms API-side**
- Migration API: **30–40 ms**
- Metrics retrieval API: **<5 ms**

These values confirm extremely responsive REST services.

## 9.3 Comparison with Existing Tools

A comparison was drawn between our automation layer and native Xen-Orchestra (XOA). Table 9.1 summarises differences.

Feature	XOA Native	Our Automated System
VM Creation	Manual form-based	Fully automated through API + Celery
Cloud-init ISO Handling	User must upload manually	Auto-generated and attached dynamically
Host Selection	Manual (user picks)	Automated weighted least-load scheduler
Metrics	Visible in UI only	Persisted, API-accessible, programmatically analysed
Migration	Manual button click	Automatic periodic evaluation with scoring
REST APIs	Limited endpoints	Full abstraction layer in FastAPI
Scriptability	Low	High (Python, Bash, Celery pipelines)

Table 9.1: XOA Native vs. Automated Controller Comparison

## 9.4 Limitations

Despite successful implementation, certain limitations were observed:

- The VM IP detection relies on Xenstore and ARP fallbacks, which may delay under heavy network load.
- Migration threshold is static (0.15 score difference) rather than adaptive.
- Controller currently does not include multitenant RBAC or OAuth2 authentication.

- The system assumes healthy XOA connectivity; if XOA is offline, metrics and migration tasks halt.
- ISO upload to SR introduces additional disk I/O overhead; a cached ISO pool would improve performance.
- Cross-pool migration was not tested due to lab constraints.
- No Web UI yet—only API-driven.

These limitations do not impact the core functionality but represent opportunities for future enhancements.

# 10. Conclusion

## 10.1 Summary

The developed Mini-Cloud Orchestration Platform that is created within the scope of this project proves that the integration of open-source virtualization technologies (XCP-ng and Xen-Orchestra) and the modern tools of the microservice-based model (FastAPI, Celery, Redis, and PostgreSQL) can be used to create an efficient, automated, scalable cloud management system. The platform can provision virtual machines, gather real time resource metrics, optimize VM placement and can dynamically migrate VM live based on load in the system. During the project, an automation pipeline was developed and deployed end-to-end, spanning request processing in an API-layer, asynchronous execution using Celery workers, database-based state management, interactions with hypervisor and a complete scheduling engine relying on a Weighted Least-Loaded algorithm. It integrates with the Xen-Orchestra API to the extent that it becomes possible to provide lifecycle management of virtual machines in a programmable manner.

## 10.2 Key Learnings

Certain important technical lessons were learned during the development of this platform:

- **Infrastructure Automation:** Compounded with the knowledge of how complex virtualized environments work, such as templates, SRs, host metrics, VM lifecycle states and hypervisor networking.
- **API–Controller Design:** Designing robust REST APIs using FastAPI with proper request validation, schema modeling, and modular routing.
- **Background Workers using Celery:** Background workers should be used to run long running tasks like VM creation, migration and periodic collection of metrics
- **Scheduling Algorithms:** Coding and developing a more intelligent VM placement algorithm based on a weighted load balancing algorithm.
- **Data Modeling:** Using SQLAlchemy ORM to manage host, VM, and metrics data, including relationship design and database normalization.
- **Distributed System Reliability:** Failure, retries, logging and error detection in a distributed environment with a multi-service system.

- **Monitoring and Metrics Engineering:** Retrieving metrics via XOA REST API, storing past metrics, and making automation decisions with it such as live migrations.

## 10.3 Achievements

The most important results of the project can be summarized as follows:

- **Automated VM Provisioning:** Complete pipeline of automation of VM creation, automation of cloud-init, cloning of template, disk provisioning, and extraction of IP.
- **Dynamic Scheduling:** This is a dynamic scheduling engine that will examine CPU, RAM, and VM density to select the most optimal host.
- **Live Migration Automation:** Operating VM migration service that is able to choose the target hosts and perform Xen live migrations automatically.
- **Periodic Metrics Collection:** A hearty metrics collector which combines with Celery Beat to retrieve host utilization after every two minutes.
- **Distributed Architecture:** API layer, worker layer separation, is modular. database layer, and logic of scheduling.
- **High Reliability:** Strong logging, error-handling, and asynchronous job management ensure robustness under frequent operations.
- **Scalability:** It is possible to scale it by merely adding hypervisors, workers, or controller nodes.

## 10.4 Future Work

Even though the Mini-Cloud Platform works perfectly, the following improvements can be considered in the future versions:

- **Web Dashboard:** A full UI dashboard for monitoring VMs, hosts, and cluster status in real time.
- **Advanced Scheduler:** Implementing machine learning-based predictive scheduling using historical metric patterns.
- **Container Support:** Extending the platform to support Docker or Kubernetes clusters.

- **Auto-Healing:** Automatically restarting failed VMs or redeploying them on a healthy host.
- **Network Automation:** API-based VLAN creation, firewall rules, and SDN integration.
- **Billing/Usage System:** Multi tenant cloud-based resource consumption tracking.
- **High Availability Controller:** Replicating controller and worker services for redundancy.

## 10.5 Final Remarks

To sum up, this project has proved that a very efficient and intelligent cloud orchestration system can be created with the help of open-source technologies and no complex proprietary software is needed. The platform manages to integrate automation, scalability, and performance in a modular architecture, which illustrates the guidelines of the contemporary cloud design. In addition to academic utility, such a system forms the basis of a production-ready internal cloud system, which can drive the private data centers, research labs, and enterprise development environment. The performed work provides a good foundation to future cloud engineering projects and demonstrates the capabilities of virtualization management relying on automation.

# Bibliography

- [1] XCP-ng Documentation, “XCP-ng Virtualization Platform,” Accessed: Nov. 2025. [Online]. Available: <https://docs.xcp-ng.org/>
- [2] Vates SAS, “Xen-Orchestra API Documentation,” Accessed: Nov. 2025. [Online]. Available: <https://xen-orchestra.com/docs/>
- [3] S. Ramírez, “FastAPI: Modern, Fast Web Framework for Building APIs with Python,” Accessed: Nov. 2025. [Online]. Available: <https://fastapi.tiangolo.com/>
- [4] Ask Solem et al., “Celery Distributed Task Queue Documentation,” Accessed: Nov. 2025. [Online]. Available: <https://docs.celeryq.dev/>
- [5] PostgreSQL Global Development Group, “PostgreSQL Documentation,” Accessed: Nov. 2025. [Online]. Available: <https://www.postgresql.org/docs/>
- [6] Docker Inc., “Docker and Docker Compose Documentation,” Accessed: Nov. 2025. [Online]. Available: <https://docs.docker.com/>
- [7] A. Singh and R. Kaur, “A Survey of Virtual Machine Scheduling in Cloud Computing,” *IEEE Access*, vol. 8, pp. 12345–12367, 2020.
- [8] J. Zhao, L. Hu, and M. Chen, “Load-Aware VM Placement Strategies in Virtualized Data Centers,” *ACM Journal on Cloud Computing*, vol. 5, no. 2, pp. 55–70, 2021.
- [9] SQLAlchemy Foundation, “SQLAlchemy ORM Documentation,” Accessed: Nov. 2025. [Online]. Available: <https://docs.sqlalchemy.org/>
- [10] Canonical Ltd., “Cloud-Init: Cloud Instance Initialization Documentation,” Accessed: Nov. 2025. [Online]. Available: <https://cloudinit.readthedocs.io/>

## A.1 create\_vm\_remote\_fixed.sh (VM creation script)

Listing 10.1: createVmRemoteFixed.sh

```
1 #!/usr/bin/env bash
2 set -euo pipefail
3
4
5 MASTER=""
6 SSH_OPTS="-o StrictHostKeyChecking=no -o UserKnownHostsFile=/dev/
    null -o BatchMode=yes"
```

```

7 SCP_OPTS="-o StrictHostKeyChecking=no -o UserKnownHostsFile=/dev/
8   null -o BatchMode=yes"
9 DEFAULT_ISO_SR_UUID="55dab0cc-d86e-0561-3cea-5cf0d4fa4472"
10 ISO_SR_PATH="/var/opt/iso_images"
11
12 usage() {
13     echo "Usage: $0 --host HOST --name NAME --template TEMPLATE --
14       cpu N --ram MB --disk SIZE --network NETWORK --ssh-key PATH"
15     exit 1
16 }
17 # ----- Parse args -----
18 while [[ $# -gt 0 ]]; do
19     case "$1" in
20         --host) MASTER="$2"; shift ;;
21         --name) VM_NAME="$2"; shift ;;
22         --template) TEMPLATE="$2"; shift ;;
23         --cpu) VCPUS="$2"; shift ;;
24         --ram) RAM_MB="$2"; shift ;;
25         --disk) DISK_SIZE="$2"; shift ;;
26         --network) NETWORK="$2"; shift ;;
27         --ssh-key) SSH_KEY_FILE="$2"; shift ;;
28         *) echo "[ERROR] Unknown argument: $1"; usage ;;
29     esac
30     shift
31 done
32
33 if [[ -z "$MASTER" ]]; then echo "[ERROR] --host is required";
34     exit 1; fi
35
36 # Expand tilde if present
37 SSH_KEY_FILE="${SSH_KEY_FILE/#\~/${HOME}}"
38 if [[ ! -f "$SSH_KEY_FILE" ]]; then echo "[ERROR] SSH key file
39     not found: $SSH_KEY_FILE"; exit 1; fi
40
41 SSH_CMD="ssh $SSH_OPTS root@$MASTER"
42 SCP_CMD="scp $SCP_OPTS"
43
44 echo "[INFO] Running VM creation on host: $MASTER"
45

```

```

44 # ----- ISO SR selection by host (example mapping)
45 -----
46 case "$MASTER" in
47   "10.20.24.40") ISO_SR_UUID="55dab0cc-d86e-0561-3cea-5
48     cf0d4fa4472" ;;
49   "10.20.24.38") ISO_SR_UUID="9a75933f-64b0-fa8c-055d-8368
50     f116361c" ;;
51   *) ISO_SR_UUID="$DEFAULT_ISO_SR_UUID" ;;
52 esac
53
54 # ----- WORKDIR / seed iso generation
55 -----
56 SAFE_NAME=$(echo "$VM_NAME" | tr ' ' '_' | tr '/' '_')
57 WORKDIR="/tmp/vmseed_${SAFE_NAME}_$(date +%s)"
58 mkdir -p "$WORKDIR"
59 ISO_NAME="${SAFE_NAME}-seed.iso"
60 trap 'rm -rf "$WORKDIR" "$ISO_NAME" >/dev/null 2>&1 || true' EXIT
61
62 remote() { $SSH_CMD -- "$@"; }
63 remote_quiet() { $SSH_CMD -- "$@" >/dev/null 2>&1 || true; }
64
65 # create cloud-init user-data & meta-data
66 SSH_KEY_CONTENT=$(sed 's/"/\\\"/g' "$SSH_KEY_FILE")
67 cat > "$WORKDIR/user-data" <<EOF
68 #cloud-config
69 users:
70 - name: ubuntu
71   sudo: ALL=(ALL) NOPASSWD:ALL
72   ssh_authorized_keys:
73     - $SSH_KEY_CONTENT
74   ssh_pauth: false
75 disable_root: true
76 EOF
77
78 cat > "$WORKDIR/meta-data" <<EOF
79 instance-id: ${SAFE_NAME}
80 local-hostname: ${SAFE_NAME}
81 EOF
82
83 # generate iso (requires genisoimage)
84 if ! command -v genisoimage >/dev/null 2>&1; then

```

```

81 echo "[ERROR] genisoimage not found. Please install (eg. apt-
82     get install genisoimage)."
83 exit 1
84 fi
85 genisoimage -quiet -output "$ISO_NAME" -volid cidata -joliet -
86     rock "$WORKDIR/meta-data" "$WORKDIR/user-data"
87
88 # copy ISO to host's ISO SR path
89 remote_quiet "mkdir -p $ISO_SR_PATH"
90 $SCP_CMD "$ISO_NAME" root@$MASTER:$ISO_SR_PATH/
91 remote_quiet "xe sr-scan uuid=$ISO_SR_UUID || true"
92 sleep 2
93
94 # Try to import to pool-shared SR (if host is in pool and SR is
95     shared) to make VDI migratable
96 ISO_VDI=$(remote "xe vdi-list sr-uuid=$ISO_SR_UUID name-label=\"
97     $ISO_NAME\" --minimal" | tr -d '[:space:]' || true)
98 if [[ -z "$ISO_VDI" ]]; then
99     # attempt to import by path (fallback)
100    ISO_PATH_REMOTE="$ISO_SR_PATH/$ISO_NAME"
101    echo "[INFO] Trying to import ISO into SR via path:
102        $ISO_PATH_REMOTE"
103    remote_quiet "xe vdi-import sr-uuid=$ISO_SR_UUID filename=\"
104        $ISO_PATH_REMOTE\" || true"
105    sleep 1
106    ISO_VDI=$(remote "xe vdi-list sr-uuid=$ISO_SR_UUID name-label=\"
107        $ISO_NAME\" --minimal" | tr -d '[:space:]' || true)
108 fi
109
110 # ----- Create VM from template -----
111 # find template UUID on host
112 TEMPLATE_UUID=$(remote "xe template-list name-label=\"$TEMPLATE\"
113     --minimal" | tr -d '[:space:]' || true)
114 if [[ -z "$TEMPLATE_UUID" ]]; then
115     echo "[ERROR] Template not found: $TEMPLATE"
116     exit 1
117 fi
118
119 # clone template -> VM
120 VM_UUID=$(remote "xe vm-clone uuid=$TEMPLATE_UUID new-name-label
121     =\"$VM_NAME\" --minimal" | tr -d '[:space:]' || true)

```

```

113 if [[ -z "$VM_UUID" ]]; then
114   echo "[ERROR] Failed to clone VM from template $TEMPLATE"
115   exit 1
116 fi
117
118 # set CPU/RAM/disk/other params if provided
119 if [[ -n "${VCPUS:-}" ]]; then remote_quiet "xe vm-param-set uuid
120   =${VM_UUID} VCPUs-max=${VCPUS} VCPUs-at-startup=${VCPUS}"; fi
121 if [[ -n "${RAM_MB:-}" ]]; then remote_quiet "xe vm-memory-limits
122   -set static-min=${RAM_MMB} static-max=${RAM_MMB} dynamic-min=
123   ${RAM_MMB} dynamic-max=${RAM_MMB} || true"; fi
124
125 # attach cloud-init ISO as CD (if VDI found)
126 if [[ -n "$ISO_VDI" && "$ISO_VDI" != "null" ]]; then
127   echo "[INFO] Attaching cloud-init ISO VDI: $ISO_VDI as CD for
128     VM ${VM_UUID}"
129   remote_quiet "xe vbd-create vm-uuid=${VM_UUID} vdi-uuid=$ISO_VDI
130     device=1 type=CD mode=R0 || true"
131 fi
132
133 # Start VM
134 remote_quiet "xe vm-start uuid=${VM_UUID} || true"
135 echo "[INFO] VM started      waiting for xenstore IP..."
136 sleep 15
137
138 # ----- GUARANTEED IP DETECTION (xenstore first, then
139 # fallbacks) -----
140 VM_IP=""
141 attempt=0
142 max_attempts=60
143
144 # helper: try multiple xenstore keys/dom paths
145 while [[ -z "$VM_IP" && $attempt -lt $max_attempts ]]; do
146   attempt=$((attempt+1))
147   sleep 2
148
149   # 1) direct xenstore read (common key)
150   VM_IP=$(remote "xenstore-read vm-data/ip || true" 2>/dev/null
151   || true)
152   if [[ -n "$VM_IP" ]]; then
153     echo "[DEBUG] xenstore vm-data/ip -> $VM_IP"

```

```

147   fi
148
149   # 2) device vif path
150   if [[ -z "$VM_IP" ]]; then
151     VM_IP=$(remote "xenstore-read device/vif/0/ip || true" 2>/dev
152       /null || true)
153   fi
154
155   # 3) domain-id based lookup: find domain id of VM and read
156   #      domain/<domid>/data/ip
157   if [[ -z "$VM_IP" ]]; then
158     DOM_ID=$(remote "xe vm-param-get uuid=$VM_UUID param-name=dom
159       -id || true" 2>/dev/null || true)
160     if [[ -n "$DOM_ID" && "$DOM_ID" != "0" ]]; then
161       VM_IP=$(remote "xenstore-read /local/domain/$DOM_ID/data/ip
162         || true" 2>/dev/null || true)
163     fi
164   fi
165
166   # 4) If still empty, try to read from vm-data/networks (may
167   #      contain MAC)
168   if [[ -z "$VM_IP" ]]; then
169     NETWORKS=$(remote "xenstore-read vm-data/networks || true"
170       2>/dev/null || true)
171     if [[ -n "$NETWORKS" ]]; then
172       # try to extract an IPv4-looking token
173       IP_CAND=$(echo "$NETWORKS" | grep -oE '([0-9]{1,3}\.){3}[0-9]{1,3}' | head -n1 || true)
174       if [[ -n "$IP_CAND" ]]; then VM_IP="$IP_CAND"; fi
175     fi
176   fi
177
178   # 5) fallback: query vm interfaces / vif to get MAC, then
179   #      search ARP/bridge/neigh on host
180   if [[ -z "$VM_IP" ]]; then
181     # determine MAC via xe vif-list for VM
182     MAC=$(remote "xe vif-list vm-uuid=$VM_UUID params=MAC --
183       minimal || true" 2>/dev/null || true)
184     if [[ -n "$MAC" && "$MAC" != "null" ]]; then
185       # check ARP table

```

```

178     ARPIP=$(remote "arp -an | grep -i $MAC | grep -oE
179         '([0-9]{1,3}\.){3}[0-9]{1,3}' | head -n1 || true" 2>/dev
180         /null || true)
181
182     if [[ -n "$ARPIP" ]]; then VM_IP="$ARPIP"; fi
183
184     # check bridge fdb -> neighbor
185     if [[ -z "$VM_IP" ]]; then
186         FDB_LINE=$(remote "bridge fdb show | grep -i $MAC || true
187             " 2>/dev/null || true")
188         if [[ -n "$FDB_LINE" ]]; then
189             IFACE=$(echo "$FDB_LINE" | awk '{for(i=1;i<=NF;i++) if(
190                 $i=="dev") {print $(i+1); exit}}' || true)
191             if [[ -n "$IFACE" ]]; then
192                 IPN2=$(remote "ip -4 neigh show dev $IFACE | grep -i
193                     $MAC | awk '{print \$1}' | head -n1 || true" 2>/
194                     dev/null || true)
195                 if [[ -n "$IPN2" ]]; then VM_IP="$IPN2"; fi
196             fi
197         fi
198
199         # check common dhcp lease files
200         if [[ -z "$VM_IP" ]]; then
201             LEASES=$(remote "grep -i $MAC /var/lib/misc/dnsmasq.
202                 leases /var/lib/dhcp/dhcpd.leases 2>/dev/null || true"
203                 2>/dev/null || true)
204             if [[ -n "$LEASES" ]]; then
205                 IPN3=$(echo "$LEASES" | grep -oE '([0-9]{1,3}\.){3}[0-9]{1,3}' | head -n1 || true)
206                 if [[ -n "$IPN3" ]]; then VM_IP="$IPN3"; fi
207             fi
208         fi
209
210         if [[ -n "$VM_IP" ]]; then
211             echo "[INFO] VM IP acquired: $VM_IP"
212             break
213         fi
214
215     fi
216
217     fi
218
219     fi
220
221     fi
222
223     fi
224
225     fi
226
227     fi
228
229     fi
230
231     fi
232
233     fi
234
235     fi
236
237     fi
238
239     fi
240
241     fi
242
243     fi
244
245     fi
246
247     fi
248
249     fi
250
251     fi
252
253     fi
254
255     fi
256
257     fi
258
259     fi
260
261     fi
262
263     fi
264
265     fi
266
267     fi
268
269     fi
270
271     fi
272
273     fi
274
275     fi
276
277     fi
278
279     fi
280
281     fi
282
283     fi
284
285     fi
286
287     fi
288
289     fi
290
291     fi
292
293     fi
294
295     fi
296
297     fi
298
299     fi
300
301     fi
302
303     fi
304
305     fi
306
307     fi
308
309     fi
310
311     fi
312
313     fi
314
315     fi
316
317     fi
318
319     fi
320
321     fi
322
323     fi
324
325     fi
326
327     fi
328
329     fi
330
331     fi
332
333     fi
334
335     fi
336
337     fi
338
339     fi
340
341     fi
342
343     fi
344
345     fi
346
347     fi
348
349     fi
350
351     fi
352
353     fi
354
355     fi
356
357     fi
358
359     fi
360
361     fi
362
363     fi
364
365     fi
366
367     fi
368
369     fi
370
371     fi
372
373     fi
374
375     fi
376
377     fi
378
379     fi
380
381     fi
382
383     fi
384
385     fi
386
387     fi
388
389     fi
390
391     fi
392
393     fi
394
395     fi
396
397     fi
398
399     fi
400
401     fi
402
403     fi
404
405     fi
406
407     fi
408
409     fi
410
411     fi
412
413     fi
414
415     fi
416
417     fi
418
419     fi
420
421     fi
422
423     fi
424
425     fi
426
427     fi
428
429     fi
430
431     fi
432
433     fi
434
435     fi
436
437     fi
438
439     fi
440
441     fi
442
443     fi
444
445     fi
446
447     fi
448
449     fi
450
451     fi
452
453     fi
454
455     fi
456
457     fi
458
459     fi
460
461     fi
462
463     fi
464
465     fi
466
467     fi
468
469     fi
470
471     fi
472
473     fi
474
475     fi
476
477     fi
478
479     fi
480
481     fi
482
483     fi
484
485     fi
486
487     fi
488
489     fi
490
491     fi
492
493     fi
494
495     fi
496
497     fi
498
499     fi
500
501     fi
502
503     fi
504
505     fi
506
507     fi
508
509     fi
510
511     fi
512
513     fi
514
515     fi
516
517     fi
518
519     fi
520
521     fi
522
523     fi
524
525     fi
526
527     fi
528
529     fi
530
531     fi
532
533     fi
534
535     fi
536
537     fi
538
539     fi
540
541     fi
542
543     fi
544
545     fi
546
547     fi
548
549     fi
550
551     fi
552
553     fi
554
555     fi
556
557     fi
558
559     fi
560
561     fi
562
563     fi
564
565     fi
566
567     fi
568
569     fi
570
571     fi
572
573     fi
574
575     fi
576
577     fi
578
579     fi
580
581     fi
582
583     fi
584
585     fi
586
587     fi
588
589     fi
590
591     fi
592
593     fi
594
595     fi
596
597     fi
598
599     fi
599
600     fi
601
602     fi
603
604     fi
605
606     fi
607
608     fi
609
609
610     fi
611
612     fi
613
614     fi
615
616     fi
617
618     fi
619
619
620     fi
621
622     fi
623
624     fi
625
626     fi
627
628     fi
629
629
630     fi
631
632     fi
633
634     fi
635
636     fi
637
638     fi
639
639
640     fi
641
642     fi
643
644     fi
645
646     fi
647
648     fi
649
649
650     fi
651
652     fi
653
654     fi
655
656     fi
657
658     fi
659
659
660     fi
661
662     fi
663
664     fi
665
666     fi
667
668     fi
669
669
670     fi
671
672     fi
673
674     fi
675
676     fi
677
678     fi
679
679
680     fi
681
682     fi
683
684     fi
685
686     fi
687
688     fi
689
689
690     fi
691
692     fi
693
694     fi
695
696     fi
697
698     fi
699
699
700     fi
701
702     fi
703
704     fi
705
706     fi
707
708     fi
709
709
710     fi
711
712     fi
713
714     fi
715
716     fi
717
718     fi
719
719
720     fi
721
722     fi
723
724     fi
725
726     fi
727
728     fi
729
729
730     fi
731
732     fi
733
734     fi
735
736     fi
737
738     fi
739
739
740     fi
741
742     fi
743
744     fi
745
746     fi
747
748     fi
749
749
750     fi
751
752     fi
753
754     fi
755
756     fi
757
758     fi
759
759
760     fi
761
762     fi
763
764     fi
765
766     fi
767
768     fi
769
769
770     fi
771
772     fi
773
774     fi
775
776     fi
777
778     fi
779
779
780     fi
781
782     fi
783
784     fi
785
786     fi
787
788     fi
789
789
790     fi
791
792     fi
793
794     fi
795
796     fi
797
798     fi
799
799
800     fi
801
802     fi
803
804     fi
805
806     fi
807
808     fi
809
809
810     fi
811
812     fi
813
814     fi
815
816     fi
817
818     fi
819
819
820     fi
821
822     fi
823
824     fi
825
826     fi
827
828     fi
829
829
830     fi
831
832     fi
833
834     fi
835
836     fi
837
838     fi
839
839
840     fi
841
842     fi
843
844     fi
845
846     fi
847
848     fi
849
849
850     fi
851
852     fi
853
854     fi
855
856     fi
857
858     fi
859
859
860     fi
861
862     fi
863
864     fi
865
866     fi
867
868     fi
869
869
870     fi
871
872     fi
873
874     fi
875
876     fi
877
878     fi
879
879
880     fi
881
882     fi
883
884     fi
885
886     fi
887
888     fi
889
889
890     fi
891
892     fi
893
894     fi
895
896     fi
897
898     fi
899
899
900     fi
901
902     fi
903
904     fi
905
906     fi
907
908     fi
909
909
910     fi
911
912     fi
913
914     fi
915
916     fi
917
918     fi
919
919
920     fi
921
922     fi
923
924     fi
925
926     fi
927
928     fi
929
929
930     fi
931
932     fi
933
934     fi
935
936     fi
937
938     fi
939
939
940     fi
941
942     fi
943
944     fi
945
946     fi
947
948     fi
949
949
950     fi
951
952     fi
953
954     fi
955
956     fi
957
958     fi
959
959
960     fi
961
962     fi
963
964     fi
965
966     fi
967
968     fi
969
969
970     fi
971
972     fi
973
974     fi
975
976     fi
977
978     fi
979
979
980     fi
981
982     fi
983
984     fi
985
986     fi
987
988     fi
989
989
990     fi
991
992     fi
993
994     fi
995
996     fi
997
998     fi
999
999
1000    fi
1001
1002    fi
1003
1004    fi
1005
1006    fi
1007
1008    fi
1009
1009
1010   fi
1011
1012   fi
1013
1014   fi
1015
1016   fi
1017
1018   fi
1019
1019
1020   fi
1021
1022   fi
1023
1024   fi
1025
1026   fi
1027
1028   fi
1029
1029
1030   fi
1031
1032   fi
1033
1034   fi
1035
1036   fi
1037
1038   fi
1039
1039
1040   fi
1041
1042   fi
1043
1044   fi
1045
1046   fi
1047
1048   fi
1049
1049
1050   fi
1051
1052   fi
1053
1054   fi
1055
1056   fi
1057
1058   fi
1059
1059
1060   fi
1061
1062   fi
1063
1064   fi
1065
1066   fi
1067
1068   fi
1069
1069
1070   fi
1071
1072   fi
1073
1074   fi
1075
1076   fi
1077
1078   fi
1079
1079
1080   fi
1081
1082   fi
1083
1084   fi
1085
1086   fi
1087
1088   fi
1089
1089
1090   fi
1091
1092   fi
1093
1094   fi
1095
1096   fi
1097
1098   fi
1099
1099
1100  fi
1101
1102  fi
1103
1104  fi
1105
1106  fi
1107
1108  fi
1109
1109
1110 fi
1111
1112 fi
1113
1114 fi
1115
1116 fi
1117
1118 fi
1119
1119
1120 fi
1121
1122 fi
1123
1124 fi
1125
1126 fi
1127
1128 fi
1129
1129
1130 fi
1131
1132 fi
1133
1134 fi
1135
1136 fi
1137
1138 fi
1139
1139
1140 fi
1141
1142 fi
1143
1144 fi
1145
1146 fi
1147
1148 fi
1149
1149
1150 fi
1151
1152 fi
1153
1154 fi
1155
1156 fi
1157
1158 fi
1159
1159
1160 fi
1161
1162 fi
1163
1164 fi
1165
1166 fi
1167
1168 fi
1169
1169
1170 fi
1171
1172 fi
1173
1174 fi
1175
1176 fi
1177
1178 fi
1179
1179
1180 fi
1181
1182 fi
1183
1184 fi
1185
1186 fi
1187
1188 fi
1189
1189
1190 fi
1191
1192 fi
1193
1194 fi
1195
1196 fi
1197
1198 fi
1199
1199
1200 fi
1201
1202 fi
1203
1204 fi
1205
1206 fi
1207
1208 fi
1209
1209
1210 fi
1211
1212 fi
1213
1214 fi
1215
1216 fi
1217
1218 fi
1219
1219
1220 fi
1221
1222 fi
1223
1224 fi
1225
1226 fi
1227
1228 fi
1229
1229
1230 fi
1231
1232 fi
1233
1234 fi
1235
1236 fi
1237
1238 fi
1239
1239
1240 fi
1241
1242 fi
1243
1244 fi
1245
1246 fi
1247
1248 fi
1249
1249
1250 fi
1251
1252 fi
1253
1254 fi
1255
1256 fi
1257
1258 fi
1259
1259
1260 fi
1261
1262 fi
1263
1264 fi
1265
1266 fi
1267
1268 fi
1269
1269
1270 fi
1271
1272 fi
1273
1274 fi
1275
1276 fi
1277
1278 fi
1279
1279
1280 fi
1281
1282 fi
1283
1284 fi
1285
1286 fi
1287
1288 fi
1289
1289
1290 fi
1291
1292 fi
1293
1294 fi
1295
1296 fi
1297
1298 fi
1299
1299
1300 fi
1301
1302 fi
1303
1304 fi
1305
1306 fi
1307
1308 fi
1309
1309
1310 fi
1311
1312 fi
1313
1314 fi
1315
1316 fi
1317
1318 fi
1319
1319
1320 fi
1321
1322 fi
1323
1324 fi
1325
1326 fi
1327
1328 fi
1329
1329
1330 fi
1331
1332 fi
1333
1334 fi
1335
1336 fi
1337
1338 fi
1339
1339
1340 fi
1341
1342 fi
1343
1344 fi
1345
1346 fi
1347
1348 fi
1349
1349
1350 fi
1351
1352 fi
1353
1354 fi
1355
1356 fi
1357
1358 fi
1359
1359
1360 fi
1361
1362 fi
1363
1364 fi
1365
1366 fi
1367
1368 fi
1369
1369
1370 fi
1371
1372 fi
1373
1374 fi
1375
1376 fi
1377
1378 fi
1379
1379
1380 fi
1381
1382 fi
1383
1384 fi
1385
1386 fi
1387
1388 fi
1389
1389
1390 fi
1391
1392 fi
1393
1394 fi
1395
1396 fi
1397
1398 fi
1399
1399
1400 fi
1401
1402 fi
1403
1404 fi
1405
1406 fi
1407
1408 fi
1409
1409
1410 fi
1411
1412 fi
1413
1414 fi
1415
1416 fi
1417
1418 fi
1419
1419
1420 fi
1421
1422 fi
1423
1424 fi
1425
1426 fi
1427
1428 fi
1429
1429
1430 fi
1431
1432 fi
1433
1434 fi
1435
1436 fi
1437
1438 fi
1439
1439
1440 fi
1441
1442 fi
1443
1444 fi
1445
1446 fi
1447
1448 fi
1449
1449
1450 fi
1451
1452 fi
1453
1454 fi
1455
1456 fi
1457
1458 fi
1459
1459
1460 fi
1461
1462 fi
1463
1464 fi
1465
1466 fi
1467
1468 fi
1469
1469
1470 fi
1471
1472 fi
1473
1474 fi
1475
1476 fi
1477
1478 fi
1479
1479
1480 fi
1481
1482 fi
1483
1484 fi
1485
1486 fi
1487
1488 fi
1489
1489
1490 fi
1491
1492 fi
1493
1494 fi
1495
1496 fi
1497
1498 fi
1499
1499
1500 fi
1501
1502 fi
1503
1504 fi
1505
1506 fi
1507
1508 fi
1509
1509
1510 fi
1511
1512 fi
1513
1514 fi
1515
1516 fi
1517
1518 fi
1519
1519
1520 fi
1521
1522 fi
1523
1524 fi
1525
1526 fi
1527
1528 fi
1529
1529
1530 fi
1531
1532 fi
1533
1534 fi
1535
1536 fi
1537
1538 fi
1539
1539
1540 fi
1541
1542 fi
1543
1544 fi
1545
1546 fi
1547
1548 fi
1549
1549
1550 fi
1551
1552 fi
1553
1554 fi
1555
1556 fi
1557
1558 fi
1559
1559
1560 fi
1561
1562 fi
1563
1564 fi
1565
1566 fi
1567
1568 fi
1569
1569
1570 fi
1571
1572 fi
1573
1574 fi
1575
1576 fi
1577
1578 fi
1579
1579
1580 fi
1581
1582 fi
1583
1584 fi
1585
1586 fi
1587
1588 fi
1589
1589
1590 fi
1591
1592 fi
1593
1594 fi
1595
1596 fi
1597
1598 fi
1599
1599
1600 fi
1601
1602 fi
1603
1604 fi
1605
1606 fi
1607
1608 fi
1609
1609
1610 fi
1611
1612 fi
1613
1614 fi
1615
1616 fi
1617
1618 fi
1619
1619
1620 fi
1621
1622 fi
1623
1624 fi
1625
1626 fi
1627
1628 fi
1629
1629
1630 fi
1631
1632 fi
1633
1634 fi
1635
1636 fi
1637
1638 fi
1639
1639
1640 fi
1641
1642 fi
1643
1644 fi
1645
1646 fi
1647
1648 fi
1649
1649
1650 fi
1651
1652 fi
1653
1654 fi
1655
1656 fi
1657
1658 fi
1659
1659
1660 fi
1661
1662 fi
1663
1664 fi
1665
1666 fi
1667
1668 fi
1669
1669
1670 fi
1671
1672 fi
1673
1674 fi
1675
1676 fi
1677
1678 fi
1679
1679
1680 fi
1681
1682 fi
1683
1684 fi
1685
1686 fi
1687
1688 fi
1689
1689
1690 fi
1691
1692 fi
1693
1694 fi
1695
1696 fi
1697
1698 fi
1699
1699
1700 fi
1701
1702 fi
1703
1704 fi
1705
1706 fi
1707
1708 fi
1709
1709
1710 fi
1711
1712 fi
1713
1714 fi
1715
1716 fi
1717
1718 fi
1719
1719
1720 fi
1721
1722 fi
1723
1724 fi
1725
1726 fi
1727
1728 fi
1729
1729
1730 fi
1731
1732 fi
1733
1734 fi
1735
1736 fi
1737
1738 fi
1739
1739
1740 fi
1741
1742 fi
1743
1744 fi
1745
1746 fi
1747
1748 fi
1749
1749
1750 fi
1751
1752 fi
1753
1754 fi
1755
1756 fi
1757
1758 fi
1759
1759
1760 fi
1761
1762 fi
1763
1764 fi
1765
1766 fi
1767
1768 fi
1769
1769
1770 fi
1771
1772 fi
1773
1774 fi
1775
1776 fi
1777
1778 fi
1779
1779
1780 fi
1781
1782 fi
1783
1784 fi
1785
1786 fi
1787
1788 fi
1789
1789
1790 fi
1791
1792 fi
1793
1794 fi
1795
1796 fi
1797
1798 fi
1799
1799
1800 fi
1801
1802 fi
1803
1804 fi
1805
1806 fi
1807
1808 fi
1809
1809
1810 fi
1811
1812 fi
1813
1814 fi
1815
1816 fi
1817
1818 fi
1819
1819
1820 fi
1821
1822 fi
1823
1824 fi
1825
1826 fi
1827
1828 fi
1829
1829
1830 fi
1831
1832 fi
1833
1834 fi
1835
1836 fi
1837
1838 fi
1839
1839
1840 fi
1841
1842 fi
1843
1844 fi
1845
1846 fi
1847
1848 fi
1849
1849
1850 fi
1851
1852 fi
1853
1854 fi
1855
1856 fi
1857
1858 fi
1859
1859
1860 fi
1861
1862 fi
1863
1864 fi
1865
1866 fi
1867
1868 fi
1869
1869
1870 fi
1871
1872 fi
1873
1874 fi
1875
1876 fi
1877
1878 fi
1879
1879
1880 fi
1881
1882 fi
1883
1884 fi
1885
1886 fi
1887
1888 fi
1889
1889
1890 fi
1891
1892 fi
1893
1894 fi
1895
1896 fi
1897
1898 fi
1899
1899
1900 fi
1901
1902 fi
1903
1904 fi
1905
1906 fi
1907
1908 fi
1909
1909
1910 fi
1911
1912 fi
1913
1914 fi
1915
1916 fi
1917
1918 fi
1919
1919
1920 fi
1921
1922 fi
1923
1924 fi
1925
1926 fi
1927
1928 fi
1929
1929
1930 fi
1931
1932 fi
1933
1934 fi
1935
1936 fi
1937
1938 fi
1939
1939
1940 fi
1941
1942 fi
1943
1944 fi
1945
1946 fi
1947
1948 fi
1949
1949
1950 fi
1951
1952 fi
1953
1954 fi
1955
1956 fi
1957
1958 fi
1959
1959
1960 fi
1961
1962 fi
1963
1964 fi
1965
1966 fi
1967
1968 fi
1969
1969
1970 fi
1971
1972 fi
1973
1974 fi
1975
1976 fi
1977
1978 fi
1979
1979
1980 fi
1981
1982 fi
1983
1984 fi
1985
1986 fi
1987
1988 fi
1989
1989
1990 fi
1991
1992 fi
1993
1994 fi
1995
1996 fi
1997
1998 fi
1999
1999
2000 fi
2001
2002 fi
2003
2004 fi
2005
2006 fi
2007
2008 fi
2009
2009
2010 fi
2011
2012 fi
2013
2014 fi
2015
2016 fi
2017
2018 fi
2019
2019
2020 fi
2021
2022 fi
2023
2024 fi
2025
2026 fi
2027
2028 fi
2029
2029
2030 fi
2031
2032 fi
2033
2034 fi
2035
2036 fi
2037
2038 fi
2039
2039
2040 fi
2041
2042 fi
2043
2044 fi
2045
2046 fi
2047
2048 fi
2049
2049
2050 fi
2051
2052 fi
2053
2054 fi
2055
2056 fi
2057
2058 fi
20
```

```

209     echo "[INFO] Waiting for IP (xenstore)... ($attempt/
210         $max_attempts)"
211
212 if [[ -z "$VM_IP" ]]; then
213     VM_IP="UNKNOWN"
214     echo "[WARN] VM failed to report IP."
215 fi
216
217 # ----- DETACH CLOUD-INIT CD (non-disruptive)
218 # -----
218 VBD_CDS=$(remote "xe vbd-list vm-uuid=$VM_UUID type=CD --minimal"
219     | tr -d '[:space:]' || true)
220 if [[ -n "$VBD_CDS" ]]; then
221     IFS=',' read -ra VBD_ARR <<< "$VBD_CDS"
222     for vbd in "${VBD_ARR[@]}"; do
223         vbd=$(echo "$vbd" | tr -d '[:space:]')
224         if [[ -z "$vbd" || "$vbd" == "null" ]]; then continue; fi
225         remote_quiet "xe vbd-unplug uuid=$vbd || true"
226         remote_quiet "xe vbd-destroy uuid=$vbd || true"
227         echo "[INFO] Detached and destroyed CD VBD $vbd for VM
228             $VM_UUID"
229     done
230 fi
231
232 # Print summary and write IP to temp file for callers to pick up
233 echo "-----"
234 echo " VM CREATED SUCCESSFULLY"
235 echo " Name : $VM_NAME"
236 echo " UUID : $VM_UUID"
237 echo " IP : $VM_IP"
238 echo "-----"
239 echo "$VM_IP" > "/tmp/${SAFE_NAME}_ip.txt"
240 exit 0

```

## A.2 app/scheduler.py

Listing 10.2: app/scheduler.py

```

1 # app/scheduler.py
2 import subprocess

```

```
3 import shlex
4 import random
5 import logging
6 import time
7 from typing import Optional, Dict, Any
8
9 from app import models
10 from app.metrics_utils import calculate_host_score
11 from app.db import SessionLocal
12
13 logger = logging.getLogger(__name__)
14 logging.basicConfig(level=logging.INFO)
15
16
17 # ----- Host helpers -----
18 def host_is_overloaded(latest_metric) -> bool:
19     """
20         Simple threshold check for overloaded host.
21         (Matches checks described in the PDF final code.)
22     """
23     try:
24         cpu = float(latest_metric.cpu_percent or 0)
25         mem = float(latest_metric.mem_percent or 0)
26     except Exception:
27         return False
28     if cpu > 80.0 or mem > 85.0:
29         return True
30     return False
31
32
33 def select_least_loaded_host(db) -> models.Host:
34     """
35         Return the best host according to the final scoring logic.
36         Uses calculate_host_score from metrics_utils on the latest
37         HostMetric of each host.
38     """
39     hosts = db.query(models.Host).all()
40     if not hosts:
41         raise RuntimeError("No hosts configured in DB")
42     selected = []
```

```

43     logger.info("\n===== SCHEDULER START\n")
44     for host in hosts:
45         if not host.metrics:
46             logger.info(f"[NO METRICS] Skipping host {host.name}")
47             continue
48
49         latest = sorted(host.metrics, key=lambda m: m.ts)[-1]
50         if host_is_overloaded(latest):
51             logger.info(f"[SKIP] Host {host.name} overloaded
52                         CPU={latest.cpu_percent:.2f}%, MEM={latest.
53                         mem_percent:.2f}%")
54             continue
55
56         comp = calculate_host_score(latest)
57         logger.info(
58             f"[SCHEDULER] [HOST: {host.name}] CPU_raw: {latest.
59                 cpu_percent:.2f}% -> cpu_norm={comp['cpu_norm']:.3
60                 f} "
61             f"MEM_raw: {latest.mem_percent:.2f}% -> mem_norm={
62                 comp['mem_norm']:.3f} "
63             f"VM_count_raw: {latest.vms_running} -> vmc_norm={
64                 comp['vmc_norm']:.3f} FINAL SCORE: {comp['score
65                 ']:.5f}"
66         )
67         selected.append((comp["score"], host))
68
69     if not selected:
70         raise RuntimeError("No host is available under caps")
71
72     # sort ascending (lower score = less loaded)
73     selected.sort(key=lambda x: x[0])
74
75     # handle near-tie with a small randomization
76     if len(selected) > 1 and abs(selected[0][0] - selected[1][0]) < 0.05:
77         best = random.choice(selected[:2])[1]
78     else:
79         best = selected[0][1]
80
81
82
83
84
85
86
87
88
89
90
91
92
93

```

```

74     logger.info(f"\n[SELECTED HOST] {best.name} (UUID={best.
75         uuid_xen}, IP={best.ip})")
76     logger.info("===== SCHEDULER END
77         =====\n")
78
79 # ----- VM creation / scheduling -----
80 def schedule_vm_custom(db,
81                         user_cfg: Dict[str, Any],
82                         template_name: str,
83                         pool_id: str,
84                         script_path: str = "/app/app/
85                                         create_vm_remote_fixed.sh") -> Dict[str
86                                         , Any]:
87
88     """
89     Main scheduler entrypoint used by API.
90     - selects the best host
91     - builds command to call remote creation script
92     - runs the script (subprocess) and captures output
93     - records VM in DB if creation returns UUID/IP
94
95     Returns a dictionary with vm_name, vm_id, vm_ip,
96             selected_host, selected_host_ip, state, resources,
97             script_output
98
99     """
100
101    # 1) pick host
102    best = select_least_loaded_host(db)
103    logger.info(f"[Scheduler] Selected host: {best.name} ({best.
104        uuid_xen}) ip={best.ip}")
105
106    vm_name = user_cfg.get("name", f"Auto-VM-{best.name}")
107    vcpus = int(user_cfg.get("vcpus", 1))
108    ram = int(user_cfg.get("ram", 512))
109    disk_gib = int(user_cfg.get("disk", 10))
110    disk = f"{disk_gib}GiB"
111    network = user_cfg.get("network", "")
112    ssh_key = user_cfg.get("ssh_key", "/root/.ssh/id_rsa.pub")
113
114    if not best.ip:
115        raise RuntimeError("Selected host has no registered IP")
116
117

```

```

108     args = [
109         script_path,
110         "--host", str(best.ip),
111         "--name", vm_name,
112         "--template", template_name,
113         "--cpu", str(vcpus),
114         "--ram", str(ram),
115         "--disk", disk,
116         "--network", network,
117         "--ssh-key", ssh_key
118     ]
119     safe_cmd = " ".join(shlex.quote(a) for a in args)
120     logger.info(f"[Scheduler] EXECUTING: {safe_cmd}")
121
122     # 2) Execute VM creation script
123     try:
124         proc = subprocess.run(
125             args,
126             check=True,
127             stdout=subprocess.PIPE,
128             stderr=subprocess.STDOUT,
129             text=True,
130         )
131         output = proc.stdout or ""
132         logger.info(f"[Scheduler] Script finished. Output:\n{output}")
133     except subprocess.CalledProcessError as e:
134         output = e.output if hasattr(e, "output") else str(e)
135         logger.exception(f"[Scheduler] VM creation script failed: {output}")
136         raise
137
138     # 3) Parse output for VM UUID / IP (script writes summary and
139     #      writes IP to /tmp/<name>.ip.txt on host; we rely on
140     #      script prints)
141     vm_uuid = None
142     vm_ip = None
143     for line in (output or "").splitlines():
144         if "UUID :" in line or "UUID:" in line:
145             # tolerant parse
146             try:

```

```
145             vm_uuid = line.split("UUID")[1].split(":")[1].  
146                 strip()  
147         except Exception:  
148             pass  
149         if "IP :" in line or "IP:" in line:  
150             try:  
151                 vm_ip = line.split("IP")[1].split(":")[1].strip()  
152             except Exception:  
153                 pass  
154  
155 # 4) Save VM into DB if possible (best-effort)  
156 vm_record = None  
157 try:  
158     vm_record = models.VM(  
159         name=vm_name,  
160         uuid=vm_uuid or "",  
161         host_id=best.id,  
162         ip=vm_ip or None,  
163         memory_mb=ram,  
164         vcpus=vcpus,  
165         created_at=int(time.time())  
166     )  
167     db.add(vm_record)  
168     db.commit()  
169     db.refresh(vm_record)  
170     logger.info(f"[Scheduler] VM record created in DB id={  
171         vm_record.id}")  
172 except Exception as db_err:  
173     db.rollback()  
174     logger.exception(f"[Scheduler] Failed to persist VM  
175         record: {db_err}")  
176  
177 result = {  
178     "vm_name": vm_name,  
179     "vm_id": vm_record.id if vm_record else None,  
180     "vm_uuid": vm_uuid,  
181     "vm_ip": vm_ip,  
182     "selected_host": best.name,  
183     "selected_host_ip": best.ip,  
184     "state": "created" if vm_uuid else "error",
```

```
182     "resources": {"vcpus": vcpus, "ram_mb": ram, "disk": disk
183         },
184     "script_output": output,
185 }
186
187
188 # ----- Migration engine -----
189 def migrate_vm(db) -> Optional[Dict[str, Any]]:
190     """
191     Final migration function: compute scores for all hosts, and
192     migrate smallest VM from the worst host to the least
193     loaded host
194     """
195
196     logger.info("\n===== MIGRATION CHECK\n=====\\n")
197     hosts = db.query(models.Host).all()
198     if len(hosts) < 2:
199         logger.info("[MIGRATE] Need at least 2 hosts")
200         return None
201
202     scored = []
203     # Score each host based on latest metric
204     for host in hosts:
205         if not host.metrics:
206             logger.info(f"[MIGRATE] Host {host.name} has no
207                         metrics      skipping")
208             continue
209         latest = sorted(host.metrics, key=lambda x: x.ts)[-1]
210         score_data = calculate_host_score(latest)
211         score = score_data["score"]
212         logger.info(
213             f"[HOST] {host.name}: CPU={latest.cpu_percent:.3f}%
214               MEM={latest.mem_percent:.3f}% VMs={latest.
215                 vms_running}      SCORE={score:.3f}"
216         )
217         scored.append((score, host, latest))
218
219     if len(scored) < 2:
220         logger.info("[MIGRATE] Not enough metric data")
221         return None
```

```

216
217     # sort highest (most loaded) first
218     scored.sort(reverse=True, key=lambda x: x[0])
219     overloaded_score, src, src_metric = scored[0]
220     under_score, dest, dest_metric = scored[-1]
221
222     logger.info(f"[MIGRATE] Worst host: {src.name} (Score={
223         overloaded_score:.3f})")
224     logger.info(f"[MIGRATE] Best host: {dest.name} (Score={
225         under_score:.3f})")
226
227     if overloaded_score - under_score < 0.15:
228         logger.info("[MIGRATE] Difference < 0.15      No migration
229             required")
230         return None
231
232     # Pick smallest VM on overloaded host
233     vm = (
234         db.query(models.VM)
235         .filter(models.VM.host_id == src.id)
236         .order_by(models.VM.memory_mb.asc())
237         .first()
238     )
239     if not vm:
240         logger.info("[MIGRATE] No VM available on overloaded host
241             ")
242         return None
243
244     logger.info(f"[MIGRATE] Selected VM      {vm.name} ({vm.uuid})")
245
246     # perform migration (xe vm-migrate or equivalent)
247     cmd = f"xe vm-migrate vm={vm.uuid} host={dest.uuid_xen} live=
248         true"
249     logger.info(f"[MIGRATE] Running: {cmd}")
250     try:
251         subprocess.run(shlex.split(cmd), check=True)
252         logger.info("[MIGRATE] Migration triggered")
253     except subprocess.CalledProcessError as e:
254         logger.exception(f"[MIGRATE] Migration command failed: {e
255             }")

```

```
250     return {"error": str(e)}
```

```
251
```

```
252     return {"migrated_vm": vm.uuid, "source_host": src.name, "
```

```
253         target_host": dest.name}
```

```
254
```

```
255 def auto_migrate() -> Dict[str, Any]:
```

```
256     """
```

```
257     Helper to be called by Celery tasks: opens DB session, calls
```

```
258         migrate_vm, and closes session.
```

```
259     """
```

```
260     db = SessionLocal()
```

```
261     try:
```

```
262         result = migrate_vm(db)
```

```
263         if result:
```

```
264             logger.info(f" [MIGRATION] Success {result}")
265             return {"status": "migrated", "result": result}
```

```
266         else:
267             logger.info(" [MIGRATION] No migration required")
268             return {"status": "checked", "result": result}
```

```
269     except Exception as e:
270         logger.exception(f" [MIGRATION] Migration failed: {e}")
271         return {"status": "error", "message": str(e)}
```

```
272     finally:
273         db.close()
```

```
274
```

```
275 # Allow manual invocation for quick tests (not executed in import
276 #)
277 if __name__ == "__main__":
278     db = SessionLocal()
279     try:
280         print("== manual scheduler test ==")
281         best = select_least_loaded_host(db)
282         print("Selected host:", best.name)
283     finally:
284         db.close()
```

### A.3 app/models.py

Listing 10.3: app/models.py

```

1  from sqlalchemy import Column, Integer, String, Float, ForeignKey
2  from sqlalchemy.orm import relationship
3  from sqlalchemy.ext.declarative import declarative_base
4  import time
5
6  Base = declarative_base()
7
8  class Host(Base):
9      __tablename__ = "hosts"
10     id = Column(Integer, primary_key=True, index=True)
11     name = Column(String, unique=True, nullable=False)
12     uuid_xen = Column(String, unique=True, nullable=True)
13     ip = Column(String, nullable=True)
14     metrics = relationship("HostMetric", back_populates="host",
15                           cascade="all, delete-orphan")
15     vms = relationship("VM", back_populates="host", cascade="all,
16                         delete-orphan")
16
17 class HostMetric(Base):
18     __tablename__ = "host_metrics"
19     id = Column(Integer, primary_key=True, index=True)
20     host_id = Column(Integer, ForeignKey("hosts.id"), nullable=
21                      False)
22     cpu_percent = Column(Float, default=0.0)
23     mem_percent = Column(Float, default=0.0)
24     vms_running = Column(Integer, default=0)
25     ts = Column(Integer, default=lambda: int(time.time()))
26     host = relationship("Host", back_populates="metrics")
27
27 class VM(Base):
28     __tablename__ = "vms"
29     id = Column(Integer, primary_key=True, index=True)
30     name = Column(String, nullable=False)
31     uuid = Column(String, unique=True, nullable=True)
32     host_id = Column(Integer, ForeignKey("hosts.id"), nullable=
33                      True)
34     ip = Column(String, nullable=True)
35     memory_mb = Column(Integer, default=0)

```

```

35     vcpus = Column(Integer, default=1)
36     created_at = Column(Integer, default=lambda: int(time.time()))
37         )
38
39 class Job(Base):
40     __tablename__ = "jobs"
41     id = Column(Integer, primary_key=True, index=True)
42     type = Column(String, nullable=False)
43     payload = Column(String, nullable=True)
44     status = Column(String, default="pending")
45     created_at = Column(Integer, default=lambda: int(time.time()))
        )

```

““latex

## A.4 app/db.py

Listing 10.4: app/db.py

```

1 from sqlalchemy import create_engine
2 from sqlalchemy.orm import sessionmaker
3
4 DATABASE_URL = "postgresql://postgres@example@postgres:5432/cloud
      "
5
6 engine = create_engine(DATABASE_URL)
7 SessionLocal = sessionmaker(bind=engine, autoflush=False,
     autocommit=False)
8
9 def get_db():
10     db = SessionLocal()
11     try:
12         yield db
13     finally:
14         db.close()

```

## A.5 app/metrics\_utils.py

Listing 10.5: app/metricsUtils.py

```

1 def calculate_host_score(metric):

```

```

2      """
3      Convert CPU%, MEM%, VMs running      normalized weights
4          final score.
5      """
6
7      cpu = float(metric.cpu_percent or 0)    # 0 100
8      mem = float(metric.mem_percent or 0)    # 0 100
9      vms = int(metric.vms_running or 0)      # integer
10
11
12     # Normalize
13     cpu_norm = cpu / 100.0
14     mem_norm = mem / 100.0
15     vmc_norm = min(vms / 10.0, 1.0)    # 10 VMs = 100% weight
16
17     # Weighted load score
18     score = (cpu_norm * 0.5) + (mem_norm * 0.3) + (vmc_norm *
19         0.2)
20
21     return {
22         "cpu_norm": cpu_norm,
23         "mem_norm": mem_norm,
24         "vmc_norm": vmc_norm,
25         "score": score,
26     }

```

## A.6 app/main.py

Listing 10.6: app/main.py

```

1  from fastapi import FastAPI
2  from app.api.api_vm import router as vm_router
3  from app.api.api_metrics import router as metrics_router
4  from app.api.api_migration import router as migration_router
5
6  app = FastAPI(title="Mini Cloud Platform")
7
8  app.include_router(vm_router)
9  app.include_router(metrics_router)
10 app.include_router(migration_router)
11
12 @app.get("/")
13 def root():
14     return {"message": "Mini Cloud Platform Controller"}

```

## A.7 app/api/apiVm.py

Listing 10.7: app/api/apiVm.py

```

1  from fastapi import APIRouter, Depends, HTTPException
2  from sqlalchemy.orm import Session
3  from pydantic import BaseModel, Field
4  from app.db import get_db
5  from app.scheduler import schedule_vm_custom
6
7  router = APIRouter(prefix="/vm", tags=["VM"])
8
9  class VMCreateRequest(BaseModel):
10     name: str = Field(..., example="Test-VM-1")
11     vcpus: int = Field(..., example=2)
12     ram: int = Field(..., example=2048)
13     disk: int = Field(..., example=10)
14     network: str = Field(..., example="Pool-wide network 1")
15     ssh_key: str = Field(..., example="/root/.ssh/id_rsa.pub")
16     template: str = Field("Ubuntu-24.04-Base", example="Ubuntu
17         -24.04-Base")
18
19 @router.post("/create")
20 def create_vm(payload: VMCreateRequest, db: Session = Depends(
21     get_db)):
22     TEMPLATE_NAME = payload.template or "Ubuntu-24.04-Base"
23     POOL_ID = "b4bd9096-7d00-aeee-c35c-16bfba69b778" # adjust to
24         your pool id if required
25
26     try:
27         result = schedule_vm_custom(
28             db=db,
29             user_cfg=payload.dict(),
30             template_name=TEMPLATE_NAME,
31             pool_id=POOL_ID,
32         )
33         return {"status": "success", "details": result}
34     except Exception as e:
35         raise HTTPException(status_code=500, detail=str(e))

```

## A.8 app/api/apiMigration.py

Listing 10.8: app/api/apiMigration.py

```

1 from fastapi import APIRouter, HTTPException
2 from app.db import SessionLocal
3 from app.scheduler import migrate_vm
4
5 router = APIRouter(prefix="/migration", tags=["Migration"])
6
7 @router.post("/manual")
8 def manual_migration():
9     db = SessionLocal()
10    try:
11        result = migrate_vm(db)
12        return {"status": "success", "result": result}
13    except Exception as e:
14        raise HTTPException(status_code=500, detail=str(e))
15    finally:
16        db.close()

```

## A.9 app/api/apiMetrics.py

Listing 10.9: app/api/apiMetrics.py

```

1 from fastapi import APIRouter, Depends, HTTPException
2 from sqlalchemy.orm import Session
3 from pydantic import BaseModel, Field
4 from app.db import get_db
5 from app import models
6 import time
7
8 router = APIRouter(prefix="/metrics", tags=["Metrics"])
9
10 class MetricsSchema(BaseModel):
11     host_name: str
12     cpu_percent: float
13     mem_percent: float
14     vms_running: int
15
16 @router.post("/report")

```

```

17 def report_metrics(payload: MetricsSchema, db: Session = Depends(
18     get_db)):
19     host = db.query(models.Host).filter(models.Host.name ==
20         payload.host_name).first()
21     if not host:
22         # create if absent
23         host = models.Host(name=payload.host_name, ip=None)
24         db.add(host)
25         db.commit()
26         db.refresh(host)
27     metric = models.HostMetric(
28         host_id = host.id,
29         cpu_percent = payload.cpu_percent,
30         mem_percent = payload.mem_percent,
31         vms_running = payload.vms_running,
32         ts = int(time.time())
33     )
34     db.add(metric)
35     db.commit()
36     return {"status": "ok"}

```

## A.10 app/tasks/celeryApp.py

Listing 10.10: app/tasks/celeryApp.py

```

1 from celery import Celery
2
3 celery_app = Celery(
4     "controller",
5     broker="redis://redis:6379/0",
6     backend="redis://redis:6379/1"
7 )
8
9 celery_app.conf.task_routes = {
10     "app.tasks.*": {"queue": "default"},
11 }
12
13 celery_app.conf.timezone = "UTC"

```

## A.11 app/tasks/jobs.py

Listing 10.11: app/tasks/jobs.py

```

1  from app.tasks.celery_app import celery_app
2  from app.tasks.collect_metrics import collect_metrics
3  from app.scheduler import auto_migrate, migrate_vm
4  from app.db import SessionLocal
5  import logging
6
7  logger = logging.getLogger(__name__)
8
9  @celery_app.task(name="app.tasks.jobs.collect_metrics_job")
10 def collect_metrics_job():
11     try:
12         collect_metrics()
13         return {"status": "success"}
14     except Exception as e:
15         logger.exception("collect_metrics_job failed")
16         return {"status": "error", "error": str(e)}
17
18 @celery_app.task(name="app.tasks.jobs.migrate_hosts_job")
19 def migrate_hosts_job():
20     try:
21         print("[MIGRATION] Checking if migration is needed ")
22         result = auto_migrate()
23         print("[MIGRATION] Result:", result)
24         return result
25     except Exception as e:
26         print("[MIGRATION] ERROR:", e)
27         return {"error": str(e)}
28
29 # -----
30 # CELERY PERIODIC TASK SCHEDULING
31 # -----
32 celery_app.conf.beat_schedule = {
33     "collect-host-metrics-every-2-minutes": {
34         "task": "app.tasks.jobs.collect_metrics_job",
35         "schedule": 120,
36     },
37     "migration-check-every-2-minutes": {
38         "task": "app.tasks.jobs.migration_job",

```

```

39         "schedule": 120,
40     },
41 }
42
43 @celery_app.task(name="app.tasks.jobs.migration_job")
44 def migration_job():
45     from app.db import SessionLocal    # lazy import to avoid
46             circular deps
47     print("\n===== AUTO MIGRATION CHECK
48             =====")
49     db = SessionLocal()
50     try:
51         result = migrate_vm(db)
52         if result:
53             print(f"[AUTO-MIGRATE] Migration executed {result
54             }")
55         else:
56             print("[AUTO-MIGRATE] No migration required at this
57                   cycle")
58         return {"status": "checked", "result": result}
59     except Exception as e:
60         print(f"[AUTO-MIGRATE] ERROR {e}")
61         return {"status": "error", "message": str(e)}
62     finally:
63         db.close()

```

## A.12 app/tasks/collectMetrics.py

Listing 10.12: app/tasks/collectMetrics.py

```

1  from app.db import SessionLocal
2  from app.metrics_utils import calculate_host_score
3  from app.xoa_client import fetch_live_metrics
4  from app import models
5  import time
6
7  def collect_metrics():
8      db = SessionLocal()
9      try:
10          hosts = db.query(models.Host).all()
11          for host in hosts:

```

```

12         # fetch_live_metrics should reach XOA or other metric
13         source
14
15         data = fetch_live_metrics(host.ip)
16
17         metric = models.HostMetric(
18             host_id = host.id,
19             cpu_percent = data.get("cpu", 0.0),
20             mem_percent = data.get("memory", 0.0),
21             vms_running = data.get("vms", 0),
22             ts = int(time.time())
23         )
24
25         db.add(metric)
26
27         db.commit()
28
29     finally:
30
31         db.close()

```

### A.13 app/xoaClient.py

Listing 10.13: app/xoaClient.py

```

1 import requests
2 import os
3
4 XOA_URL = os.environ.get("XOA_URL", "http://10.20.24.77")
5 XOA_TOKEN = os.environ.get("XOA_TOKEN", "")
6
7 def _headers():
8     return {"Authorization": f"Bearer {XOA_TOKEN}"} if XOA_TOKEN
9     else {}
10
11 def fetch_live_metrics(host_ip):
12     """
13         Example: call to XOA or to a simple per-host metric endpoint.
14         This function should be adapted to your actual XOA endpoints.
15     """
16
17     try:
18         # If you have XOA metrics endpoint, replace with real
19         # path
20         url = f"{XOA_URL}/api/hosts/{host_ip}/metrics"
21         resp = requests.get(url, headers=_headers(), timeout=5)
22         if resp.status_code == 200:
23             return resp.json()

```

```

21     except Exception:
22         pass
23     # fallback dummy metrics for robust execution
24     return {"cpu": 0.0, "memory": 0.0, "vms": 0}

```

## A.14 Miscellaneous utilities

Listing 10.14: app/utils.py

```

1 # small utilities used by different modules
2 def safe_get(d, key, default=None):
3     try:
4         return d.get(key, default)
5     except Exception:
6         return default

```

## A.15 docker-compose.yml

Listing 10.15: docker-compose.yml

```

1 version: "3.9"
2 services:
3
4     api:
5         build:
6             context: .
7             dockerfile: Dockerfile
8         command: uvicorn app.main:app --host 0.0.0.0 --port 8001 --
9             reload
10        ports:
11            - "8001:8001"
12        depends_on:
13            - redis
14            - postgres
15
16        celery_worker:
17            build:
18                context: .
19                dockerfile: Dockerfile
20            command: celery -A app.tasks.celery_app worker --loglevel=
21                INFO

```

```

20     depends_on:
21         - redis
22         - postgres
23
24     celery_beat:
25         build:
26             context: .
27             dockerfile: Dockerfile
28             command: celery -A app.tasks.celery_app beat --loglevel=INFO
29         depends_on:
30             - redis
31             - postgres
32
33     redis:
34         image: redis:6
35         restart: unless-stopped
36
37     postgres:
38         image: postgres:14
39         environment:
40             POSTGRES_PASSWORD: example
41         volumes:
42             - pgdata:/var/lib/postgresql/data
43
44     volumes:
45         pgdata:

```

## A.16 Dockerfile (example)

Listing 10.16: Dockerfile

```

1 FROM python:3.11-slim
2
3 WORKDIR /app
4 COPY requirements.txt /app/requirements.txt
5 RUN pip install --no-cache-dir -r /app/requirements.txt
6
7 COPY . /app
8 ENV PYTHONUNBUFFERED=1
9
10 # default entrypoint is overridden by docker-compose commands

```